

Parallel Computing with GPUs: OpenMP

Dr Mozhgan Kabiri Chimeh

<http://mkchimeh.staff.shef.ac.uk/teaching/COM4521>



The
University
Of
Sheffield.



GPU
RESEARCH
CENTER


Last Lecture

- ❑ We looked at how to make programs fast on a single core
- ❑ **But** we didn't consider parallelism
- ❑ Guess what we are going today?

- ❑ Multicore systems and OpenMP
- ❑ Parallelising Loops
- ❑ Critical Sections and Synchronisation
- ❑ Scoping
- ❑ Data vs Task Parallelism

Multicore systems

- ❑ Multi-core CPUs are a shared memory system
 - ❑ Each CPU has access to main memory
 - ❑ Each CPU can access all of the memory (hence shared)
 - ❑ Each CPU cores have their own L2 and L3 cache
 - ❑ This can cause a lack of coherence
 - ❑ If one core modifies its cache this might need to be synchronised on other cores




OpenMP

❑ Open Multi-Processing Standard

- ❑ An API that supports shared memory programming in C, C++ and FORTRAN
- ❑ Cross platform support using native threading
 - ❑ Higher level than OS models and portable
- ❑ Is not suitable for distributed computing (look at MPI)

❑ It is not an automatic parallel programming language

- ❑ Parallelism is explicitly defined and controlled by the programmer
- ❑ Requires compiler directives, a runtime, environment variables



OpenMP Compiler Directives

❑ Use of #pragmas

- ❑ If not understood by the compiler then they are ignored
- ❑ Does not require serial code to be changed
- ❑ Allows behaviour to be specified which are not part of the C specification

```
#include <stdio.h>
#include <omp.h>

int main()
{
    #pragma omp parallel
    {
        printf("Hello World\n");
    }
    return 0;
}
```

Extending OpenMP Hello World


```
#include <stdio.h>
#include <omp.h>

int main()
{
    #pragma omp parallel
    {
        int thread = omp_get_thread_num();
        int max_threads = omp_get_max_threads();
        printf("Hello World (Thread %d of %d)\n", thread, max_threads);
    }
    return 0;
}
```


```
Hello World (Thread 5 of 8)
Hello World (Thread 6 of 8)
Hello World (Thread 2 of 8)
Hello World (Thread 7 of 8)
Hello World (Thread 1 of 8)
Hello World (Thread 0 of 8)
Hello World (Thread 3 of 8)
Hello World (Thread 4 of 8)
```

Fork and Join


- ❑ OpenMP uses a fork a join model
 - ❑ Fork: Creates a number of parallel threads from a master thread
 - ❑ Master thread is always thread 0
 - ❑ No guarantee of order
 - ❑ Join: Synchronises thread termination and returns program control to master



Terminology



Terminology



- ❑ Multicore systems and OpenMP
- ❑ Parallelising Loops
- ❑ Critical Sections and Synchronisation
- ❑ Scoping
- ❑ Data vs Task Parallelism

OpenMP Syntax

❑ Parallel region directive

- ❑ `#pragma omp parallel [clause list] {structured block}`

- ❑ Spawns a number of parallel threads

❑ Clauses

- ❑ Are used to specify modifications to the parallel directive e.g.

- ❑ Control scoping of variables in multiple threads

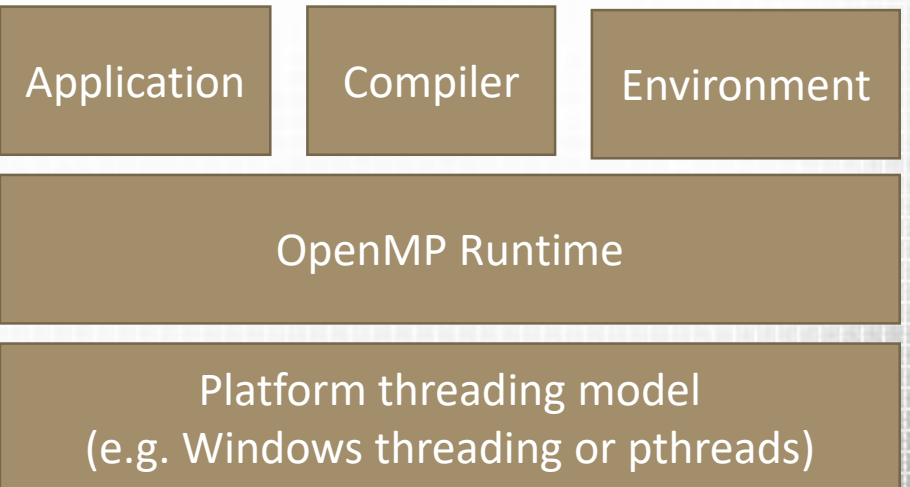
- ❑ Dictate the number of parallel threads (example below)

- ❑ Conditional parallelism

```
#pragma omp parallel num_threads(16)
{
    int thread = omp_get_thread_num();
    int max_threads = omp_get_max_threads();
    printf("Hello World (Thread %d of %d)\n", thread, max_threads);
}
```

num_threads()

- ❑ Without this clause `OMP_NUM_THREADS` will be used
 - ❑ This is an environment variable
 - ❑ Set to the number of cores (or hyperthreads) on your machine
 - ❑ This can be set globally by `omp_set_num_threads(int)`
 - ❑ Value can be queried by `int omp_get_num_threads();`
- ❑ `num_threads` takes precedence over the environment variable
- ❑ `num_threads()` does not guarantee that the number requested will be created
 - ❑ System limitations may prevent this
 - ❑ However: It almost always will





parallel for

- ❑ #pragma omp for
 - ❑ Assigns work units to the team
 - ❑ Divides loop iterations between threads
- ❑ For can be combined e.g. #pragma omp parallel for
 - ❑ Threads are spawned and then assigned to loop iterations

```
int n;  
#pragma omp parallel for  
for (n = 0; n < 8; n++){  
    int thread = omp_get_thread_num();  
    printf("Parallel thread %d \n", thread);  
}
```

```
#pragma omp parallel  
{  
    int n;  
    for (n = 0; n < 8; n++){  
        int thread = omp_get_thread_num();  
        printf("Parallel thread %d \n", thread);  
    }  
}
```

```
#pragma omp parallel  
{  
    int n;  
    #pragma omp for  
    for (n = 0; n < 8; n++){  
        int thread = omp_get_thread_num();  
        printf("Parallel thread %d \n", thread);  
    }  
}
```

Which is the odd one out?

parallel for

- ❑ `#pragma omp for`
 - ❑ Assigns work units to the team
 - ❑ Divides loop iterations between threads
- ❑ For can be combined e.g. `#pragma omp for`
 - ❑ Threads are spawned and then assigned

```
#pragma omp parallel
{
    int n;
    for (n = 0; n < 8; n++){
        int thread = omp_get_thread_num();
        printf("Parallel thread %d \n", thread);
    }
}
```

```
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 5
Parallel thread 5
Parallel thread 5
Parallel thread 5
Parallel thread 4
Parallel thread 4
Parallel thread 3
Parallel thread 3
Parallel thread 1
...
```

- ❑ Multicore systems and OpenMP
- ❑ Parallelising Loops
- ❑ Critical Sections and Synchronisation
- ❑ Scoping
- ❑ Data vs Task Parallelism



What is wrong with this code?

❑ Consider a problem such as Taylor series expansion for \cos function

$$\square \cos(x) = \sum_{n=0}^{\infty} (-1)^n \frac{x^{2n}}{(2n)!}$$

$$\square \cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} \dots$$

```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for  
for (n = 0; n < EXPANSION_STEPS; n++) {  
    double r = pow(-1, n) * pow(x, 2 * n) / fac(2 * n);  
    result += r;  
}  
  
printf("Approximation is %f, value is %f\n", result, cos(x));
```


Critical sections

❑ Consider a problem such as Taylor series expansion for *cos* function

$$\square \cos(x) = \sum_{n=0}^{\infty} (-1)^{n-1} \frac{x^{2n-1}}{(2n)!}$$

$$\square \cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} \dots$$

```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for  
for (n = 0; n < EXPANSION_STEPS; n++) {  
    double r = pow(-1, n - 1) * pow(x, 2 * n - 1) / fac(2 * n);  
    result -= r;  
}  
  
printf("Approximation is %f, value is %f\n", result, cos(x));
```

Multiple threads try to write to the same value!
(undefined behaviour and unpredictable results)

Critical sections

❑ Consider a problem such as Taylor series expansion for *cos* function

$$\square \cos(x) = \sum_{n=0}^{\infty} (-1)^n \frac{x^{2n}}{(2n)!}$$

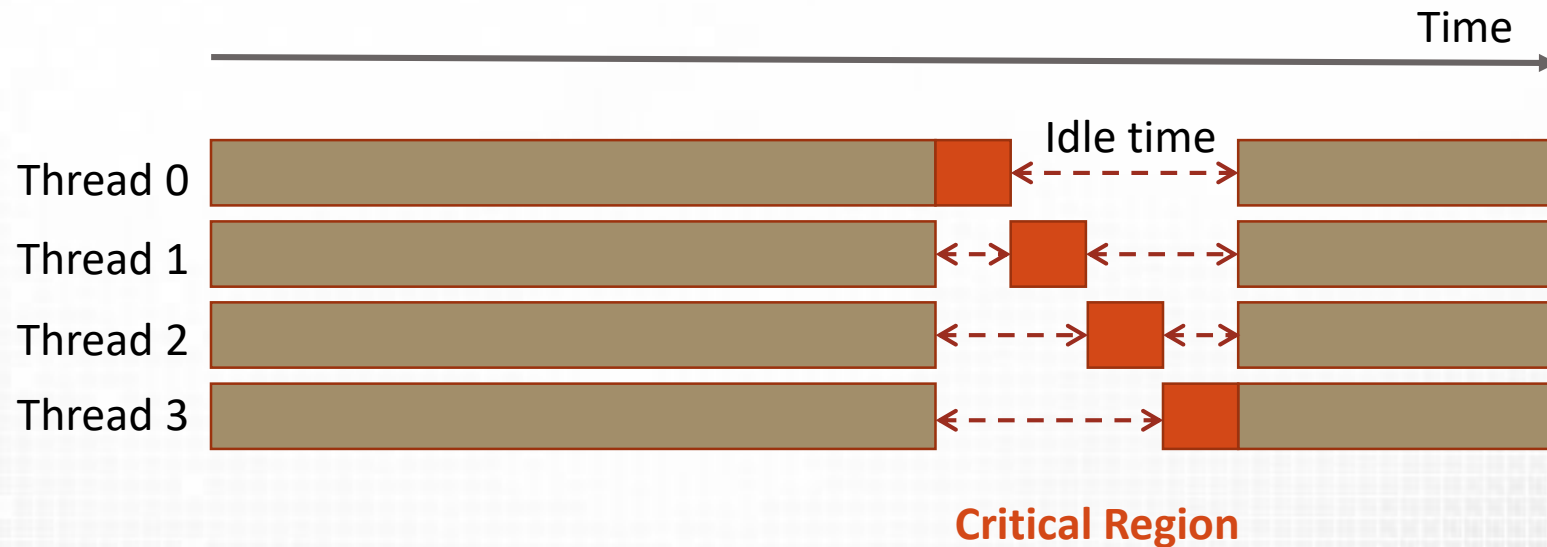
$$\square \cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} \dots$$

```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for  
for (n = 0; n < EXPANSION_STEPS; n++) {  
    double r = pow(-1, n) * pow(x, 2 * n) / fac(2 * n);  
    #pragma omp critical  
    {  
        result -= r;  
    }  
}  
  
printf("Approximation is %f, value is %f\n", result, cos(x));
```

Define as a critical section

Critical sections

- ❑ `#pragma omp critical [name]`
 - ❑ Ensures mutual exclusions when accessing a shared value
 - ❑ Prevents race conditions
 - ❑ A thread will wait until no other thread is executing a critical region (with the same name) before beginning
 - ❑ Unnamed critical regions map to the same unspecified name



Atomics

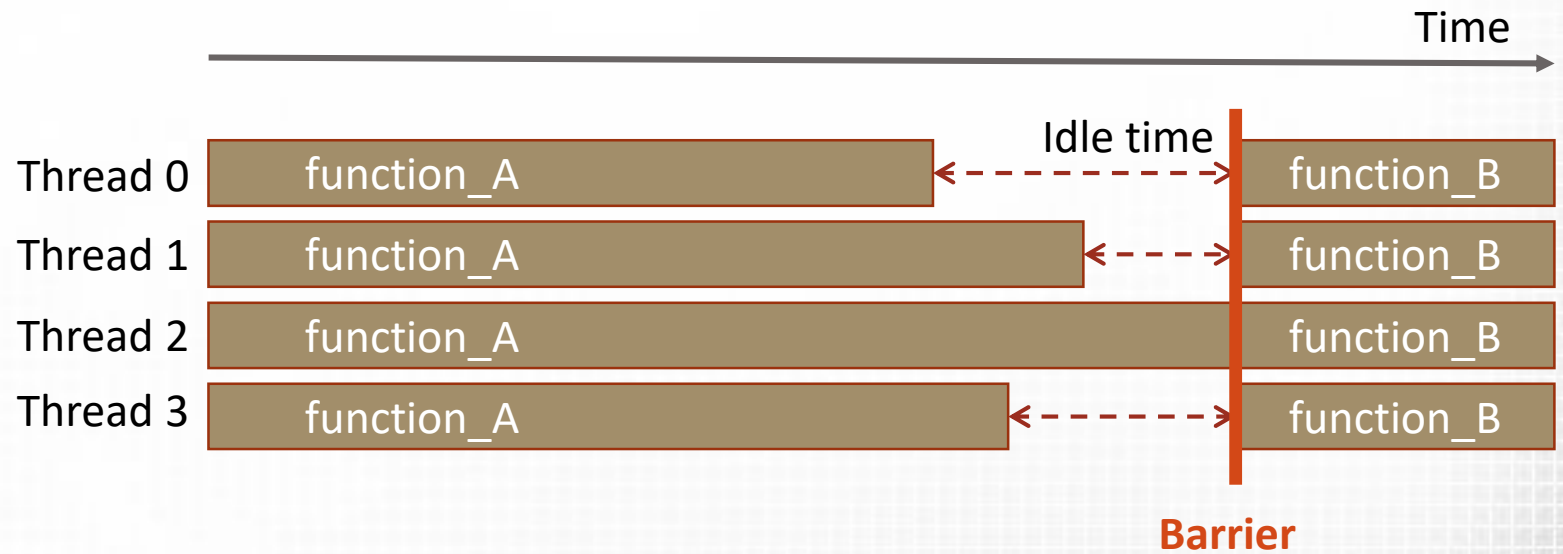
- ❑ Atomic operations can be used to safely increment a shared numeric value
 - ❑ For example summation
 - ❑ Atomics only apply to the immediate assignment
- ❑ Atomics are usually faster than critical sections
 - ❑ Critical sections can be applied to general blocks of code (atomics can not)
- ❑ Example
 - ❑ Compute histogram of random values for a given range
 - ❑ Random is an `int` array of size `NUM_VALUES` with random value within `0 : RANGE`
 - ❑ Histogram is an `int` array of size `RANGE` with 0 values;

```
#pragma omp parallel
{
    int i;
    #pragma omp for
    for (i = 0; i < NUM_VALUES; i++){
        int value = randoms[i];
        #pragma omp atomic
        histogram[value]++;
    }
}
```

Barriers

- ❑ `#pragma omp barrier`
 - ❑ Synchronises threads at a barrier point
 - ❑ Parallel regions have an implicit barrier
 - ❑ Can be used to ensure execution of particular code is complete
 - ❑ E.g. data read by `function_B`

```
#pragma omp parallel
{
    function_A()
    #pragma omp barrier
    function_B();
}
```



Single and Master Sections

- ❑ `#pragma omp single { ... }`
 - ❑ Used to ensure that only a single thread executes a region of a structured block
 - ❑ Useful for I/O and initialisation
 - ❑ First available thread will execute the defined region
 - ❑ No control over which this is
 - ❑ Will cause an implicit barrier unless a `nowait` clause is used
 - ❑ E.g. `#pragma omp single nowait`
 - ❑ `nowait` will remove an implied barrier and can also be applied to parallel for loops
- ❑ `#pragma omp master { ... }`
 - ❑ Similar to `single` but will always use the master thread
 - ❑ Is equivalent to using an `IF` clause
 - ❑ E.g. `#pragma omp parallel IF(omp_get_thread_num() == 0) nowait`
 - ❑ The `IF` clause makes the spawning of parallel threads conditional
 - ❑ Preferable to `single`
 - ❑ Does not require an implicit barrier

Master example

```
int t, r;
int local_histogram[THREADS][RANGE];

zero_histogram(local_histogram);

#pragma omp parallel num_threads(THREADS)
{
    int i;
    #pragma omp for
    for (i = 0; i < NUM_VALUES; i++){
        int value = randoms[i];
        local_histogram[omp_get_thread_num()][value]++;
    }
    #pragma omp barrier
    #pragma omp master
    for (t = 0; t < THREADS; t++){
        for (r = 0; r < RANGE; r++){
            histogram[r] += local_histogram[t][r];
        }
    }
}
```

Same result as the atomic version

- ❑ Multicore systems and OpenMP
- ❑ Parallelising Loops
- ❑ Critical Sections and Synchronisation
- ❑ Scoping
- ❑ Data vs Task Parallelism

Scoping

- ❑ Scope refers to the part of the program in which a variable can be used
- ❑ OpenMP has different scoping to serial programming
 - ❑ We must define if a variable is private or shared between threads
- ❑ **Shared:** A variable can be accessed by all threads in the team
 - ❑ All variables outside of a parallel loop are shared by default
- ❑ **Private:** A Variable is local to a single thread and can only be accessed by this thread within the structured block it is defined
 - ❑ All variables inside a structured block are private by default

Scoping

```
int t, r;
int local_histogram[THREADS][RANGE];

zero_histogram(local_histogram);

#pragma omp parallel num_threads(THREADS)
{
    int i;
    #pragma omp for
    for (i = 0; i < NUM_VALUES; i++){
        int value = randoms[i];
        local_histogram[omp_get_thread_num()][value]++;
    }
    #pragma omp barrier
    #pragma omp master
    for (t = 0; t < THREADS; t++){
        for (r = 0; r < RANGE; r++){
            histogram[r] += local_histogram[t][r];
        }
    }
}
```

Shared

But what about *i*?

Private

Scoping

```
int t, r;
int local_histogram[THREADS][RANGE];

zero_histogram(local_histogram);

#pragma omp parallel num_threads(THREADS)
{
    int i;
    #pragma omp for
    for (i = 0; i < NUM_VALUES; i++){
        int value = randoms[i];
        local_histogram[omp_get_thread_num()][value]++;
    }
    #pragma omp barrier
    #pragma omp master
    for (t = 0; t < THREADS; t++){
        for (r = 0; r < RANGE; r++){
            histogram[r] += local_histogram[t][r];
        }
    }
}
```

Shared

i is private as it is
the counter of the
parallel for loop

Private

Explicit scoping

❑ Why is explicit scoping required?

❑ It is possible to use implicit scoping as in previous example

❑ Although it is good practice to use shared for any shared variables

❑ The clause default(shared or none) is helpful in ensuring you have defined variables scope correctly

❑ By changing the default scope from shared to none it enforces explicit scoping of variables and will give errors if scoping is not defined

❑ `const` variables can not be explicitly scoped (always shared) - [more](#)

❑ Not enforced in windows but this is against the spec

```
int a, b = 0;
#pragma omp parallel default(none) shared(b)
{
    b += a;
}
```

error C3052: 'a' : variable doesn't appear in a data-sharing clause under a default(none) clause

Explicit scoping

- ❑ Why is explicit scoping required?
 - ❑ Older C programming (C89) style has variable declarations before definitions and statements (including loops)
 - ❑ Requires declarations to be made explicitly private for the parallel structured block
 - ❑ E.g. Consider our atomic histogram example

```
void calculate_histogram()  
{  
    int i;  
    int value;  
    #pragma omp parallel for private(value)  
    for (i = 0; i < NUM_VALUES; i++){  
        value = randoms[i];  
    #pragma omp atomic  
        histogram[value]++;  
    }  
}
```

Advanced private scoping

- ❑ If you want to pass the value of a variable outside of a parallel structured block then you must use the `firstprivate` clause
 - ❑ Private variables will be initialised with the value of the master thread before the parallel directive
- ❑ If you want to pass a private value to a variable outside of the parallel for loop you can use the `lastprivate` clause
 - ❑ This will assign the value of the last iteration of the loop

```
int i = 10;
#pragma omp parallel private(i)
{
    printf("Thread %d: i = %d\n", omp_get_thread_num(), i);
}
```

```
Thread 0: i = 0
Thread 2: i = 0
Thread 1: i = 0
Thread 3: i = 0
```

```
int i = 10;
#pragma omp parallel firstprivate(i)
{
    printf("Thread %d: i = %d\n", omp_get_thread_num(), i);
}
```

```
Thread 0: i = 10
Thread 2: i = 10
Thread 1: i = 10
Thread 3: i = 10
```

- ❑ Multicore systems and OpenMP
- ❑ Parallelising Loops
- ❑ Critical Sections and Synchronisation
- ❑ Scoping
- ❑ Data vs Task Parallelism



Data vs Task Parallelism

□ Is an OpenMP parallel for clause data or task parallel?

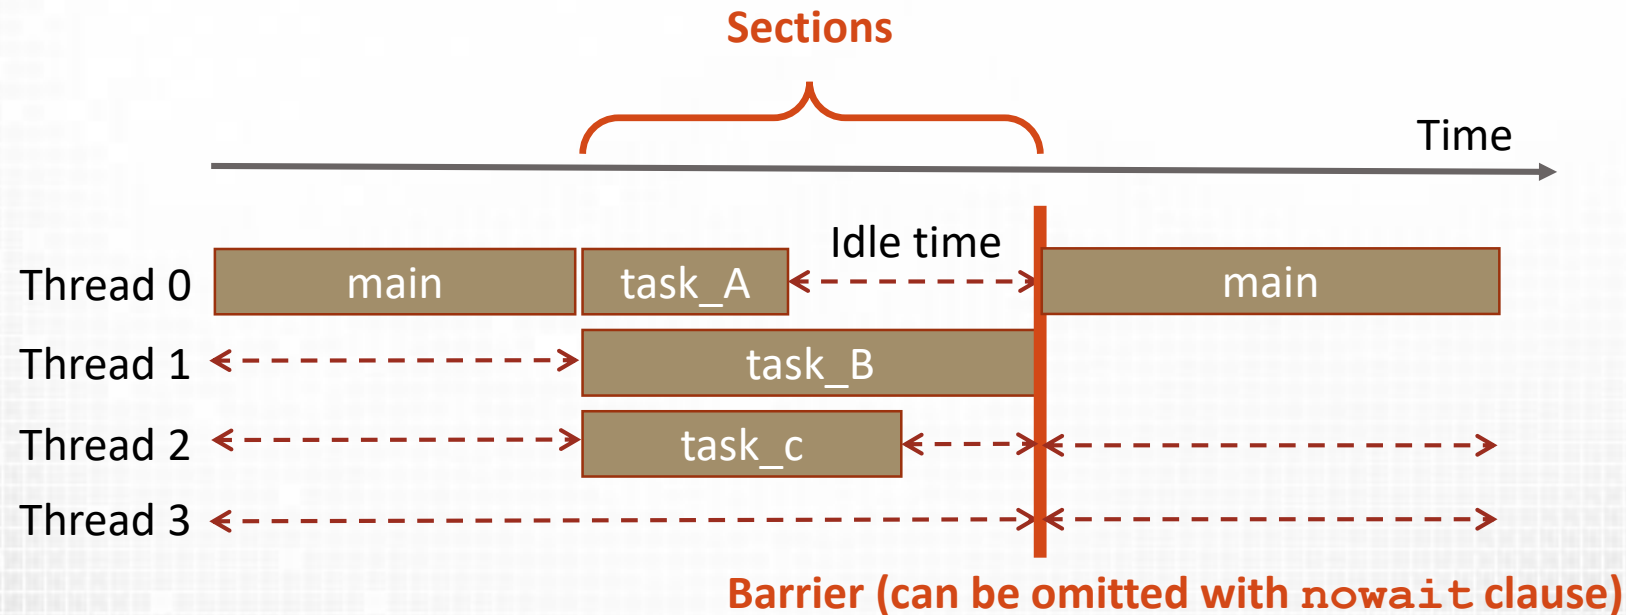
Data vs Task Parallelism

- ❑ Parallelism over loops is **data parallelism**. i.e.
 - ❑ The task is the same (the loop)
 - ❑ Parallelism is over the data elements the loop refers to
- ❑ What about task parallelism?
 - ❑ Task Parallelism: Divide a set of tasks between threads
 - ❑ This is supported by sections
 - ❑ Further task parallelism is supported by OpenMP tasks
 - ❑ This is OpenMP 3.0 spec and not supported in Visual Studio 2017
 - ❑ Very similar to sections

Sections

- ❑ `#pragma omp sections [clauses]`
 - ❑ Defines a code region where individual sections can be assigned to individual threads
 - ❑ Each section is executed exactly once by one thread
 - ❑ Unused threads wait for **implicit barrier**

```
#pragma omp parallel
#pragma omp sections
{
    #pragma omp section
        task_A();
    #pragma omp section
        task_B();
    #pragma omp section
        task_C();
}
```

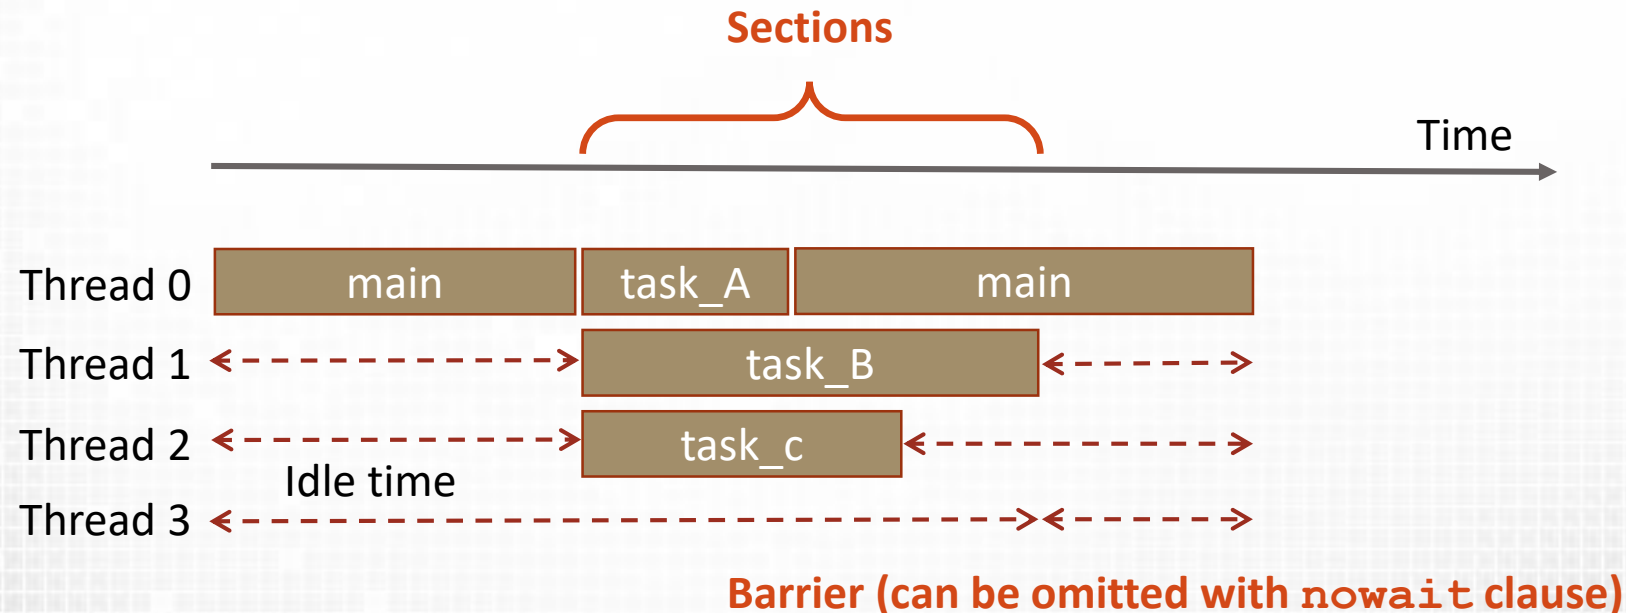


Sections

❑ If `nowait` clause is used then sections omit the barrier

❑ will immediately enter other parallel sections

```
#pragma omp parallel
{
    #pragma omp sections nowait
    {
        #pragma omp section
        task_A();
        #pragma omp section
        task_B();
        #pragma omp section
        task_C();
    }
}
```



Summary

- ❑ OpenMP lets us add explicit parallelism to serial code
 - ❑ We can parallelise loops or tasks
- ❑ OpenMP uses directives to modify the code
 - ❑ This enables portability (serial and parallel code is the same)
- ❑ OpenMP exposes both data and task parallelism using a fork and join model
- ❑ Care must be taken on parallel blocks which require access to shared variables
 - ❑ There is a distinction between private and shared variables within a parallel blocks scope