

Parallel Computing with GPUs: OpenMP 2

Dr Mozhgan Kabiri Chimeh

<http://mkchimeh.staff.shef.ac.uk/teaching/COM4521>



The
University
Of
Sheffield.



GPU
RESEARCH
CENTER

- ❑ OpenMP Timing
- ❑ Parallel Reduction
- ❑ Scheduling
- ❑ Nesting
- ❑ Summary

The problem with clock()

❑ clock() function behaviour

- ❑ In windows: represents a measure of real time (wall clock time)
- ❑ Linux: represents a cumulative measure of time spent executing instructions
 - ❑ Cumulative over core = not good for measuring parallel performance

❑ Open MP timing

- ❑ omp_get_wtime() – cross platform wall clock timing

```
double begin, end, seconds;  
begin = omp_get_wtime();  
  
some_function();  
  
end = omp_get_wtime();  
seconds = (end - begin);  
  
printf("Sum Time was %.2f seconds\n", seconds);
```

- ❑ OpenMP Timing
- ❑ Parallel Reduction
- ❑ Scheduling
- ❑ Nesting
- ❑ Summary

Parallel Reduction

- ❑ A Reduction is the combination of local copies of a variable into a single copy
 - ❑ Consider a case where we want to sum the values of a function operating on a vector of values;

```
void main() {  
    int i;  
    float vector[N];  
    float sum;  
  
    init_vector_values(vector);  
    sum = 0;  
  
    for (i = 0; i < N; i++) {  
        float v = some_func(vector[i]);  
        sum += v;  
    }  
    printf("Sum of values is %f\n", sum);  
}
```

Candidate for
parallel
reduction...

NBody calculation with OpenMP

```
void main(){
    int i;
    float vector[N];
    float sum;

    init_vector_values(vector);
    sum = 0;

    #pragma omp parallel for reduction(+: sum);
    for (i = 0; i < N; i++){
        float v = some_func(vector[i]);
        sum += v;
    }
    printf("Sum of values is %f\n", sum);
}
```

Without reduction
we would need a
critical section to
update the shared
variable!

OpenMP Reduction

- ❑ Reduction is supported with the reduction clause which requires a reduction variable
 - ❑ E.g. `#pragma omp parallel reduction(+: sum_variable) {...}`
 - ❑ Reduction variable is implicitly private to other threads
- ❑ OpenMP implements this by;
 - ❑ Creating a local (private) copy of the (shared) reduction variable
 - ❑ Combining local copies of the variable at the end of the structured block
 - ❑ Saving the reduced value to the shared variable in the master thread.
- ❑ Reduction operators are `+`, `-`, `*`, `&`, `|`, `&&` and `||`
 - ❑ `&`: bitwise and
 - ❑ `|`: bitwise or
 - ❑ `&&`: logical and
 - ❑ `||`: logical or

- ❑ OpenMP Timing
- ❑ Parallel Reduction
- ❑ Scheduling
- ❑ Nesting
- ❑ Summary



Scheduling

- ❑ OpenMP by default uses static scheduling
 - ❑ Static: schedule is determined at compile time
 - ❑ E.g. `#pragma omp parallel for schedule(static)`
- ❑ In general: `schedule(type [, chunk size])`
 - ❑ `type=static`: iterations assigned to threads before execution (preferably at compile time)
 - ❑ *`type=dynamic`: iterations are assigned to threads as they become available*
 - ❑ *`type=guided`: iterations are assigned to threads as they become available (with reducing chunk size)*
 - ❑ `type=auto`: compiler and runtime determine the schedule
 - ❑ `type=runtime`: schedule is determined at runtime by an environment variable, illegal to specify chunk size

What would be a use case where static scheduling is a bad choice?

Static scheduling chunk size

❑ chunk size

- ❑ Refers to the amount of work assigned to each thread
- ❑ By default chunk size is to divide the work by the number of threads
 - ❑ Low overhead (no going back for more work)
 - ❑ Not good for uneven workloads
 - ❑ E.g. consider our last lectures Taylor series example (updated to use reduction)

```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for reduction(-: result)  
for (n = 0; n < EXPANSION_STEPS; n++){  
    double r = pow(-1, n - 1) * pow(x, 2 * n - 1) / fac(2 * n);  
    result -= r;  
}  
  
printf("Approximation is %f, value is %f\n", result, cos(x));
```

Uneven workload


Scheduling Workload

```
long long int factorial(int n)
{
    if (n == 0)
        return 1;
    else
        return(n * factorial(n - 1));
}
```

❑ Uneven workload amongst threads

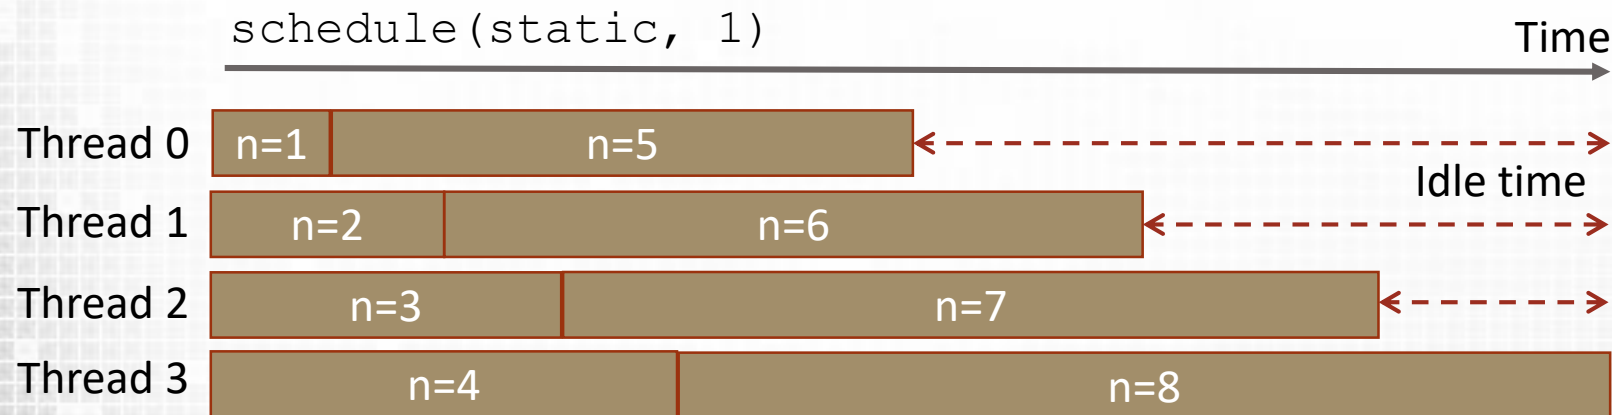

❑ Increase in n leads to increased computation

❑ E.g. `EXPANSION_STEPS=8, num_threads(4), schedule(static)`



Cyclic Scheduling

- ❑ It would be better to partition the workload more evenly
 - ❑ E.g. Cyclic scheduling via chunk size



Cyclic Scheduling

```
#pragma omp for num_threads(4)
for (i = 0; i < 16; i++)
```

schedule(static, 1)

Thread 0	0	4	8	12
Thread 1	1	5	9	13
Thread 2	2	6	10	14
Thread 3	3	7	11	15

schedule(static, 2)

Thread 0	0	1	8	9
Thread 1	2	3	10	11
Thread 2	4	5	12	13
Thread 3	6	7	14	15

schedule(static, 4)

Thread 0	0	1	2	3
Thread 1	4	5	6	7
Thread 2	8	9	10	11
Thread 3	12	13	14	15

Default case

❑ Default chunk size is $n/\text{threads}$

❑ where n is the number of iterations

Dynamic and Guided Scheduling

☐ Dynamic (med overhead)

- ☐ **Iterations are broken down by chunk size**

- ☐ Threads request chunks of work from a runtime queue when they are free

- ☐ Default chunk size is 1

☐ Guided (high overhead)

- ☐ **Chunks of the workload grow exponentially smaller**

- ☐ Threads request chunks of work from a runtime queue when they are free

- ☐ Chunk size is the size which the workloads decrease to

 - ☐ with the exception of last chunk which may have remainder

☐ Both

- ☐ Requesting work dynamically creates overhead

 - ☐ Not well suited if iterations are balanced

- ☐ Overhead vs. imbalance: How do I decide which is best?

 - ☐ **Benchmark all to find the best solution**

- ❑ OpenMP Timing
- ❑ Parallel Reduction
- ❑ Scheduling
- ❑ Nesting
- ❑ Summary

Nesting

- ❑ Consider the following example...
- ❑ How should we parallelise this example?

```
for (i = 0; i < OUTER_LOOPS; i++) {  
    for (j = 0; j < INNER_LOOPS; j++) {  
        printf("Hello World (Thread %d)\n", omp_get_thread_num());  
    }  
}
```

Nesting

- ❑ Consider the following example...
- ❑ How should we parallelise this example?

```
#pragma omp parallel for
for (i = 0; i < OUTER_LOOPS; i++){
    for (j = 0; j < INNER_LOOPS; j++){
        printf("Hello World (Thread %d)\n", omp_get_thread_num());
    }
}
```

- ❑ What if $\text{OUTER_LOOPS} \ll \text{number of threads}$
- ❑ E.g. $\text{OUTER_LOOPS} = 2$

Nesting

❑ We can use parallel nesting

- ❑ Nesting is turned off by default so we must use `omp_set_nested()`
- ❑ When inner loop is met each outer thread creates a new team of threads
- ❑ Allows us to expose higher levels of parallelism
 - ❑ *Only useful when outer loop does not expose enough*

```
omp_set_nested(1);


#define OUTER_LOOPS 2
#define INNER_LOOPS 4

#pragma omp parallel for
for (i = 0; i < OUTER_LOOPS; i++){
    int outer_thread = omp_get_thread_num();
    #pragma omp parallel for
    for (j = 0; j < INNER_LOOPS; j++){
        int inner_thread = omp_get_thread_num();
        printf("Hello World (i T=%d j T=%d)\n", outer_thread, inner_thread);
    }
}
```

```
Hello World (i T=0 j T=0)
Hello World (i T=0 j T=1)
Hello World (i T=0 j T=3)
Hello World (i T=1 j T=2)
Hello World (i T=1 j T=1)
Hello World (i T=1 j T=0)
Hello World (i T=0 j T=2)
Hello World (i T=1 j T=3)
```


Nesting Fork and Join

- ❑ Every parallel directive creates a fork (new team)
 - ❑ In this case each `omp parallel` is used to fork a new parallel region



Collapse

- ❑ Only available in OpenMP 3.0 and later (**not VS2017**)
- ❑ Can automatically collapse multiple loops
- ❑ Loops must not have statements or expressions between them

```
#pragma omp parallel for collapse(2)
for (i = 0; i < OUTER_LOOPS; i++){
    for (j = 0; j < INNER_LOOPS; j++){
        int thread = omp_get_thread_num();
        printf("Hello World (T=%d)\n", thread);
    }
}
```

Work around...

```
#pragma omp parallel for
for (i = 0; i < OUTER_LOOPS* INNER_LOOPS; i++){
    int thread = omp_get_thread_num();
    printf("Hello World (T=%d)\n", thread);
}
```

- ❑ OpenMP Timing
- ❑ Parallel Reduction
- ❑ Scheduling
- ❑ Nesting
- ❑ Summary

Clauses usage summary

Clause	Directive: #pragma omp ...					
	parallel	for	sections	single	parallel for	parallel sections
if						
private						
shared						
default						
firstprivate						
lastprivate						
reduction						
schedule						
nowait						

Performance

- ❑ Remember ideas for general C performance
 - ❑ Have good data locality (good cache usage)
 - ❑ Combine loops where possible
- ❑ Additional performance criteria
 - ❑ Minimise the use of barriers
 - ❑ Use `nowait` but only if it is safe to do so!
 - ❑ Minimise critical sections
 - ❑ High overhead. Can you use reduction or atomics?

Summary

- ❑ Parallel reduction is very helpful in combining data
 - ❑ It will use the OS most efficient method to implement the combination
- ❑ Scheduling can be static or dynamic
 - ❑ Static is good for fixed work sizes
 - ❑ Dynamic is good for varying work sizes
 - ❑ Benchmarking is important to find the best approach
- ❑ Nested parallelism can improve performance for outer loops with poor parallelism
- ❑ To get good performance try to avoid critical sections and barriers

Further reading

❑ <https://software.intel.com/en-us/articles/32-openmp-traps-for-c-developers>