

# Parallel Computing with GPUs: CUDA Memory

Dr Mozhgan Kabiri Chimeh

<http://mkchimeh.staff.shef.ac.uk/teaching/COM4521>



The  
University  
Of  
Sheffield.



GPU  
RESEARCH  
CENTER

## ❑ Memory Hierarchy Overview

❑ Global Memory

❑ Constant Memory

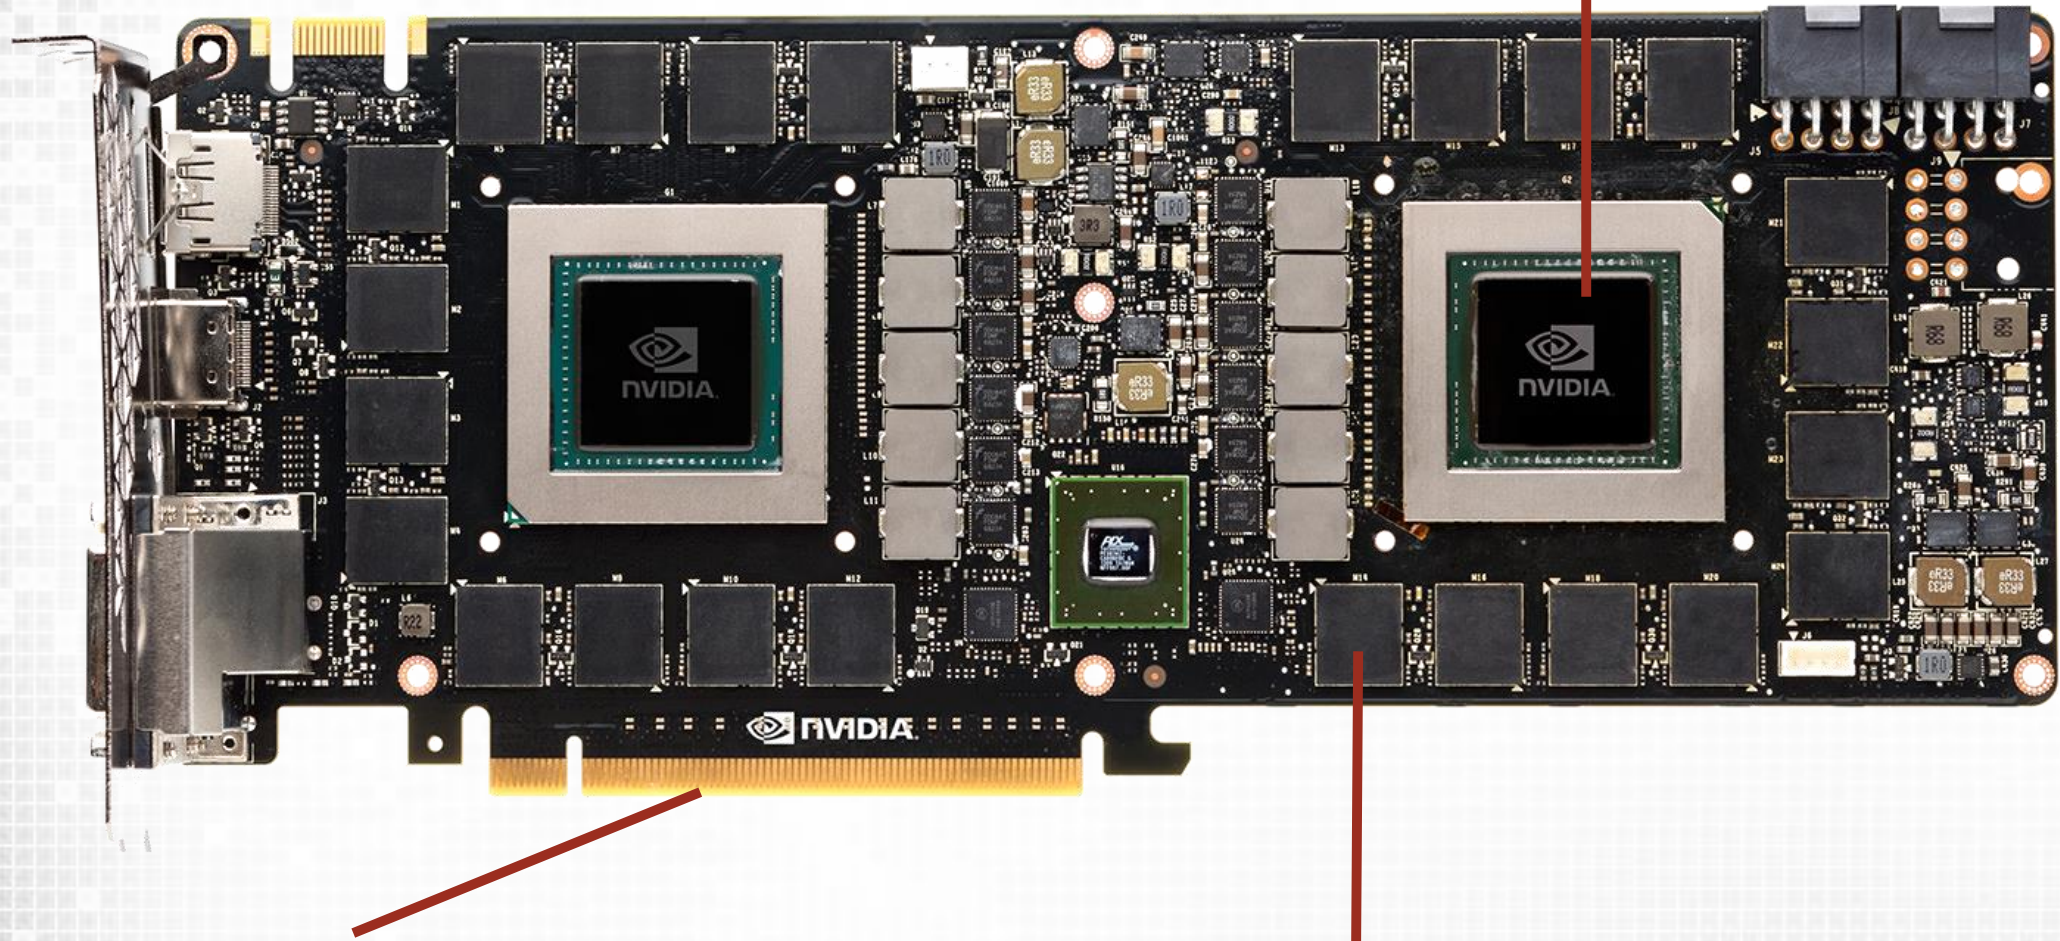
❑ Texture and Read-only Memory

❑ Roundup & Performance Timing



# GPU Memory (GTX Titan Z)

Shared Memory, cache and registers



Host Memory (via PCIe)

GPU DRAM Memory

# Simple Memory View

❑ Threads have access to;

❑ **Registers**

❑ Read/Write **per thread**

❑ **Local memory**

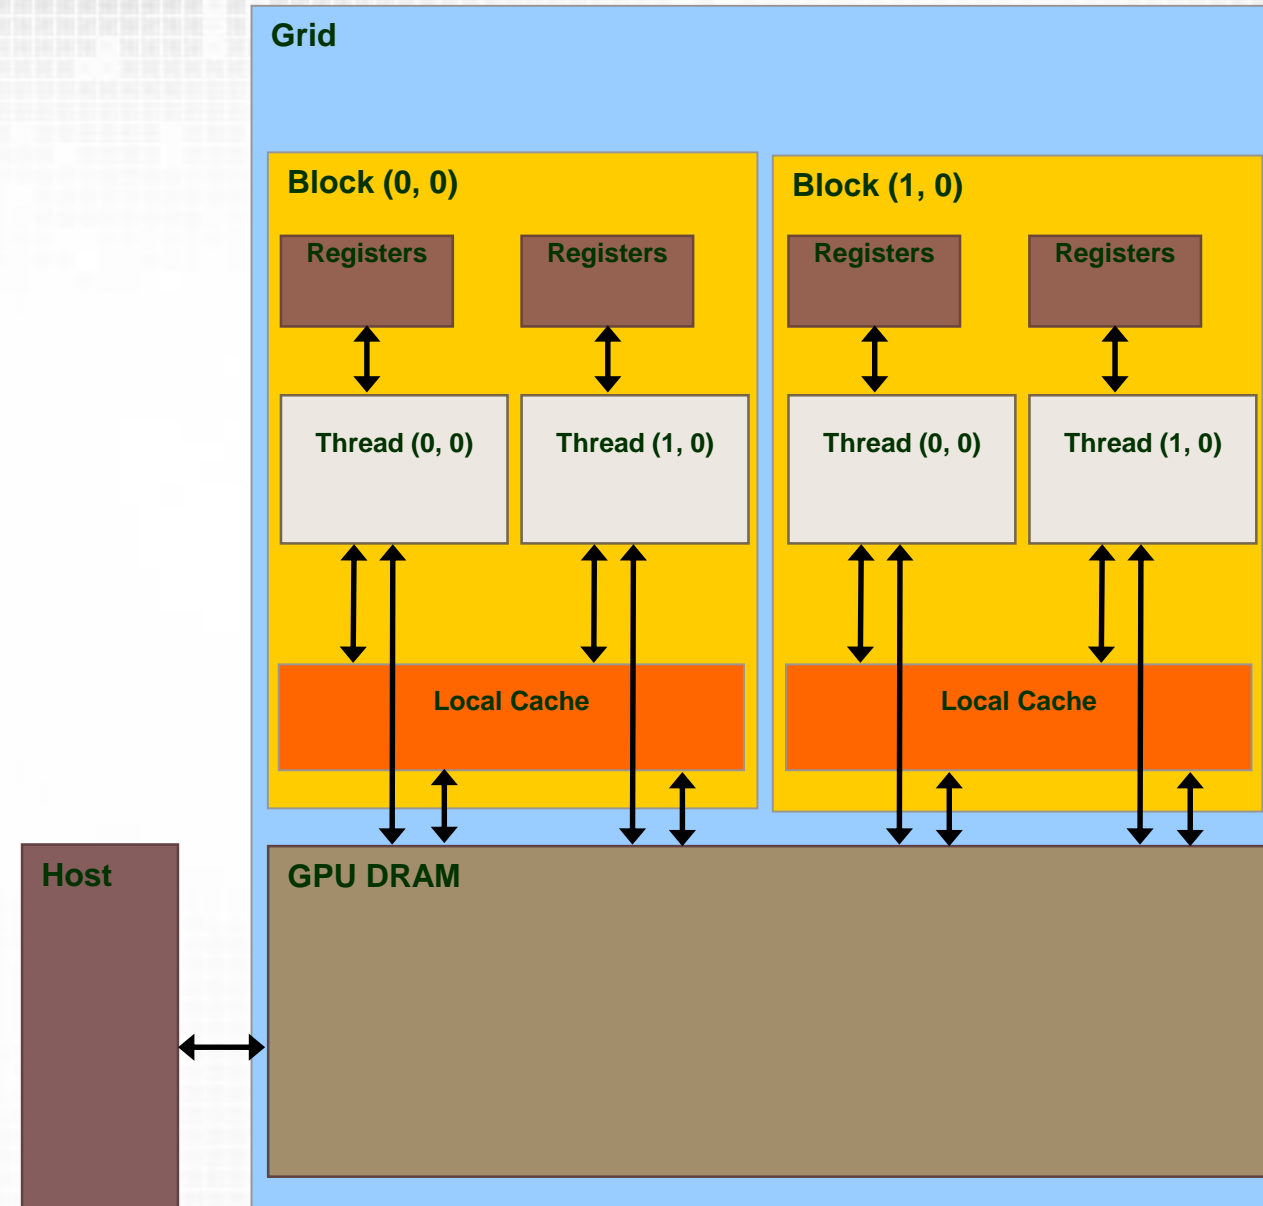
❑ Read/Write **per thread**

❑ **Local Cache**

❑ Read/Write **per block**

❑ **Main DRAM Memory**

❑ Read/Write **per grid**





# Local Memory

## ❑ Local memory (Thread-Local Global Memory)

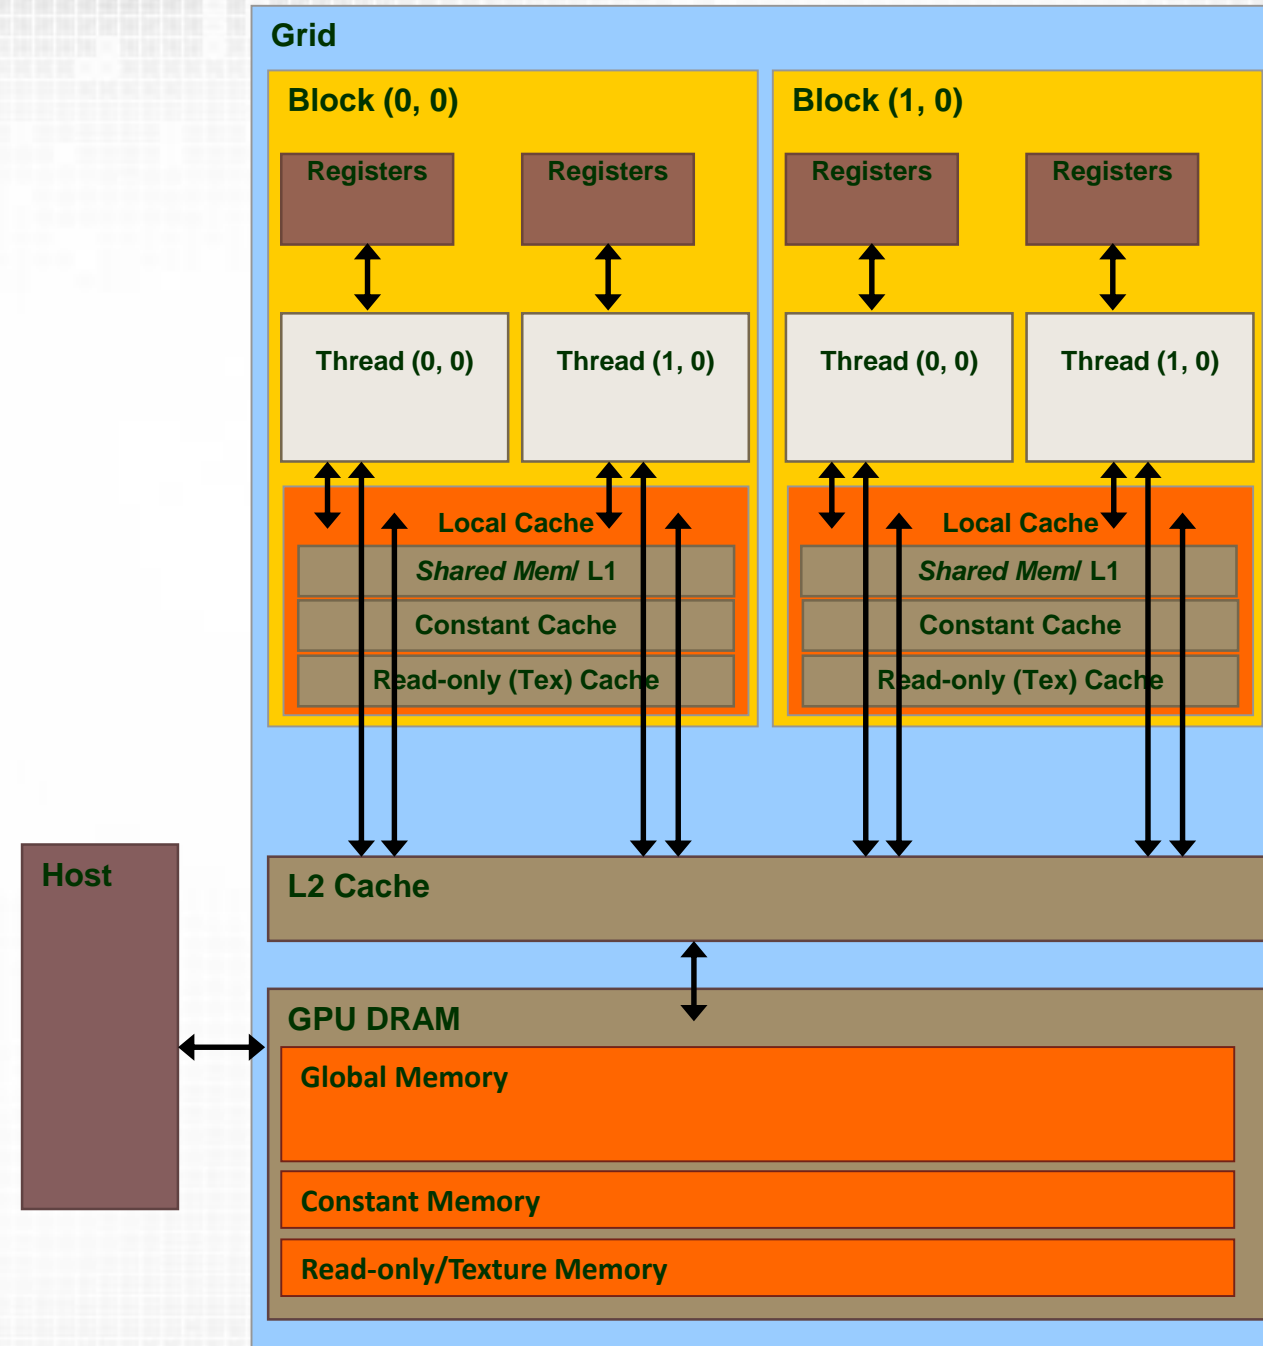
- ❑ Read/Write per thread
- ❑ Does not physically exist (reserved area in global memory)
- ❑ Cached locally
- ❑ Used for variables if you exceed the number of registers available
  - ❑ Very bad for perf!
- ❑ Arrays go in local memory if they are indexed with non constants

```
__global__ void localMemoryExample  
(int * input)  
{  
  
    int a;  
    int b;  
    int index;  
  
    int myArray1[4];  
    int myArray2[4];  
    int myArray3[100];  
  
    index = input[threadIdx.x];  
    a = myArray1[0];  
    b = myArray2[index];  
  
}
```

non constant index

# Kepler Memory View

- ❑ Each Thread has access to
  - ❑ Registers
  - ❑ Local memory
  - ❑ Main DRAM Memory via cache
    - ❑ Global Memory
      - ❑ Via **L2 cache** and configurable per block **Shared Memory cache**
    - ❑ Constant Memory
      - ❑ Via **L2 cache** and per block **Constant cache**
    - ❑ Read-only/Texture Memory
      - ❑ Via **L2 cache** and per block **Read-only cache**



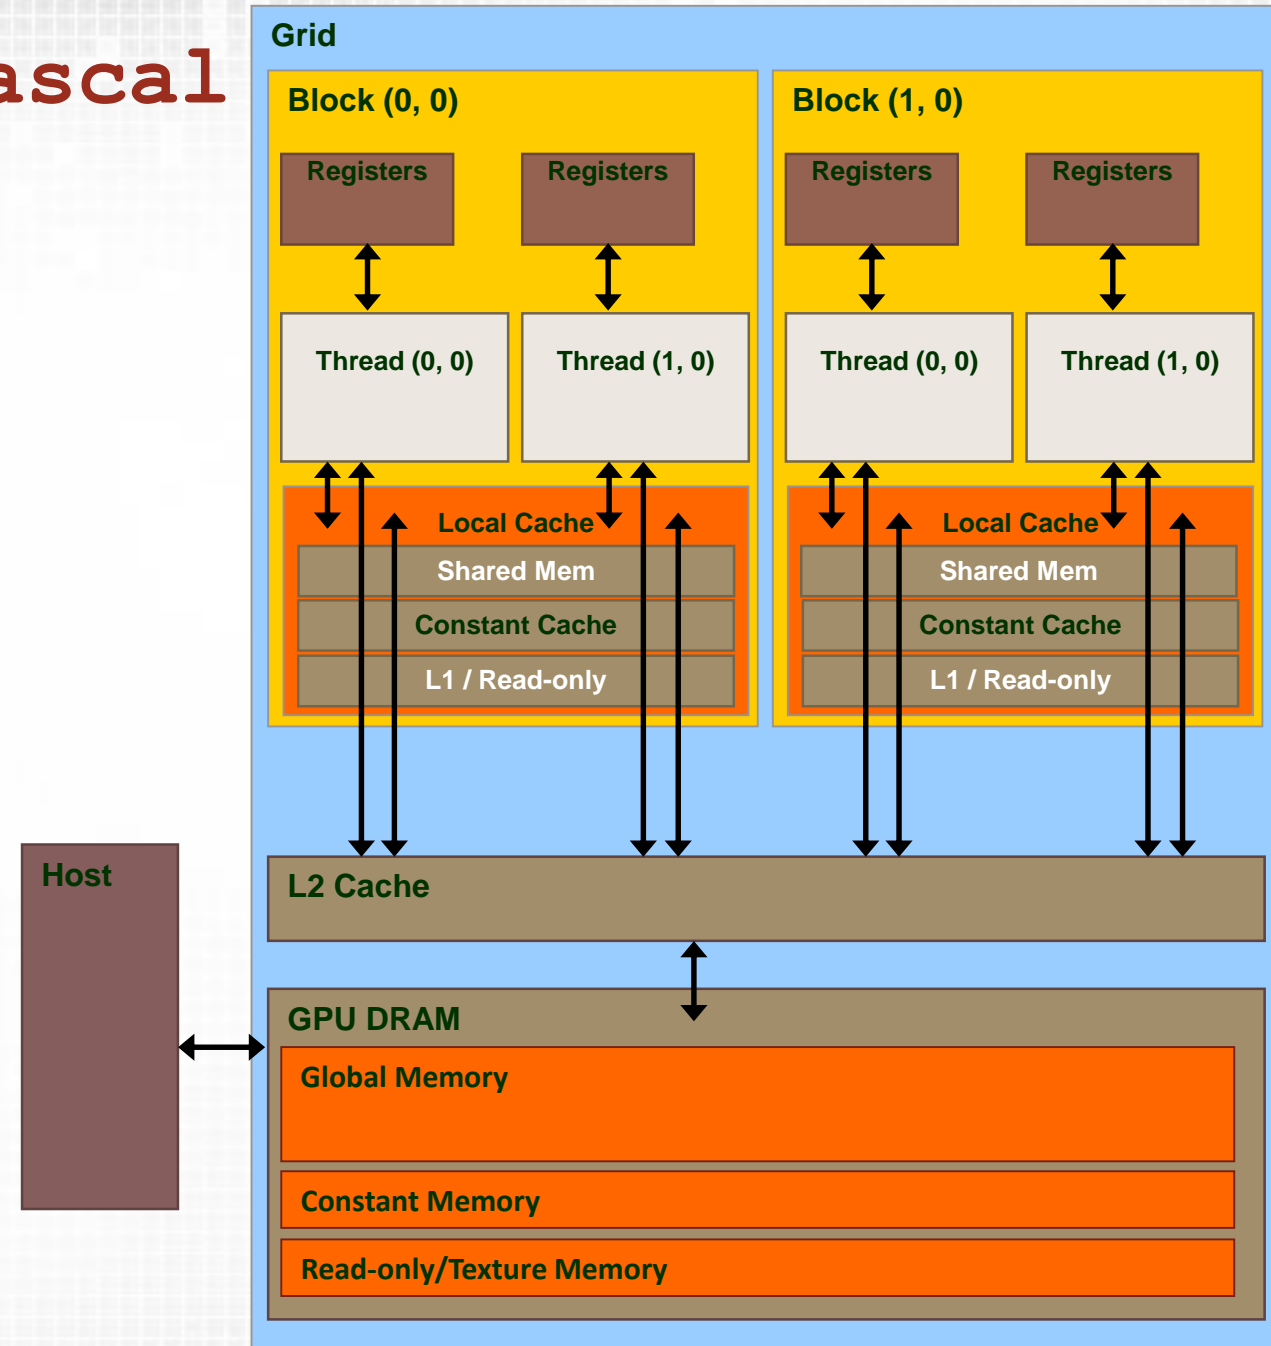
# Memory Latencies

- ❑ What is the cost of accessing each area of memory?
  - ❑ On chip caches are MUCH lower latency

	Cost (cycles)
Register	1
Global	200-800
Shared memory	~1
L1	1
Constant	~1 (if cached)
Read-only (tex)	1 if cached (same as global if not)

# Changes in Maxwell/Pascal

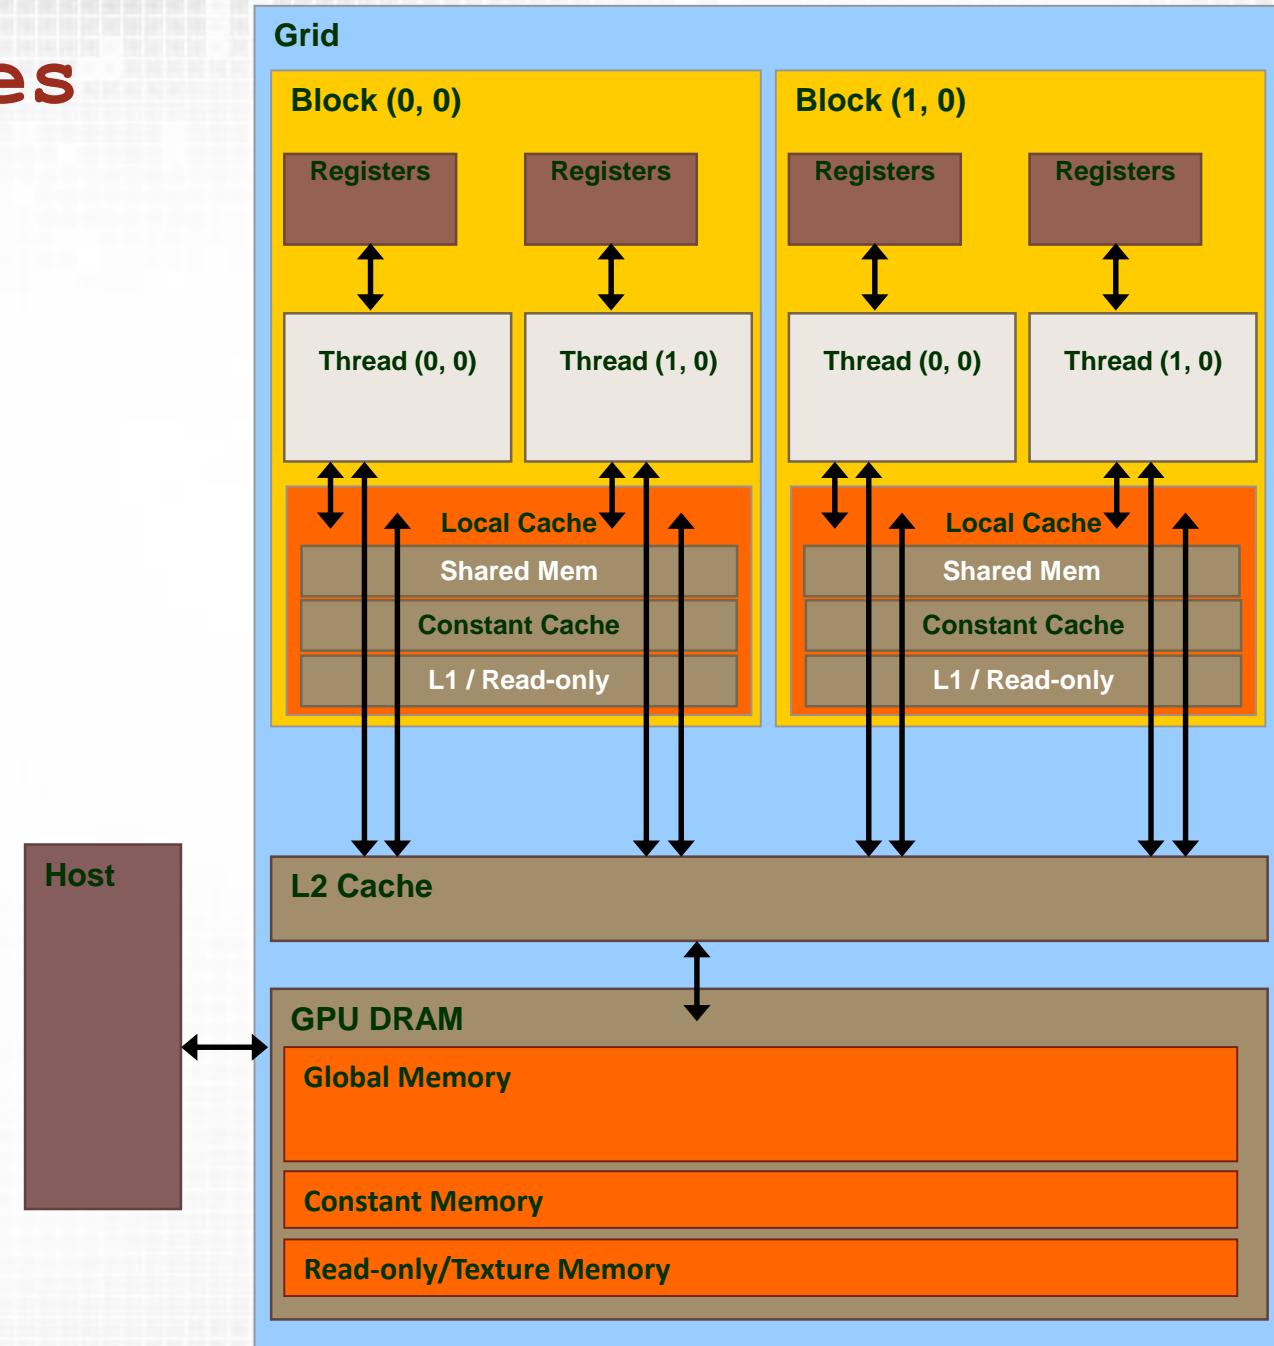
- ❑ Shared memory has dedicated Cache
  - ❑ No longer shared with L1 as in Fermi and Kepler
- ❑ Read-only (texture) cache unified with L1





# Cache and Memory Sizes

	Kepler	Maxwell
<b>Registers</b>	<i>64k 32 bit registers per SM</i>	<i>64k 32 bit registers per SM</i>
<b>Max Registers / thread</b>	63	255
<b>Shared Memory</b>	<i>16KB / 48KB Configurable SM and L1</i>	<i>64KB Dedicated</i>
<b>Constant Memory</b>	<i>64KB DRAM 8KB Cache per SM</i>	<i>64KB DRAM 8KB Cache per SM</i>
<b>Read Only Memory</b>	<i>48KB per SM</i>	<i>48KB per SM Shared with L1</i>
<b>Device Memory</b>	<i>Varying 12GB Max</i>	<i>Varying 12GB Max</i>



# Device Query

❑ What are the specifics of my GPU?

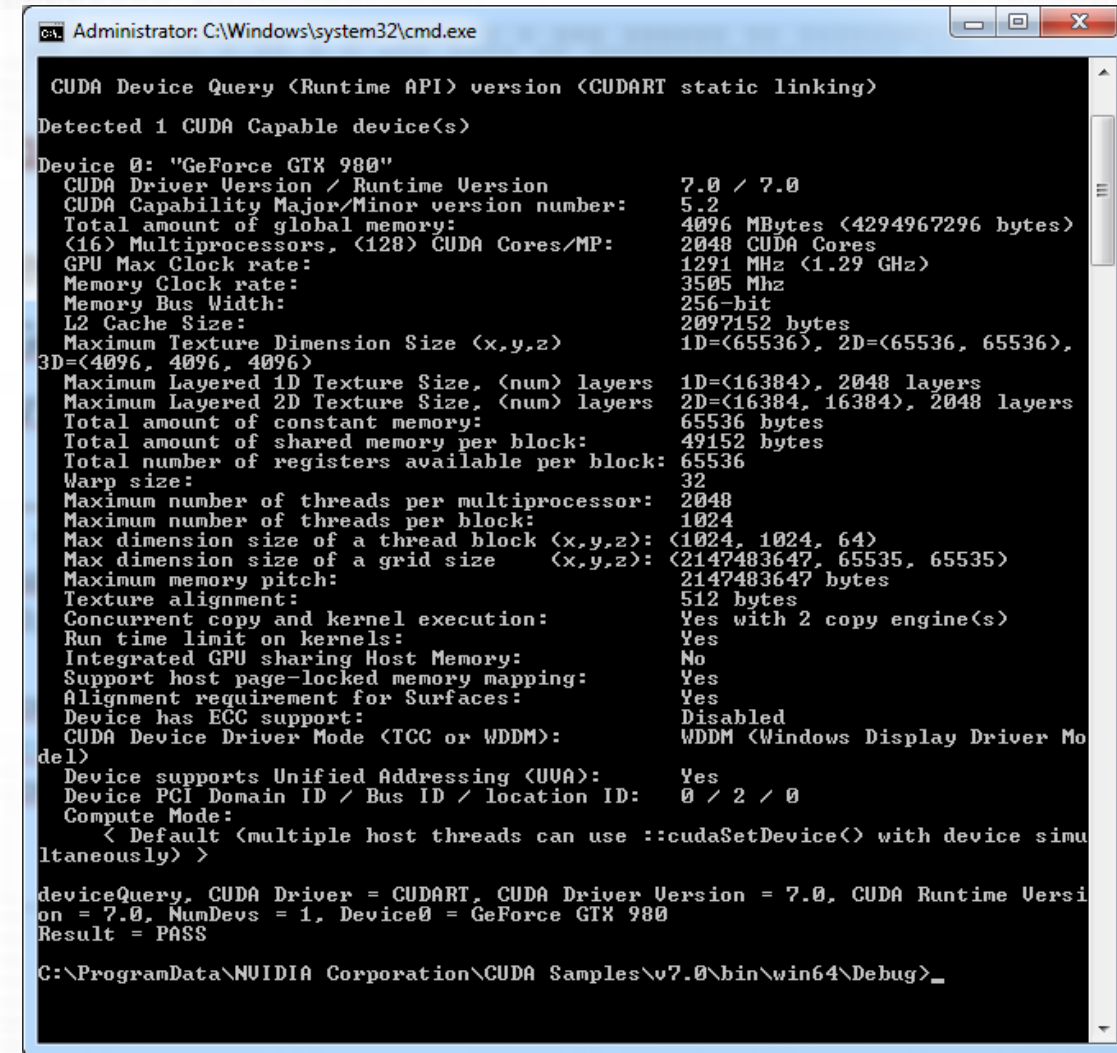
❑ Use `cudaGetDeviceProperties`

❑ E.g.

❑ `deviceProp.sharedMemPerBlock`

❑ CUDA SDK `deviceQuery` example

```
int deviceCount = 0;
cudaGetDeviceCount(&deviceCount);
for (int dev = 0; dev < deviceCount; ++dev)
{
    cudaSetDevice(dev);
    cudaDeviceProp deviceProp;
    cudaGetDeviceProperties(&deviceProp, dev);
    ...
}
```



```
Administrator: C:\Windows\system32\cmd.exe

CUDA Device Query (Runtime API) version (CUDA static linking)

Detected 1 CUDA Capable device(s)

Device 0: "GeForce GTX 980"
  CUDA Driver Version / Runtime Version      7.0 / 7.0
  CUDA Capability Major/Minor version number: 5.2
  Total amount of global memory:             4096 MBytes (4294967296 bytes)
  (16) Multiprocessors, (128) CUDA Cores/MP: 2048 CUDA Cores
  GPU Max Clock rate:                        1291 MHz (1.29 GHz)
  Memory Clock rate:                         3505 MHz
  Memory Bus Width:                          256-bit
  L2 Cache Size:                             2097152 bytes
  Maximum Texture Dimension Size (x,y,z)     1D=(65536), 2D=(65536, 65536),
  3D=(4096, 4096, 4096)
  Maximum Layered 1D Texture Size, (num) layers 1D=(16384), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(16384, 16384), 2048 layers
  Total amount of constant memory:            65536 bytes
  Total amount of shared memory per block:    49152 bytes
  Total number of registers available per block: 65536
  Warp size:                                 32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z):  (2147483647, 65535, 65535)
  Maximum memory pitch:                      2147483647 bytes
  Texture alignment:                         512 bytes
  Concurrent copy and kernel execution:       Yes with 2 copy engine(s)
  Run time limit on kernels:                  Yes
  Integrated GPU sharing Host Memory:          No
  Support host page-locked memory mapping:    Yes
  Alignment requirement for Surfaces:         Yes
  Device has ECC support:                     Disabled
  CUDA Device Driver Mode (TCC or WDDM):       WDDM (Windows Display Driver Mo
del)
  Device supports Unified Addressing (UVA):    Yes
  Device PCI Domain ID / Bus ID / location ID: 0 / 2 / 0
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simul
taneously) >
deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 7.0, CUDA Runtime Versi
on = 7.0, NumDevs = 1, Device0 = GeForce GTX 980
Result = PASS

C:\ProgramData\NVIDIA Corporation\CUDA Samples\v7.0\bin\win64\Debug>
```

- ❑ Memory Hierarchy Overview
- ❑ Global Memory
- ❑ Constant Memory
- ❑ Texture and Read-only Memory
- ❑ Roundup & Performance Timing

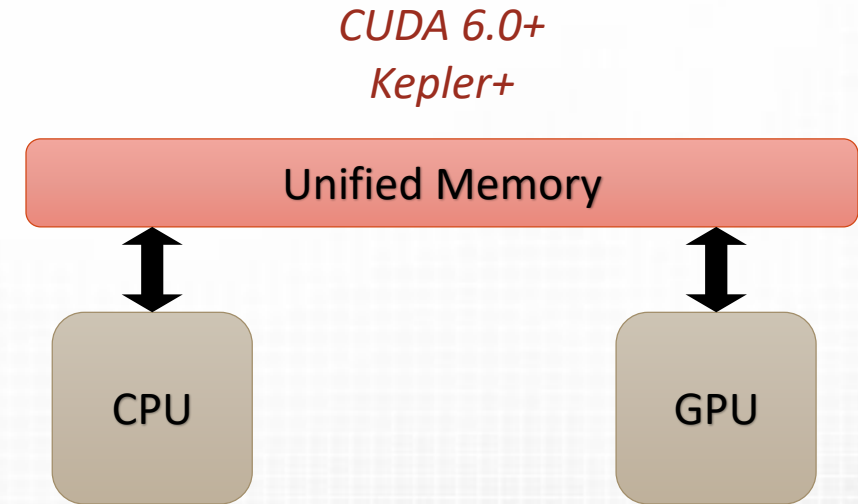



# Dynamic vs Static Global Memory

- ❑ In the previous lecture we learned about dynamically defined GPU memory
  - ❑ Using `cudaMalloc()`
- ❑ You can also statically define (and allocate) GPU global memory
  - ❑ Using `__device__` qualifier
  - ❑ Requires memory copies are performed using `cudaMemcpyToSymbol` or `cudaMemcpyFromSymbol`
  - ❑ See example from last weeks lecture
- ❑ This is the difference between the following in C
  - ❑ `int my_static_array[1024];`
  - ❑ `int *my_dynamic_array = (int*) malloc(1024*sizeof(int));`

# Unified Memory

- ❑ So far the developer view is that GPU and CPU have separate memory
  - ❑ Memory must be explicitly copied
  - ❑ Deep copies required for complex data structures
- ❑ Unified Memory changes that view



# Unified Memory Example

## C Code

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    data = (char *)malloc(N);  
  
    fread(data, 1, N, fp);  
  
    qsort(data, N, 1, compare);  
  
    use_data(data);  
  
    free(data);  
}
```

## CUDA (6.0+) Code

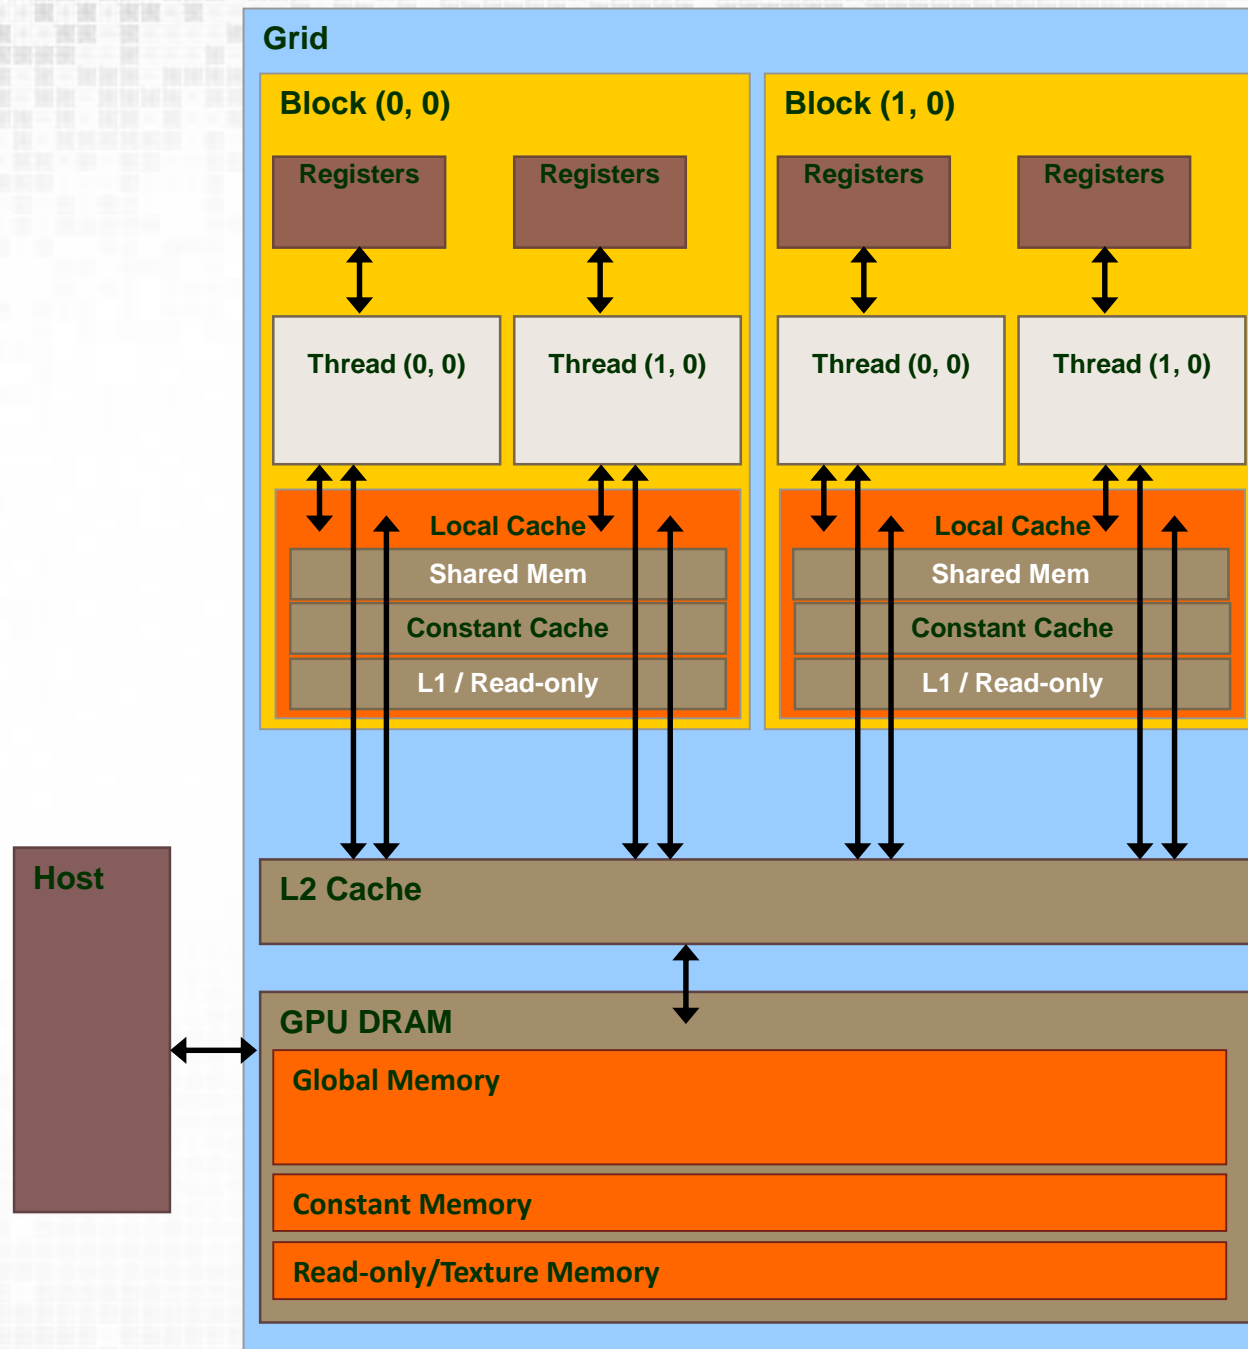
```
void sortfile(FILE *fp, int N) {  
    char *data;  
    cudaMallocManaged(&data, N);  
  
    fread(data, 1, N, fp);  
  
    gpu_qsort(data, N, 1, compare);  
    cudaDeviceSynchronize();  
  
    use_data(data);  
  
    cudaFree(data);  
}
```



# Implications of CUDA Unified Managed Memory

- ❑ Simpler porting of code
- ❑ Memory is only *virtually* unified
  - ❑ GPU still has discrete memory
  - ❑ It still has to be transferred via PCIe (or NVLINK)
- ❑ Easier management of data to and from the device
  - ❑ Explicit memory movement is not required
  - ❑ Similar to the way the OS handles virtual memory
- ❑ Issues
  - ❑ Requires look ahead and paging to ensure memory is in the correct place (and synchronised)
  - ❑ It is not as fast as hand tuned code which has finer explicit control over transfers
- ❑ *We will manage memory movement ourselves!*

- ❑ Memory Hierarchy Overview
- ❑ Global Memory
- ❑ **Constant Memory**
- ❑ Texture and Read-only Memory
- ❑ Roundup & Performance Timing





# Constant Memory

## ☐ Constant Memory

- ☐ Stored in the devices global memory
- ☐ Read through the per SM constant cache
- ☐ Set at runtime
- ☐ When using correctly only 1/16 of the traffic compared to global loads

## ☐ When to use it?

- ☐ When small amounts of data are **read only**
- ☐ When values are **broadcast** to threads in a half warp (of 16 threads)
- ☐ Very fast when cache hit
- ☐ Very slow when no cache hit

## ☐ How to use

- ☐ Must be **statically** (compile-time) defined as a symbol using `__constant__` qualifier
- ☐ Value(s) must be copied using **cudaMemcpyToSymbol**.



# Constant Memory Broadcast

❑.... When values are **broadcast** to threads in a half warp (groups of 16 threads)

```
__constant__ int my_const[16];

__global__ void vectorAdd() {
    int i = blockIdx.x;

    int value = my_const[i % 16];
}
```

```
__constant__ int my_const[16];

__global__ void vectorAdd() {
    int i = blockIdx.x * blockDim.x + threadIdx.x;

    int value = my_const[i % 16];
}
```

*Which is good use of constant memory?*

# Constant Memory Broadcast

❑.... When values are **broadcast** to threads in a half warp (groups of 16 threads)

```
__constant__ int my_const[16];

__global__ void constant_test() {
    int i = blockIdx.x;

    int value = my_const[i % 16];
}
```

```
__constant__ int my_const[16];

__global__ void constant_test() {
    int i = blockIdx.x * blockDim.x + threadIdx.x;

    int value = my_const[i % 16];
}
```

*Which is good use of constant memory?*

- ❑ Best possible use of constant memory
- ❑ Every thread in half warp reads the same
  - ❑ Index based on `blockIdx`
- ❑ No serialisation
  - ❑ 1 read request for every thread!
- ❑ Other threads in the block will also hit cache

- ❑ Worst possible use of constant memory
- ❑ Every thread in half warp reads different value
  - ❑ Index based on `threadIdx`
- ❑ Each access will be serialised
  - ❑ 16 different read requests!
- ❑ Other threads in the block will likely miss the cache



# Constant Memory

❑ Question: Should I convert `#define` to constants?

❑ E.g. `#define MY_CONST 1234`

❑ Answer: No

❑ Leave alone

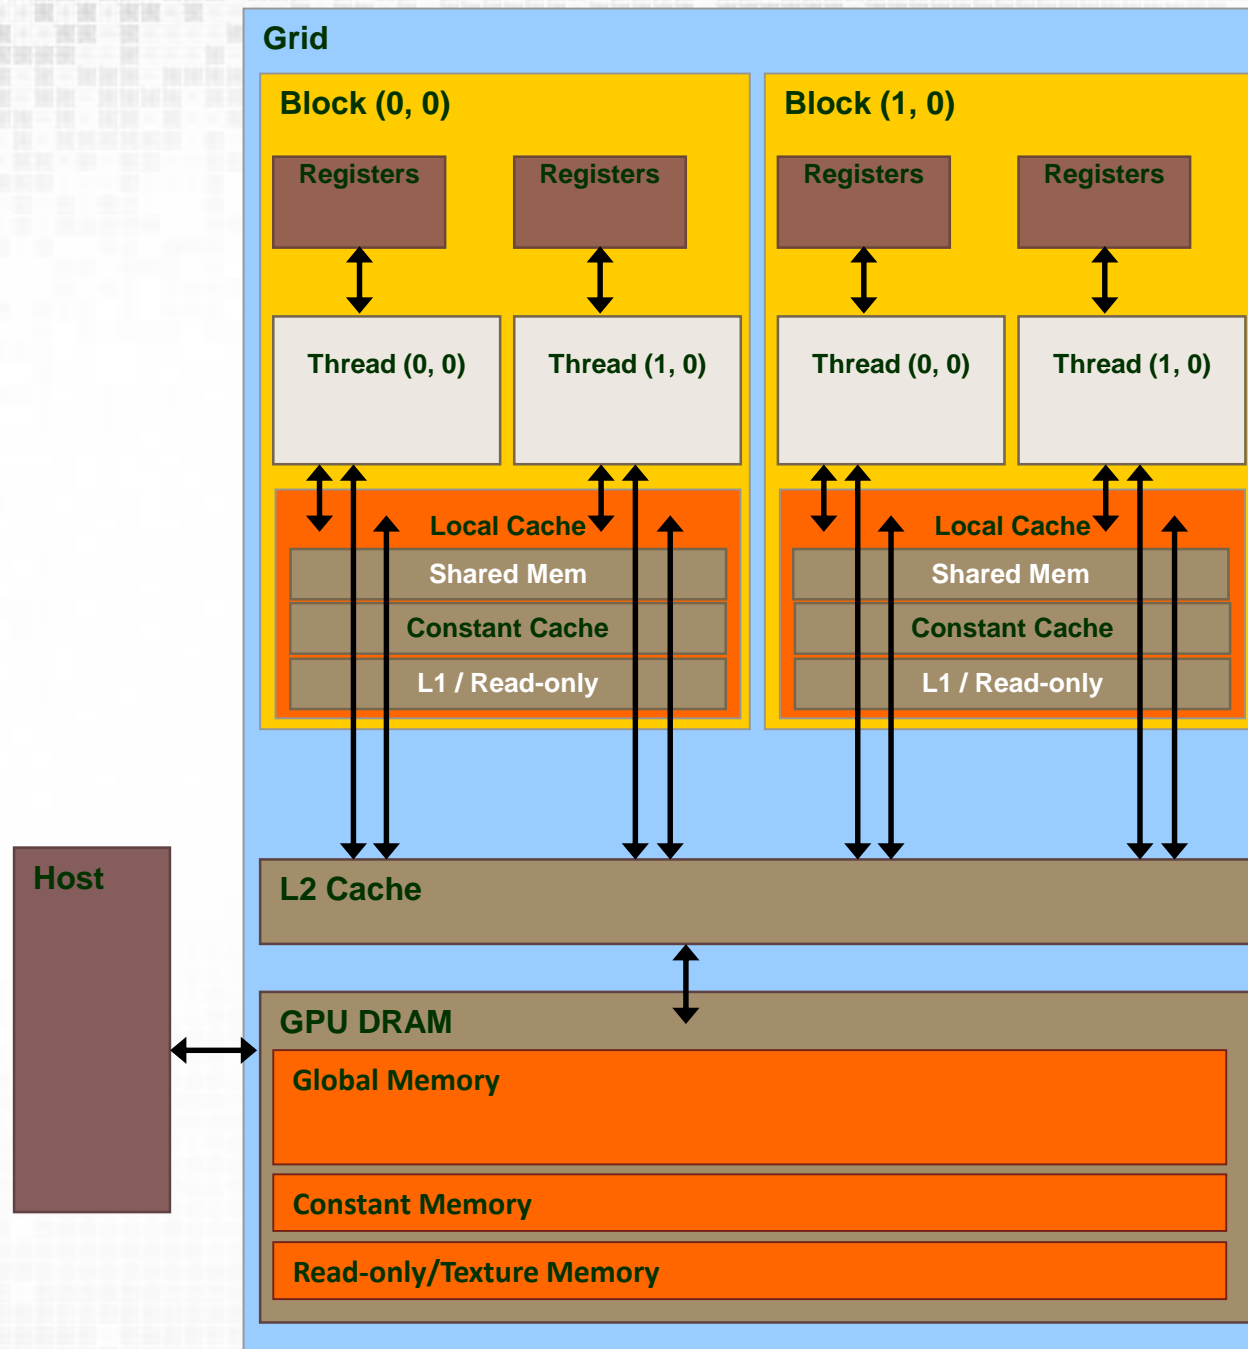
❑ `#defines` are embed in the code by pre-processors

❑ They don't take up registers as they are embed within the instruction space

❑ i.e. are replaced with literals by the pre-processor

❑ Only replace constants that may change at runtime (but not during the GPU programs)

- ❑ Memory Hierarchy Overview
- ❑ Global Memory
- ❑ Constant Memory
- ❑ Texture and Read-only Memory
- ❑ Roundup & Performance Timing



# Read-only and Texture Memory

- ❑ Separate in Kepler but unified thereafter
  - ❑ Same use case but used in different ways
- ❑ When to use read-only or texture
  - ❑ When data is read only
  - ❑ Good for bandwidth limited kernels
  - ❑ Regular memory accesses with good locality (think about the way textures are accessed)
- ❑ Two Methods for utilising Read-only/Texture Memory
  - ❑ Bind memory to texture (or use advanced bindless textures in CUDA 5.0+)
  - ❑ Hint the compiler to load via read-only cache



# Texture Memory Binding

□ Known as bound texture (or texture reference method)

```
#define N 1024
texture<float, 1, cudaReadModeElementType> tex;

__global__ void kernel() {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x = tex1Dfetch(tex, i);
}

int main() {
    float *d_buffer;
    cudaMalloc(&d_buffer, N*sizeof(float));
    cudaBindTexture(0, tex, d_buffer, N*sizeof(float));
    kernel << <grid, block >> >();
    cudaUnbindTexture(tex);
    cudaFree(d_buffer);
}
```

# Texture Memory Binding

❑ Known as bound texture (or texture reference method)

```
#define N 1024
texture<float> 1, cudaReadModeElementType> tex;

__global__ void kernel() {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x = tex1Dfetch(tex, i);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, N*sizeof(float));
    cudaBindTexture(0, tex, buffer, N*sizeof(float));
    kernel << <grid, block >> >();
    cudaUnbindTexture(tex);
    cudaFree(buffer);
}
```

Must be either;

❑ char, short, long,  
long long, float or  
double

Vector Equivalents are also  
permitted e.g.

❑ uchar4

# Texture Memory Binding

❑ Known as bound texture (or texture reference method)

```
#define N 1024
texture<float, 1, cudaReadModeElementType> tex;

__global__ void kernel() {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x = tex1Dfetch(tex, i);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, N*sizeof(float));
    cudaBindTexture(0, tex, buffer, N*sizeof(float));
    kernel << <grid, block >> >();
    cudaUnbindTexture(tex);
    cudaFree(buffer);
}
```

Dimensionality:

- ❑ cudaTextureType1D (1)
- ❑ cudaTextureType2D (2)
- ❑ cudaTextureType3D (3)
- ❑ cudaTextureType1DLayered (4)
- ❑ cudaTextureType2DLayered (5)
- ❑ cudaTextureTypeCubemap (6)
- ❑ cudaTextureTypeCubemapLayered (7)

# Texture Memory Binding

❑ Known as bound texture (or texture reference method)

```
#define N 1024
texture<float, 1, cudaReadModeElementType> tex;

__global__ void kernel() {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x = tex1Dfetch(tex, i);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, N*sizeof(float));
    cudaBindTexture(0, tex, buffer, N*sizeof(float));
    kernel << <grid, block >> >();
    cudaUnbindTexture(tex);
    cudaFree(buffer);
}
```

Value normalization:

- ❑ cudaReadModeElementType
- ❑ cudaReadModeNormalizedFloat
  - ❑ Normalises values across range



# Texture Memory Binding on 2D Arrays

```
#define N 1024
texture<float, 2, cudaReadModeElementType> tex;

__global__ void kernel() {
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;
    float v = tex2D (tex, x, y);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, W*H*sizeof(float));
    cudaChannelFormatDesc desc = cudaCreateChannelDesc<float>();
    cudaBindTexture2D(0, tex, buffer, desc, W,
                      H, W*sizeof(float));
    kernel << <grid, block >> >();
    cudaUnbindTexture(tex);
    cudaFree(buffer);
}
```

❑ Use tex2D rather than tex1Dfetch for CUDA arrays

❑ Note that last arg of **cudaBindTexture2D** is pitch

❑ Row size not != total size

# Read-only Memory

- ❑ No textures required
- ❑ Hint to the compiler that the data is read-only without pointer aliasing
  - ❑ Using the `const` and `__restrict__` qualifiers
  - ❑ Suggests the compiler should use `__ldg` but does not guarantee it
- ❑ Not the same as `__constant__`
  - ❑ Does not require broadcast reading

```
#define N 1024

__global__ void kernel(float const* __restrict__ buffer) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x = __ldg(buffer[i]);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, N*sizeof(float));
    kernel << <grid, block >> >(buffer);
    cudaFree(buffer);
}
```

- ❑ Memory Hierarchy Overview
- ❑ Global Memory
- ❑ Constant Memory
- ❑ Texture and Read-only Memory
- ❑ Roundup & Performance Timing

# CUDA qualifiers summary

## ❑ Where can a variable be accessed?

### ❑ Is declared inside the kernel?

❑ Then the host can not access it

❑ Lifespan ends after kernel execution

### ❑ Is declared outside the kernel

❑ Then the host can access it (via  
cudaMemcpyToSymbol)

## ❑ What about pointers?

❑ They can point to anything

❑ BUT are not typed on memory space

❑ Be careful not to confuse the compiler

Remember!

```
const int *p != int * const p
```

```
__device__ int my_global;  
__constant__ int my_constant;  
  
__global__ void my_kernel() {  
    int my_local;  
  
    int *ptr1 = &my_global;  
    int *ptr2 = &my_local;  
    const int *ptr3 = &my_constant;  
}
```

```
if (something)  
    ptr1 = &my_global;  
else  
    ptr1 = &my_local;
```



# Performance Measurements

- ❑ How can we benchmark our CUDA code?
- ❑ Kernel Calls are asynchronous
  - ❑ If we use a standard CPU timer it will measure only launch time not execution time
  - ❑ We could call `cudaDeviceSynchronise()` but this will stall the entire GPU pipeline
- ❑ Alternative: CUDA Events
  - ❑ Events are created with `cudaEventCreate()`
  - ❑ Timestamps can be set using `cudaEventRecord()`
  - ❑ `cudaEventElapsedTime()` sets the time in *ms* between the two events.

```
cudaEvent_t start, stop;
cudaEventCreate(&start);
cudaEventCreate(&stop);

cudaEventRecord(start);
my_kernel <<<(N /TPB), TPB >>>();
cudaEventRecord(stop);

cudaEventSynchronize(stop);
float milliseconds = 0;
cudaEventElapsedTime(&milliseconds,
                    start, stop);

cudaEventDestroy(start);
cudaEventDestroy(stop);
```

## Summary

- ❑ The CUDA Memory Hierarchy varies between hardware generations
- ❑ Utilisation of local caches can have a big impact on the expected performance (1 cycle vs. 100s)
- ❑ Global memory can be declared statically or dynamically
- ❑ Constant cache good for small read only data accessed in broadcast by *nearby* threads
- ❑ Read-Only cache is larger than constant cache but does not have broadcast performance of constant cache
- ❑ Kernel variables are not available outside of the kernel



# Mole Quiz – March 5th

- ☐ NO lectures for next 2 weeks
- ☐ Labs will be as usual
- ☐ 2 Weeks from NOW (17:00) will be a MOLE quiz.
  - ☐ Where? Alfred Denny PC Room (ADB-A04)
  - ☐ When? March 5th
  - ☐ How Long: 45 mins (25 Questions)
  - ☐ What? Everything up to the end of last weeks lectures...

☐ E.g.

```
int a[5] = {1, 2, 3, 4, 5};  
x = &a[3];
```

☐ What is x?

1. a pointer to an integer with value of 3
2. a pointer to an integer with value of 4
3. a pointer to an integer with a value of the address of the third element of a
4. an integer with a value of 4