

Parallel Computing with GPUs: Shared Memory

Dr Mozhgan Kabiri Chimeh

<http://mkchimeh.staff.shef.ac.uk/teaching/COM4521>

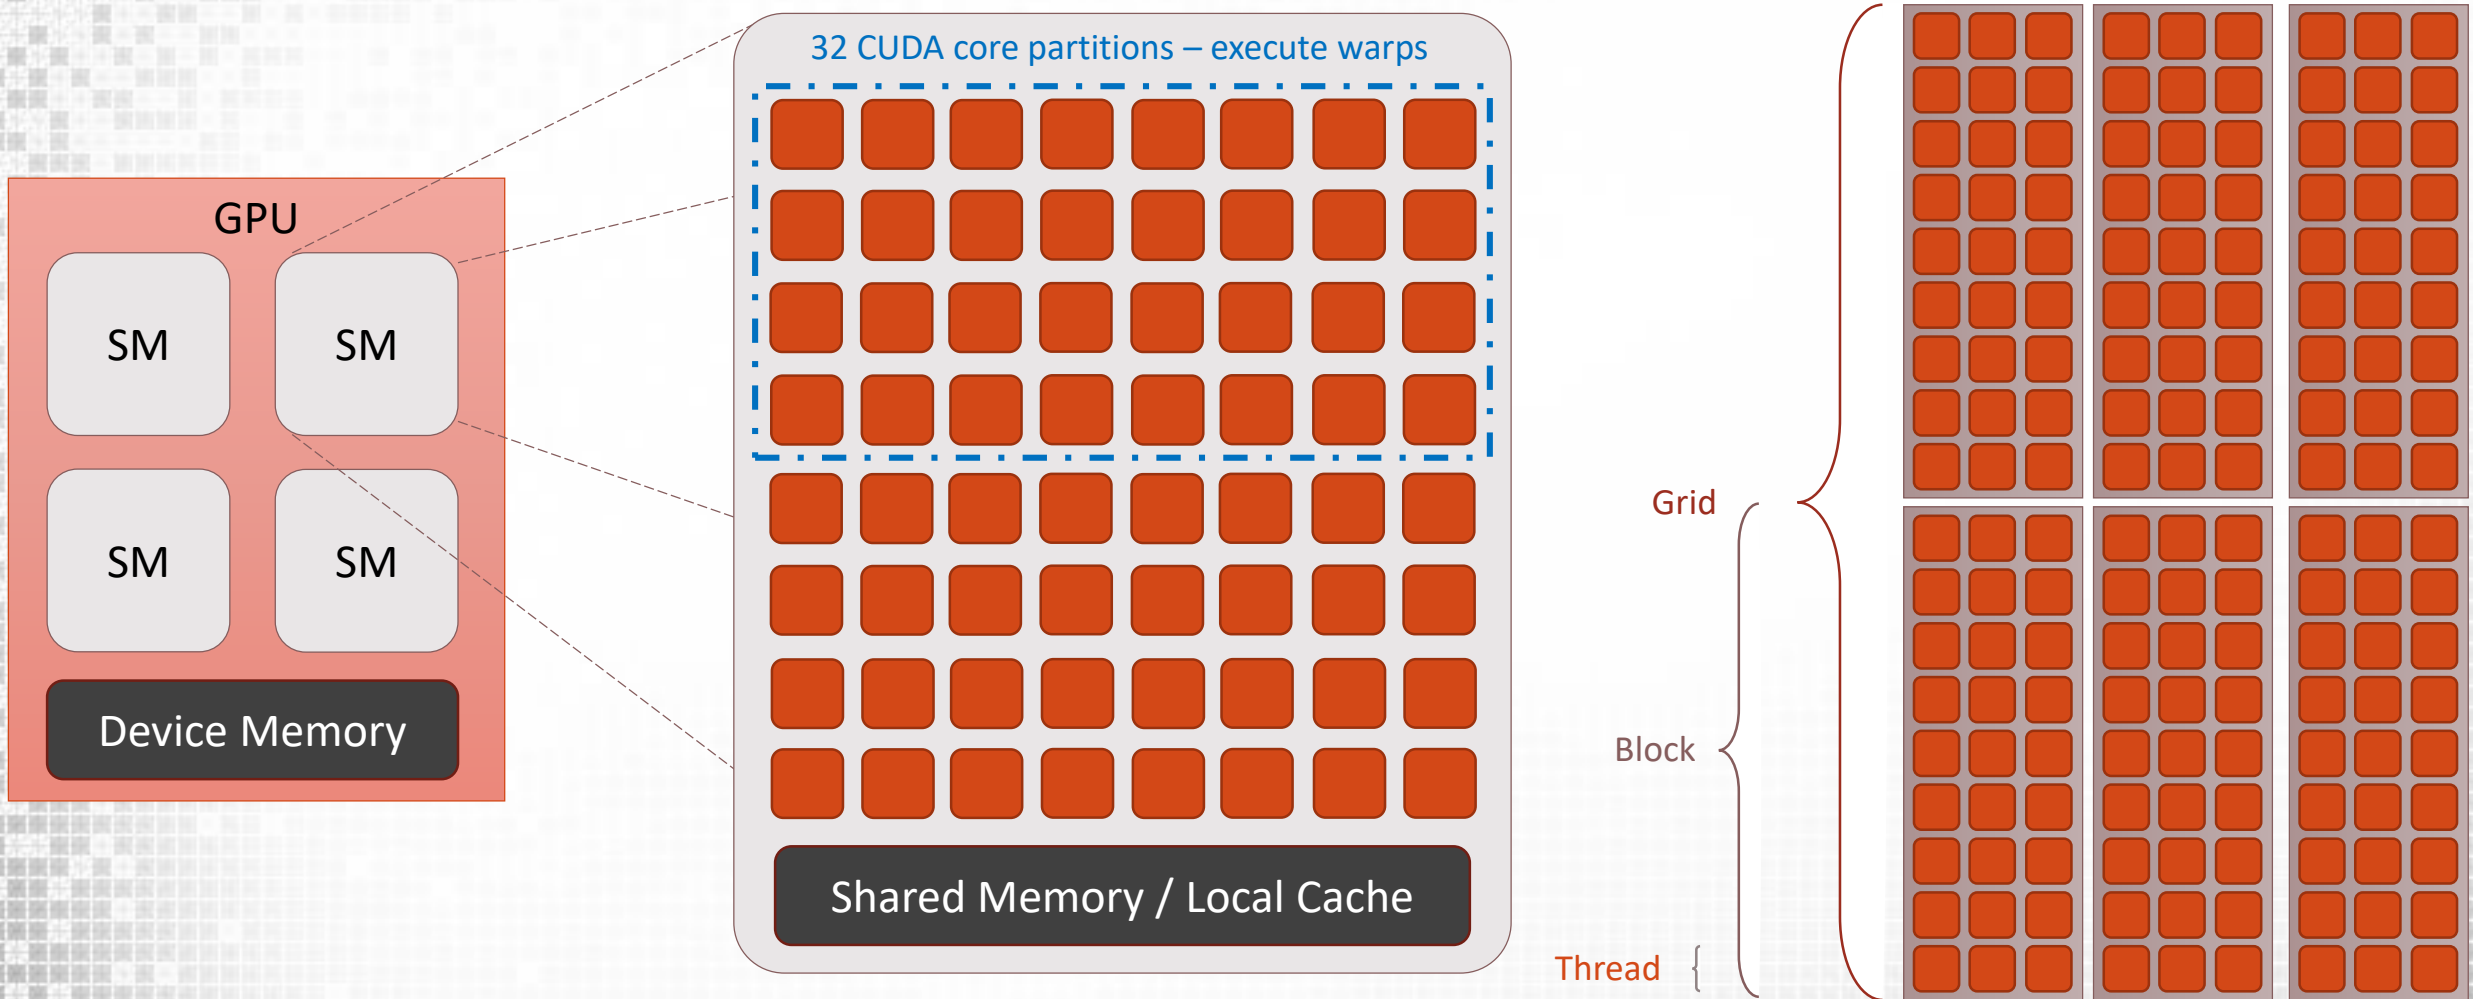


The
University
Of
Sheffield.

 NVIDIA

GPU
RESEARCH
CENTER

Grids, Blocks, Warps & Threads



Grids, Blocks, Warps & Threads

- ❑ Blocks map to SMs
 - ❑ SMs may have more than one block
 - ❑ Blocks are split into warps by hardware (always pick block size multiple of 32)
 - ❑ Blocks do not migrate between SMs
 - ❑ No guarantee of order of block execution
 - ❑ No communication or synchronisation between blocks
- ❑ Threads map to CUDA cores
 - ❑ Executed in partitions of 32, called warps
 - ❑ Lots of warps means lots of opportunity to hide memory movement

Review

- ❑ We have seen the importance of different types of memory
 - ❑ And observed the performance improvement from read-only and constant cache usage
- ❑ So far we have seen how CUDA can be used for performing thread local computations; e.g.
 - ❑ Load data from memory to registers
 - ❑ Perform thread-local computations
 - ❑ Store results back to global memory
- ❑ We will now consider another important type of memory
 - ❑ Shared memory

❑ Shared Memory

- ❑ Shared Memory Bank Conflicts

- ❑ 2D Shared Memory Bank Conflicts

- ❑ Boundary Conditions for Shared Memory Loading

- ❑ Host-side Configurations for Shared Memory

Shared Memory

❑ Architecture Details

- ❑ In Kepler (64KB) of Shared Memory is split between Shared Memory and L1 cache

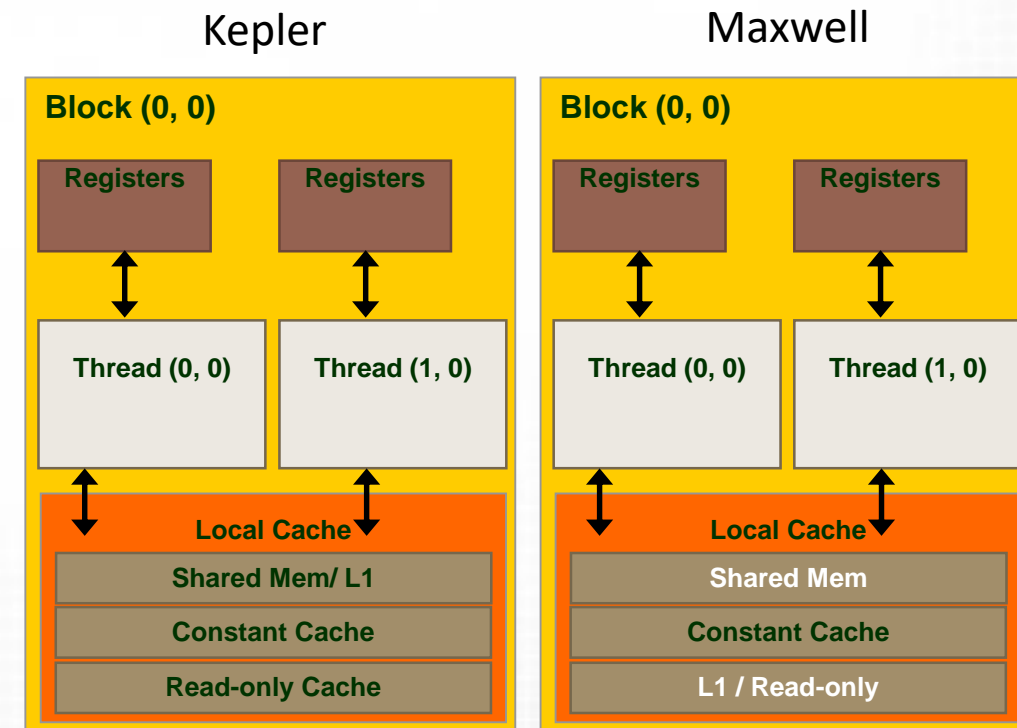
 - ❑ The ratio to SM and L1 can be configured

- ❑ In Maxwell 64KB of Shared Memory is dedicated

❑ Its just another Cache, right?

- ❑ User configurable

- ❑ Requires manually loading and synchronising data



Shared Memory

❑ Performance

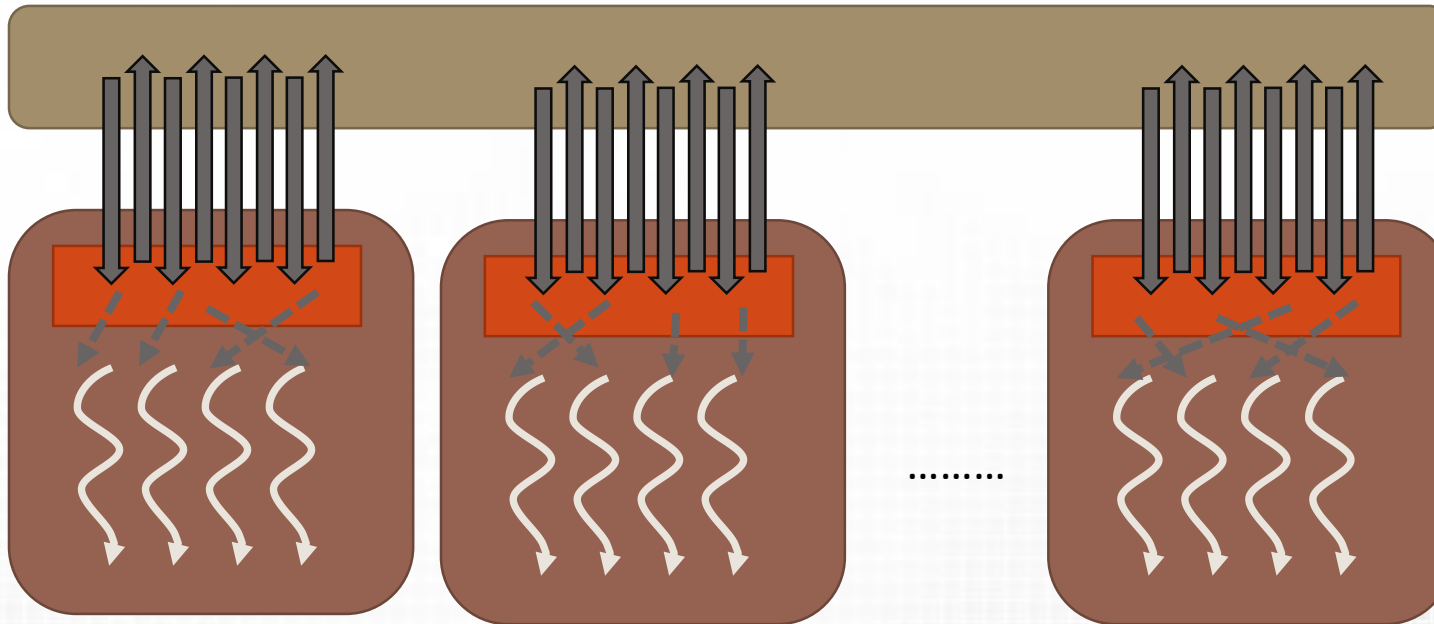
- ❑ Shared memory is very fast
- ❑ Bandwidth > 1 TB/s

❑ Block level computation

- ❑ Challenges the thread level view...
- ❑ Allows data to be shared between threads in the same block
- ❑ User configurable cache at the thread block level
- ❑ Still no broader synchronisation beyond the level of thread blocks

Block Local Computation

- ❑ Partition data into groups that fit into shared memory
- ❑ Load subset of data into shared memory
- ❑ Perform computation on the subset
- ❑ Copy subset back to global memory



Move, execute, move

- ☐ From Host view
 - ☐ Move: Data to GPU memory
 - ☐ Execute: Kernel
 - ☐ Move: Data back to host
- ☐ From Device view
 - ☐ Move: Data from device memory to registers
 - ☐ Execute: instructions
 - ☐ Move: Data back to device memory

Thread level parallelism

- ☐ From Host view
 - ☐ Move: Data to GPU memory
 - ☐ Execute: Kernel
 - ☐ Move: Data back to host
- ☐ From Device view
 - ☐ Move: Data from device memory to local cache
 - ☐ Execute: subset of kernel (reusing cached values)
 - ☐ Move: Data back to device memory
- ☐ From Block View
 - ☐ Move: Data from local cache
 - ☐ Execute: instructions
 - ☐ Move: Data back to local cache (or device memory)

Block level parallelism



A Case for Shared Memory

```
__global__ void sum3_kernel(int *c, int *a)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int left, right;

    //load value at i-1
    left = 0;
    if (i > 0)
        left = a[i - 1];

    //load value at i+1
    right = 0;
    if (i < (N - 1))
        right = a[i + 1];

    c[i] = left + a[i] + right; //sum three values
}
```

Do we have a candidate for block level parallelism using shared memory?

A Case for Shared Memory

```
__global__ void sum3_kernel(int *c, int *a)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int left, right;

    //load value at i-1
    left = 0;
    if (i > 0)
        left = a[i - 1];

    //load value at i+1
    right = 0;
    if (i < (N - 1))
        right = a[i + 1];

    c[i] = left + a[i] + right; //sum three values
}
```

- ❑ Currently: Thread-local computation
- ❑ Bandwidth limited
 - ❑ Requires three loads per thread (at index $i-1$, i , and $i+1$)
- ❑ Block level solution: load each value only once!

CUDA Shared memory

- ❑ Shared memory between threads in the same block can be defined using `__shared__`
- ❑ Shared variables are only accessible from within device functions
 - ❑ Not addressable in host code
- ❑ Must be careful to avoid race conditions
 - ❑ Multiple threads writing to the same shared memory variable
 - ❑ Results in undefined behaviour
 - ❑ Typically write to shared memory using `threadIdx`
 - ❑ Thread level synchronisation is available through `__syncthreads()`
 - ❑ Synchronises threads in the block

```
__shared__ int s_data[BLOCK_SIZE];
```




Example

```
__global__ void sum3_kernel(int *c, int *a)
{
    __shared__ int s_data[BLOCK_SIZE];

    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int left, right;

    s_data[threadIdx.x] = a[i];
    __syncthreads();

    //load value at i-1
    left = 0;
    if (i > 0){
        left = s_data[threadIdx.x - 1];
    }

    //load value at i+1
    right = 0;
    if (i < (N - 1)){
        right = s_data[threadIdx.x + 1];
    }

    c[i] = left + s_data[threadIdx.x] + right; //sum
}
```

- ❑ Allocate a shared array
 - ❑ One integer element per thread
- ❑ Each thread loads a single item to shared memory
- ❑ Call `__syncthreads` to ensure shared memory data is populated by all threads
- ❑ Load all elements through shared memory

What is wrong with this code?

Example

```
__global__ void sum3_kernel(int *c, int *a)
{
    __shared__ int s_data[BLOCK_SIZE];

    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int left, right;

    s_data[threadIdx.x] = a[i];
    __syncthreads();

    //load value at i-1
    left = 0;
    if (i > 0){
        if (threadIdx.x > 0)
            left = s_data[threadIdx.x - 1];
        else
            left = a[i - 1];
    }

    //load value at i+1
    right = 0;
    if (i < (N - 1)){
        if (threadIdx.x < (BLOCK_SIZE-1))
            right = s_data[threadIdx.x + 1];
        else
            right = a[i + 1];
    }

    c[i] = left + s_data[threadIdx.x] + right; //sum
}
```

❑ Additional step required!

❑ Check boundary conditions for the edge of the block

Problems with Shared memory

- ❑ In the example we saw the introduction of boundary conditions
 - ❑ Global loads still present at boundaries
 - ❑ We have introduced divergence in the code (remember the SIMD model)
 - ❑ This is even more prevalent in 2D examples where we *tile* data into shared memory

```
//boundary condition
left = 0;
if (i > 0){
    if (threadIdx.x > 0)
        left = s_data[threadIdx.x - 1];
    else
        left = a[i - 1];
}
```

❑ Shared Memory

❑ Shared Memory Bank Conflicts

❑ 2D Shared Memory Bank Conflicts

❑ Boundary Conditions for Shared Memory Loading

❑ Host-side Configurations for Shared Memory

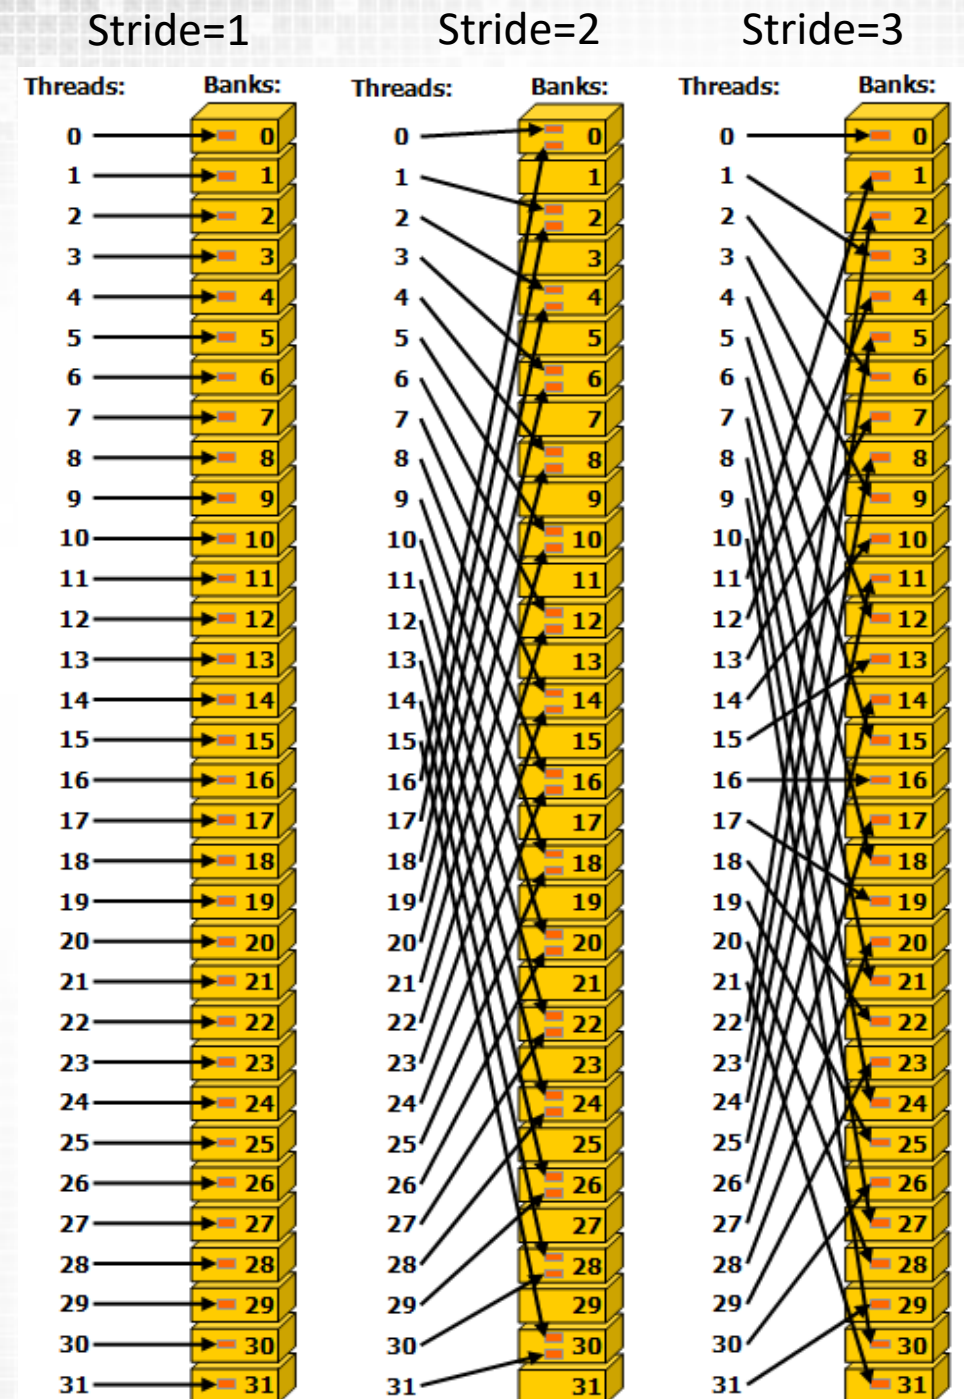
Shared Memory Bank Conflicts

- ❑ Shared memory is arranged into 4byte (32bit banks)
 - ❑ A load or store of N addresses spanning N distinct banks can be serviced simultaneously
 - ❑ Overall bandwidth of $\times N$ a single module
 - ❑ Kepler+ can also serve broadcast accesses simultaneously
- ❑ A bank conflict occurs when two threads request addresses from the same bank
 - ❑ Results in serialisation of the access
- ❑ Bank conflicts only occur between threads in a warp
 - ❑ There are 32 banks and 32 threads per warp (device of compute capability 2.0)
 - ❑ If two threads in a warp access the same bank this is said to be a 2-way bank conflict

Think about you block sized array of floats
 $\text{bank} = (\text{index} * \text{stride}) \% 32$

Access Strides

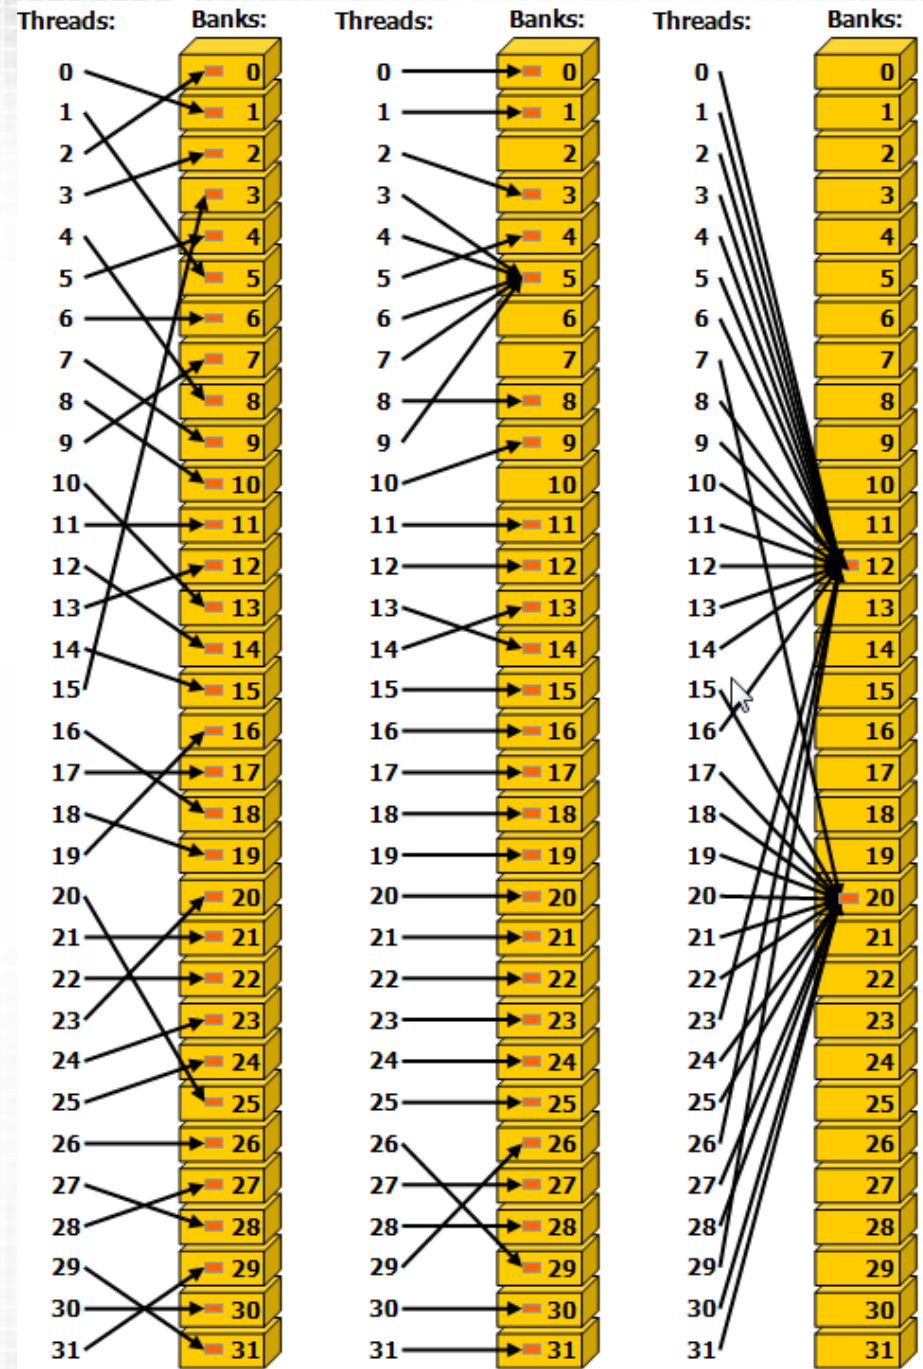
- ❑ Stride refers to the size (in increments of the bank size) between each threads memory access pattern
 - ❑ If threads access consecutive 4 byte values (e.g. `int` or `float`) then the stride is 1.
 - ❑ No conflicts
 - ❑ If a thread accesses consecutive 8 bytes values (e.g. `double`) then the stride is 2.
 - ❑ 2 way conflicts
 - ❑ In general odd strides result in no conflicts



More on SM banks

- ❑ Irregular access is fine as long as no bank conflicts
- ❑ Multiple threads can access the same bank conflict free **if** they access addresses in broadcast
- ❑ Broadcast can be to any number of threads in a warp

```
__shared__ float s_data[??];  
//read from shared memory using broadcast  
  
some_thread_value = s_data[0] ;
```





Strided access example

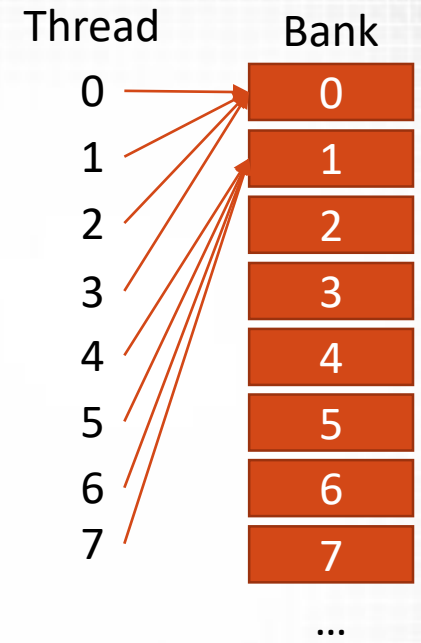
```
__shared__ char s_data[BLOCK_SIZE];  
  
//load or calculate some_thread_value  
  
s_data[threadIdx.x] = some_thread_value;  
__syncthreads();
```

Thread	Bank
0	0
1	1
2	2
3	3
4	4
5	5
6	6
7	7
...	...
31	31

- ❑ What is the stride?
- ❑ What is the level of conflict?
- ❑ How can this be improved?

Strided access example

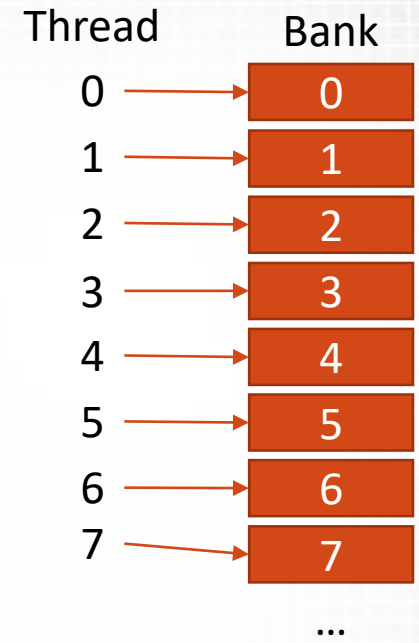
```
__shared__ char s_data[BLOCK_SIZE];  
  
//load or calculate some_thread_value  
  
s_data[threadIdx.x] = some_thread_value;  
__syncthreads();
```



- ❑ What is the stride? **Less than 1 (0.25)**
- ❑ What is the level of conflict? **4 way**
- ❑ How can this be improved? **Increase the stride**

Increase the stride (OK solution)

```
__shared__ char s_data[BLOCK_SIZE*4];  
  
//load or calculate some_thread_value  
  
s_data[threadIdx.x*4] = some_thread_value;  
__syncthreads();
```



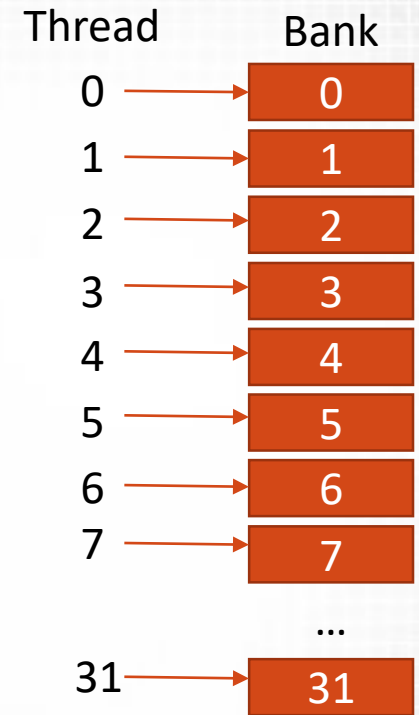
- ❑ What is the stride? **1**
- ❑ What is the level of conflict? **1 way (no conflict)**
- ❑ How can this be improved? **Use less memory!**

Increase the stride (good solution)

```
__shared__ char s_data[BLOCK_SIZE+1];  
  
//load or calculate some_thread_value  
  
s_data[CONFLICT_FREE(threadIdx.x)] = some_thread_value;  
__syncthreads();
```

where

```
#define CHAR_MULTIPLIER 4  
#define CONFLICT_FREE(x) (x*CHAR_MULTIPLIER % (BLOCK_SIZE+1))
```



- ❑ What is the stride? **1**
- ❑ What is the level of conflict? **1 way (no conflict)**
- ❑ How much shared memory is required? **BLOCK_SIZE+1**

Increase the stride (good solution)

stride	0.25			
multiplier	4			
block_size	128			
threadIdx.x			adjusted index	bank
0			0	0
1			4	1
2			8	2
3			12	3
4			16	4
5			20	5
6			24	6
7			28	7
8			32	8
9			36	9
10			40	10
...		
127			121	30
		Banks Used		32
		Max Conflicts		1

```
=(tid*multiplier) % (block_size+1)
```

$$= (\text{index} * \text{stride}) \% 32$$

$\square_{\text{BLOCK_SIZE}+1}$ unique indices

❑ Much better than
BLOCK SIZE * 4 unique indices

- ❑ Shared Memory
- ❑ Shared Memory Bank Conflicts
- ❑ 2D Shared Memory Bank Conflicts
- ❑ Boundary Conditions for Shared Memory Loading
- ❑ Host-side Configurations for Shared Memory

Bank conflicts with 2D tiles

```
__global__ void image_kernel(float *image)
{
    __shared__ float s_data[BLOCK_DIM][BLOCK_DIM];

    for (int i = 0; i < BLOCK_DIM; i++) {
        some_thread_value += f(s_data[threadIdx.x][i]);
    }
}
```

bank = threadIdx.x * stride % 32

- ❑ Example where each thread (of 2D block) operates on a row
- ❑ Loads values by column

Shared Memory Bank

0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		0

Bank conflicts with 2D tiles

```
__global__ void image_kernel(float *image)
{
    __shared__ float s_data[BLOCK_DIM][BLOCK_DIM];

    for (int i = 0; i < BLOCK_DIM; i++) {
        some_thread_value += f(s_data[threadIdx.x][i]);
    }
}
```

bank = threadIdx.x * stride % 32

BLOCK_DIM=32

i=0

Shared Memory Bank

0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		0

- ❑ Example where each thread (of 2D block) operates on a row
 - ❑ Loads values by column

- ❑ 32 way bank conflicts!
 - ❑ Very bad
- ❑ Stride = 32

Bank conflicts with 2D tiles

```
__global__ void image_kernel(float *image)
{
    __shared__ float s_data[BLOCK_DIM][BLOCK_DIM];

    for (int i = 0; i < BLOCK_DIM; i++) {
        some_thread_value += f(s_data[threadIdx.x][i]);
    }
}
```

bank = threadIdx.x * stride % 32

BLOCK_DIM=32

i=0

Shared Memory Bank

0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		0

- ❑ Example where each thread (of 2D block) operates on a row
 - ❑ Loads values by column

- ❑ How to fix
 - ❑ Memory padding
 - ❑ Transpose the matrix
 - ❑ Or operate on columns (loading by row) if possible

Bank conflicts with 2D tiles

```
__global__ void image_kernel(float *image)
{
    __shared__ float s_data[BLOCK_DIM][BLOCK_DIM+1];

    for (int i = 0; i < BLOCK_DIM; i++){
        some_thread_value += f(d_data[threadIdx.x][i]);
    }
}
```

Memory Padding Solution

bank = threadIdx.x * stride % 32

BLOCK_DIM+1=33

Shared Memory Bank

0	1	2	3	4	5	6	7		0
1	2	3	4	5	6	7	8		1
2	3	4	5	6	7	8	9		2
3	4	5	6	7	8	9	10		3
4	5	6	7	8	9	10	11		4
5	6	7	8	9	10	11	12		5
6	7	8	9	10	11	12	13		6
7	8	9	10	11	12	13	14		7
31	0	1	2	3	4	5	6		31

Bank conflicts with 2D tiles

```
__global__ void image_kernel(float *image)
{
    __shared__ float s_data[BLOCK_DIM][BLOCK_DIM+1];

    for (int i = 0; i < BLOCK_DIM; i++){
        some_thread_value += f(d_data[threadIdx.x][i]);
    }
}
```

bank = threadIdx.x * stride % 32

BLOCK_DIM+1=33

i=0

Shared Memory Bank

0	1	2	3	4	5	6	7		0
1	2	3	4	5	6	7	8		1
2	3	4	5	6	7	8	9		2
3	4	5	6	7	8	9	10		3
4	5	6	7	8	9	10	11		4
5	6	7	8	9	10	11	12		5
6	7	8	9	10	11	12	13		6
7	8	9	10	11	12	13	14		7
31	0	1	2	3	4	5	6		31

❑ Memory Padding Solution

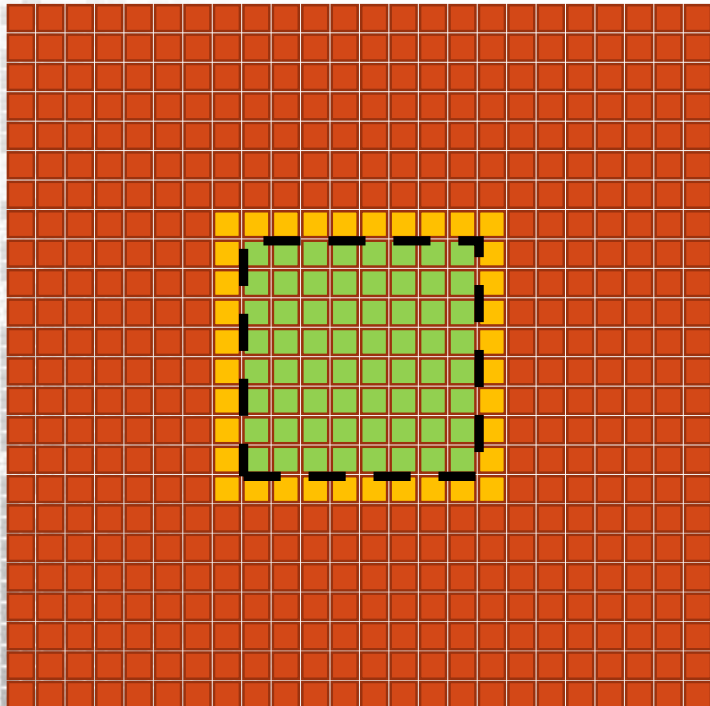
❑ Every thread in warp reads from different bank

❑ *Alternative: Transpose solution left to you!*

- ❑ Shared Memory
- ❑ Shared Memory Bank Conflicts
- ❑ 2D Shared Memory Bank Conflicts
- ❑ Boundary Conditions for Shared Memory Loading
- ❑ Host-side Configurations for Shared Memory

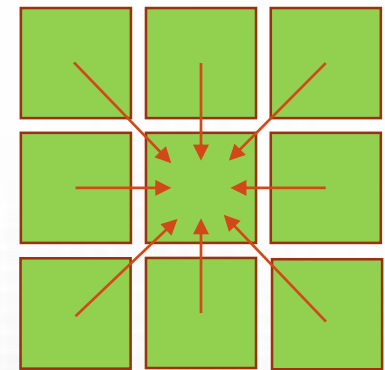
Boundary Conditions & Shared Memory Tiling

- ❑ Consider a 2D problem where data is gathered from neighbouring cells
 - ❑ Each cell reads 8 values (gather pattern)
 - ❑ Sounds like a good candidate for shared memory
 - ❑ We can tile data into memory



Thread Block size is 8x8

- Data tiled into shared memory
- Data not tiled into shared memory



Gather pattern

Problem with our tiling approach

1	2	3	3	3	3	3	3	2	1
2	4	6	6	6	6	6	6	4	2
3	6	8	8	8	8	8	8	6	3
3	6	8	8	8	8	8	8	6	3
3	6	8	8	8	8	8	8	6	3
3	6	8	8	8	8	8	8	6	3
3	6	8	8	8	8	8	8	6	3
3	6	8	8	8	8	8	8	6	3
2	4	6	6	6	6	6	6	4	2
1	2	3	3	3	3	3	3	2	1

- ❑ Memory access pattern is good for values inside the boundary
 - ❑ 448 cached reads
 - ❑ 64 loads
- ❑ Memory outside of boundary is loaded multiple times
 - ❑ 92 un-cached reads
 - ❑ 92 loads

Boundary Condition Improvements

❑ Launch more threads

- ❑ Launch thread block of $DIM+2 \times DIM+2$
- ❑ Allocate one element of space per thread in SM
- ❑ Every thread loads one value
- ❑ Only threads in inner $DIM \times DIM$ compute values
 - ❑ Causes under utilisation

$$Utilisation = \frac{DIM^2}{(DIM + 2)^2}$$

DIM	Utilisation
8	64%
12	73%
16	79%
20	83%
24	85%
28	87%
32	89%
36	90%
40	91%
44	91%
48	92%

❑ Use more shared memory per thread

- ❑ Launch same $DIM \times DIM$ threads
- ❑ Allocate $DIM+2 \times DIM+2$ elements of space in SM
- ❑ Threads on boundary load multiple elements
 - ❑ Causes unbalanced loads
- ❑ All threads perform compute values

- ❑ Shared Memory
- ❑ Shared Memory Bank Conflicts
- ❑ 2D Shared Memory Bank Conflicts
- ❑ Boundary Conditions for Shared Memory Loading
- ❑ Host-side Configurations for Shared Memory

Dynamically Assigned Shared Memory

- ❑ It is possible to dynamically assign shared memory at runtime.
- ❑ Requires both a host and device modification to code
 - ❑ Device: Must declare shared memory as extern
 - ❑ Host: Must declare shared memory size in kernel launch parameters

```
unsigned int sm_size = sizeof(float)*DIM*DIM;
image_kernel<<<blocksPerGrid, threadsPerBlock, sm_size >>>(d_image);

__global__ void image_kernel(float *image)
{
    extern __shared__ float s_data[];
}
```

Is equivalent to

```
image_kernel<<<blocksPerGrid, threadsPerBlock>>>(d_image);

__global__ void image_kernel(float *image)
{
    __shared__ float *s_data[DIM][DIM];
}
```


Summary

- ❑ Shared Memory introduces the idea of block level computation rather than just thread level computation
- ❑ Shared Memory is a limited resource but can be very useful for reducing global memory bandwidth
 - ❑ Where data is reused
- ❑ Shared Memory requires user synchronisation unlike other general purpose caches (i.e. L1, texture)
- ❑ For optimal performance memory banks must be considered and boundary conditions must be handled
- ❑ There are hardware specific options for configuring how Shared Memory is used

Acknowledgements and Further Reading

- ❑ <http://cuda-programming.blogspot.co.uk/2013/02/bank-conflicts-in-shared-memory-in-cuda.html>
- ❑ <http://acceleware.com/blog/maximizing-shared-memory-bandwidth-nvidia-kepler-gpus>