

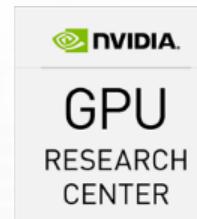
Parallel Computing with GPUs

Dr Mozhgan Kabiri Chimeh

<http://mkchimeh.staff.shef.ac.uk/teaching/COM4521>



The
University
Of
Sheffield.



Last Week

- ❑ We learnt about warp level CUDA
- ❑ How threads are scheduled and executed
 - ❑ Impacts of divergence
- ❑ Atomics: Good and bad...
- ❑ Do the warp shuffle!
- ❑ Parallel primitives
- ❑ Scan and Reduction



Credits

- ❑ The code and much of the content from this lecture is based on the GTC2016 Talk by C. Angerer and J. Progsch (NVIDIA)
 - ❑ [S6112 – CUDA Optimisation with NVIDIA Nsight for Visual Studio](#)
 - ❑ Provided by NVIDIA with thanks to Joe Bungo
- ❑ Content has been adapted to use Visual Profiler Guided Analysis where possible
- ❑ Additional steps and analysis have been added

Learning Objectives

- ❑ Understand the key performance metrics of GPU code.
- ❑ Understand profiling metrics and relate this to approaches which they have already learnt to address limiting factors in their code.
- ❑ Appreciate memory vs compute bound code and be able to recognise factors which contribute to this.

- ❑ Profiling Introduction
- ❑ The Problem
- ❑ Visual Profiler Guided Analysis
 - ❑ Iteration 1
 - ❑ Iteration 2
 - ❑ Iteration 3
 - ❑ Iteration 4



The
University
Of
Sheffield.



NVIDIA
GPU
RESEARCH
CENTER

The APOD Cycle

4. Deploy
and Test



1. Assess

- Identify Performance Limiter
- Analyze Profile
- Find Indicators

2. Parallelize

3. Optimize



3b. Build Knowledge

<https://devblogs.nvidia.com/assess-parallelize-optimize-deploy/>

CUDA Profiling Options

Visual Profiler (1st choice)

- Stand alone cross platform (java on Eclipse) program
- Guided performance analysis
- Links to CUDA best practice guide

NVProf

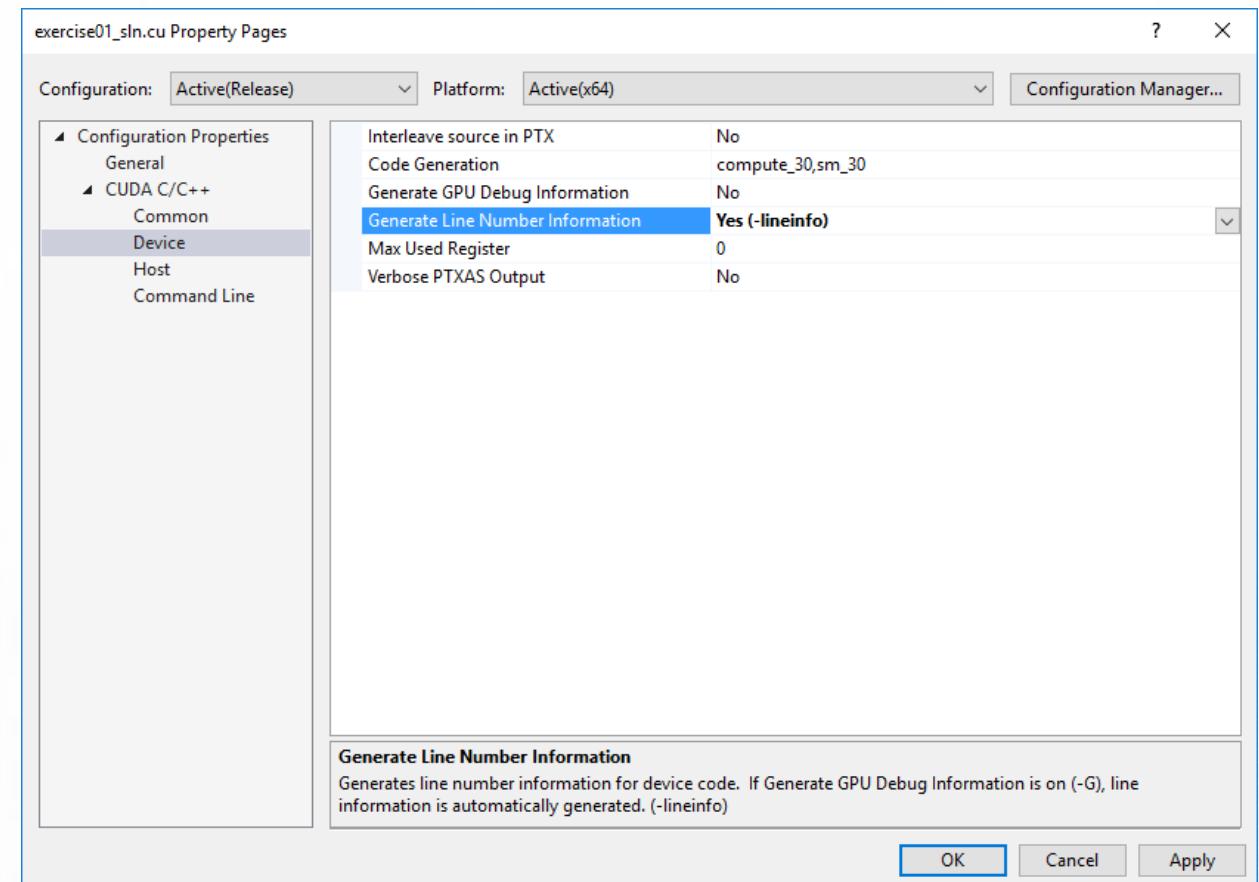
- Command line profiler
- Results can be visualised in Visual Profiler

Visual Studio Nsight Profiler

- Built into visual studio
- Detailed kernel and source level analysis (more than Visual Profiler)
- Unguided

Changes to your code

- ❑ If you want to associate profile information with source line
 - ❑ --lineinfo argument
 - ❑ Works in release mode
- ❑ Must flush GPU buffers
 - ❑ cudaDeviceReset ()
 - ❑ At end of program





Conveyor belt model

- ❑ Our GPU program is like a factory assembly line
 - ❑ Data in and data out (in a new form)
 - ❑ Skilled operators (multi processors) doing stuff with chunks of the data
 - ❑ Both the belt and people have maximum operating speed

- ❑ Ideal situation
 - ❑ Conveyor belt runs at full speed
 - ❑ Skilled operators always 100% busy

What is likely to effect this model?

Potential Performance Limiters

Memory

- Program limited by memory bandwidth
- Can't get data to the registers on the device fast enough
- Are you using lots of global memory but not faster local memory caches?*
- Have you exceed the amount of cache available?*

Compute

- Memory bandwidth well below peak
- GPU is too busy performing compute
- Have you got high levels of divergence with low warp execution efficiency ?*

Latency

- Poor occupancy = not enough active threads
- Instruction execution stalls due to poor memory access patterns (sparse or poorly used data)
- Is problem size or block size too small? Are you using the memory bandwidth effectively (cache line utilisation)?*

- ❑ Profiling Introduction
- ❑ The Problem
- ❑ Visual Profiler Guided Analysis

- ❑ Iteration 1
- ❑ Iteration 2
- ❑ Iteration 3
- ❑ Iteration 4



The
University
Of
Sheffield.

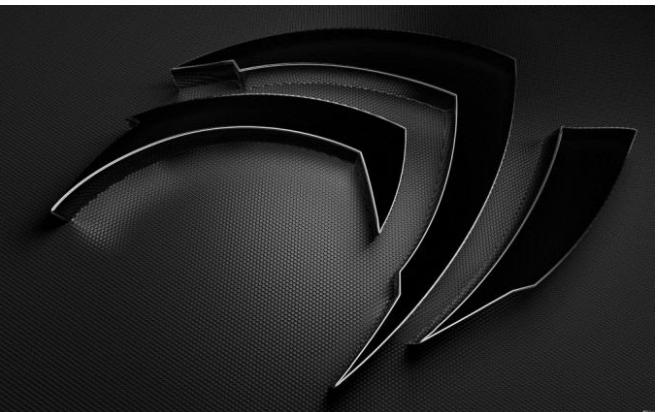


NVIDIA
GPU
RESEARCH
CENTER

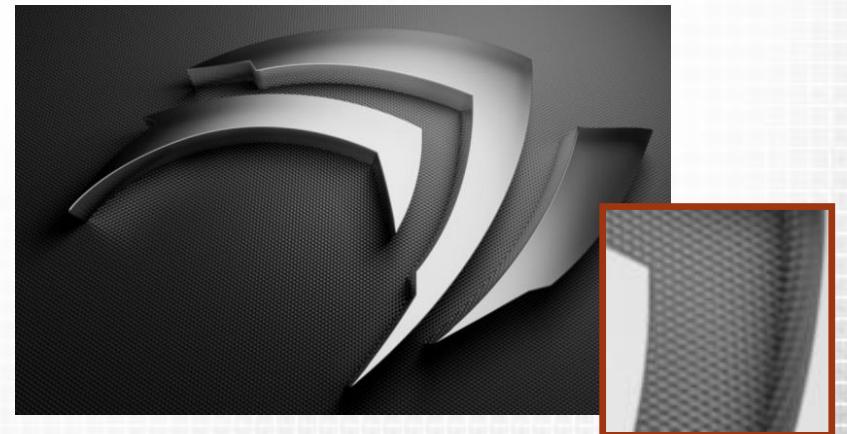
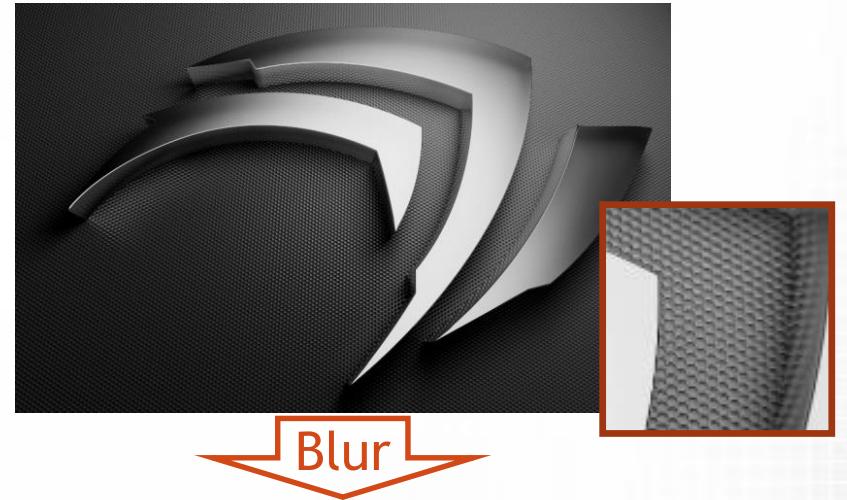
Introducing the Application



Grayscale



Edges



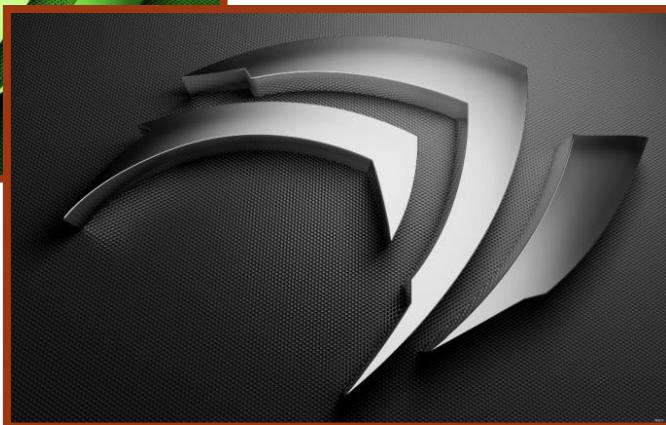
The
University
Of
Sheffield.



GPU
RESEARCH
CENTER

Introducing the Application

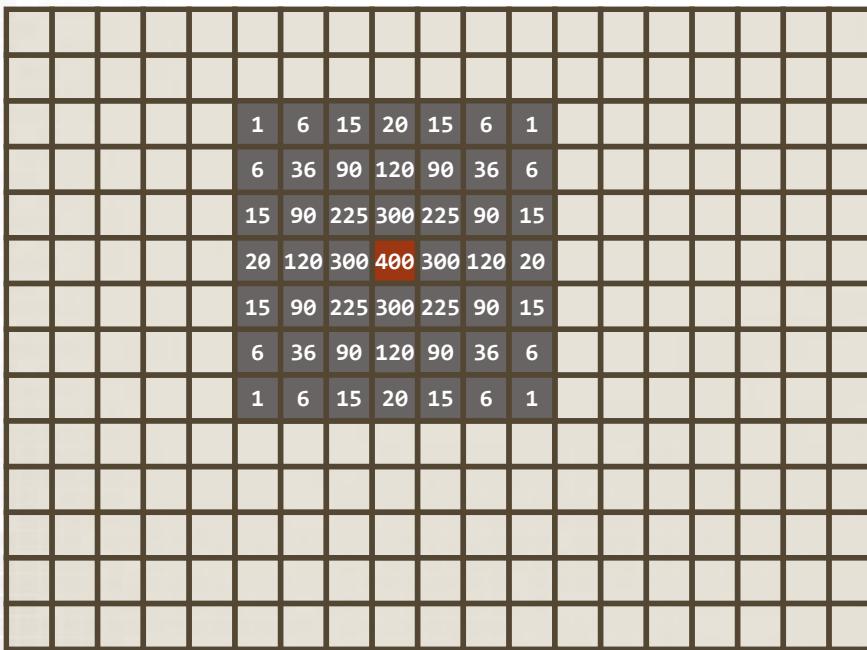
❑ Grayscale Conversion



```
// r, g, b: Red, green, blue components of the pixel p  
foreach pixel p:  
    p = 0.298839f*r + 0.586811f*g + 0.114350f*b;
```

Introducing the Application

❑ Blur: 7x7 Gaussian Filter



foreach pixel p:

$G = \text{weighted sum of } p \text{ and its 48 neighbors}$
 $p = G/256$

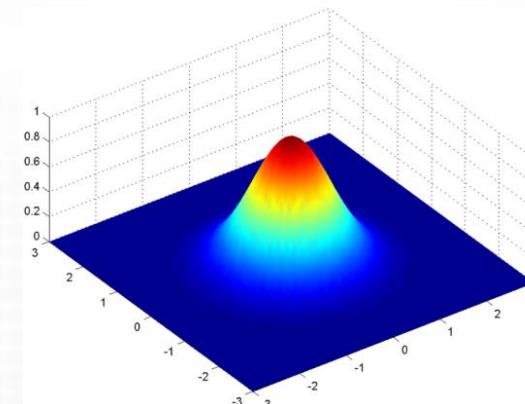


Image from Wikipedia



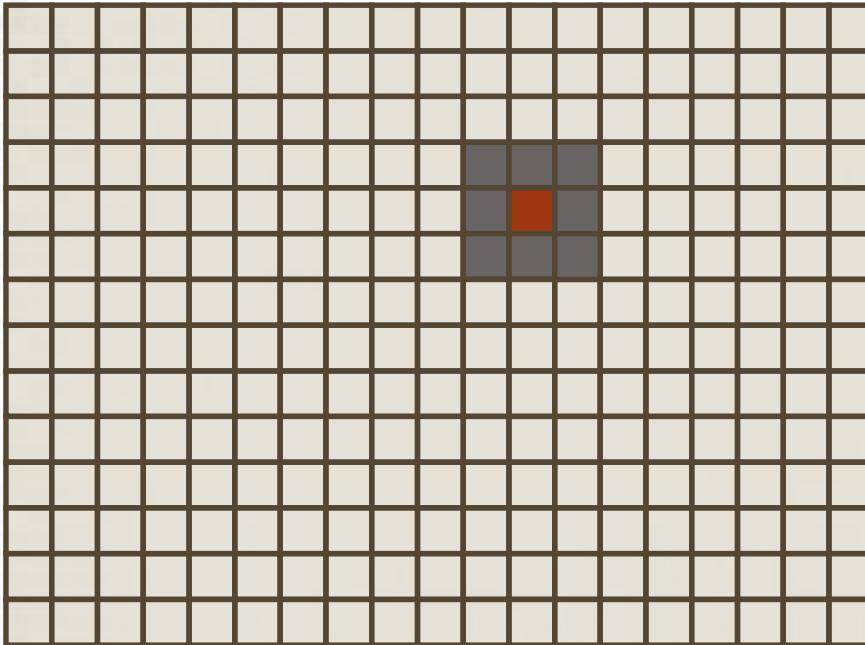
The
University
Of
Sheffield.



NVIDIA
GPU
RESEARCH
CENTER

Introducing the Application

□ Edges: 3x3 Sobel Filters



foreach pixel p :

G_x = weighted sum of p and its 8 neighbors

G_y = weighted sum of p and its 8 neighbors

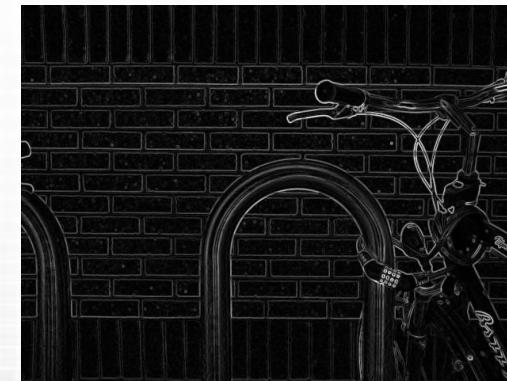
p = **sqrt**($G_x + G_y$)

Weights for G_x :

-1	0	1
-2	0	2
-1	0	1

Weights for G_y :

1	2	1
0	0	0
-1	-2	-1



The
University
Of
Sheffield.

 NVIDIA
GPU
RESEARCH
CENTER



The Starting Code

```
void gaussian_filter_7x7_v0(int w, int h, const uchar *src, uchar *dst)
{
    // Position of the thread in the image.
    const int x = blockIdx.x*blockDim.x + threadIdx.x;
    const int y = blockIdx.y*blockDim.y + threadIdx.y;

    // Early exit if the thread is not in the image.
    if( !in_img(x, y, w, h) )
        return;

    // Load the 48 neighbours and myself.
    int n[7][7];
    for( int j = -3 ; j <= 3 ; ++j )
        for( int i = -3 ; i <= 3 ; ++i )
            n[j+3][i+3] = in_img(x+i, y+j, w, h) ? (int) src[(y+j)*w + (x+i)] : 0;

    // Compute the convolution.
    int p = 0;
    for( int j = 0 ; j < 7 ; ++j )
        for( int i = 0 ; i < 7 ; ++i )
            p += gaussian_filter[j][i] * n[j][i];

    // Store the result.
    dst[y*w + x] = (uchar) (p / 256);
}
```

What is good and
what is bad?

❑ <https://github.com/chmaruni/nsight-gtc>



The Starting Code

```
void gaussian_filter_7x7_v0(int w, int h, const uchar *src, uchar *dst)
{
    // Position of the thread in the image.
    const int x = blockIdx.x*blockDim.x + threadIdx.x;
    const int y = blockIdx.y*blockDim.y + threadIdx.y;

    // Early exit if the thread is not in the image.
    if( !in_img(x, y, w, h) )
        return;

    // Load the 48 neighbours and myself.
    int n[7][7];
    for( int j = -3 ; j <= 3 ; ++j )
        for( int i = -3 ; i <= 3 ; ++i )
            n[j+3][i+3] = in_img(x+i, y+j, w, h) ? (int) src[(y+j)*w + (x+i)] : 0;

    // Compute the convolution.
    int p = 0;
    for( int j = 0 ; j < 7 ; ++j )
        for( int i = 0 ; i < 7 ; ++i )
            p += gaussian_filter[j][i] * n[j][i];

    // Store the result.
    dst[y*w + x] = (uchar) (p / 256);
}
```

What is good and
what is bad?

☐ <https://github.com/chmaruni/nsight-gtc>

Profiling Machine

- ❑ NVIDIA GeForce GTX980
 - ❑ GM200
 - ❑ Compute Capability SM5.2
- ❑ CUDA 7.0
- ❑ Windows 7
- ❑ Visual Studio 2013
- ❑ Nsight Visual Studio Edition 5.0



The
University
Of
Sheffield.

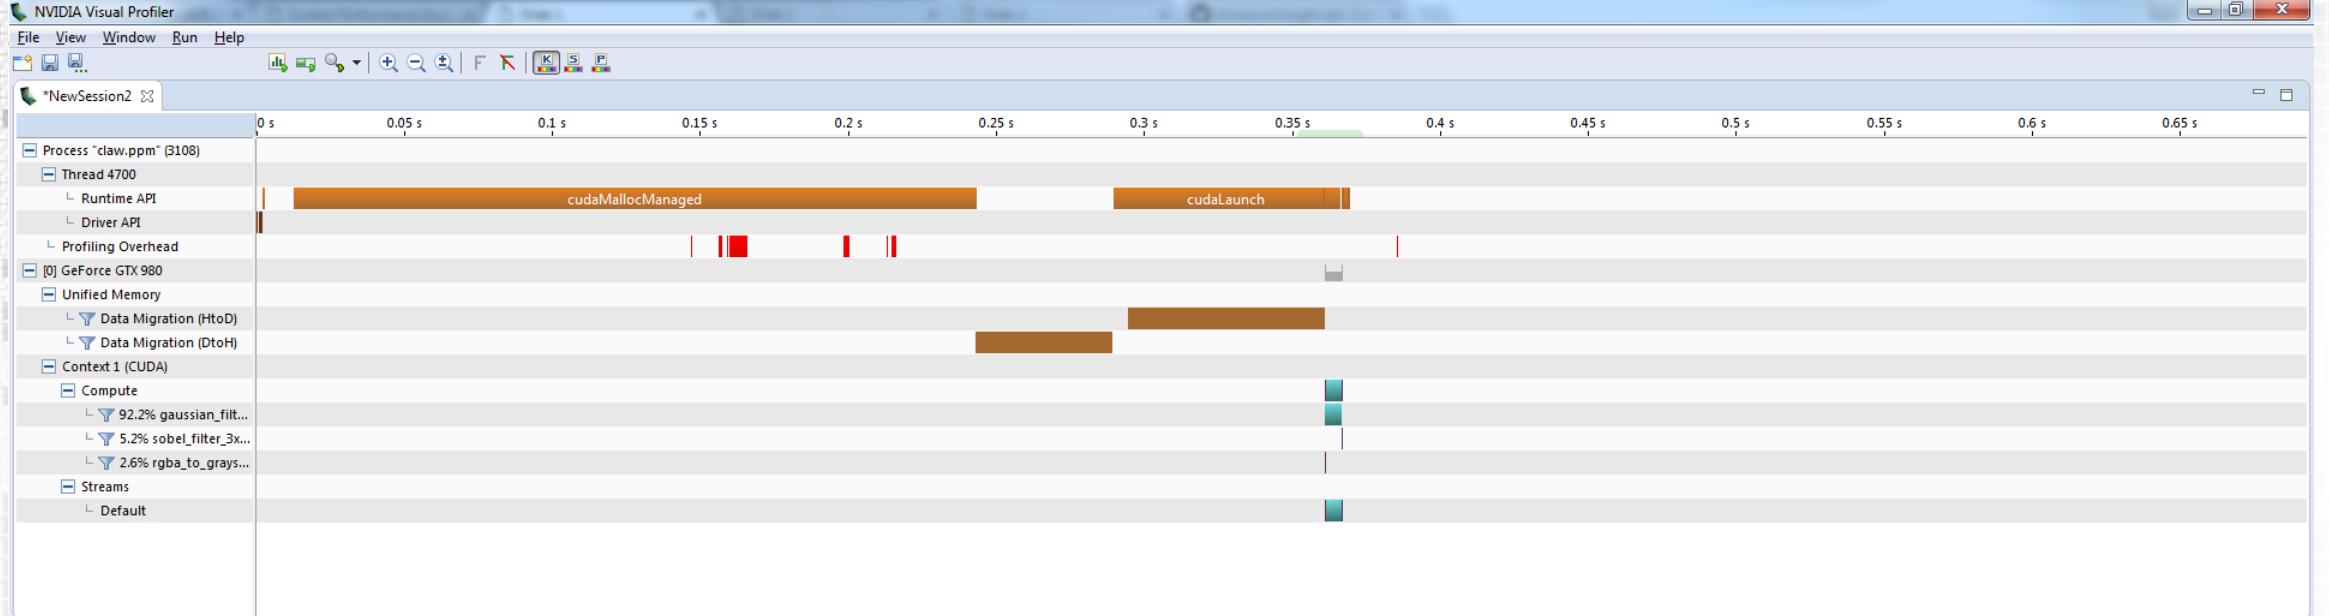


- ❑ Profiling Introduction
- ❑ The Problem
- ❑ Visual Profiler Guided Analysis
 - ❑ Iteration 1



The
University
Of
Sheffield.





Analysis Details Console Settings

Export PDF Report

1. CUDA Application Analysis

2. Check Overall GPU Usage

The analysis results on the right indicate potential problems in how your application is taking advantage of the GPU's available compute and data movement capabilities. You should examine the information provided with each result to determine if you can make changes to your application to increase GPU utilization.

Examine Individual Kernels

You can also examine the performance of individual kernels to expose additional optimization opportunities.

Results

Low Kernel Concurrency [0 ns / 5.94 ms = 0%]
The percentage of time when two kernels are being executed in parallel is low.
[More...](#)

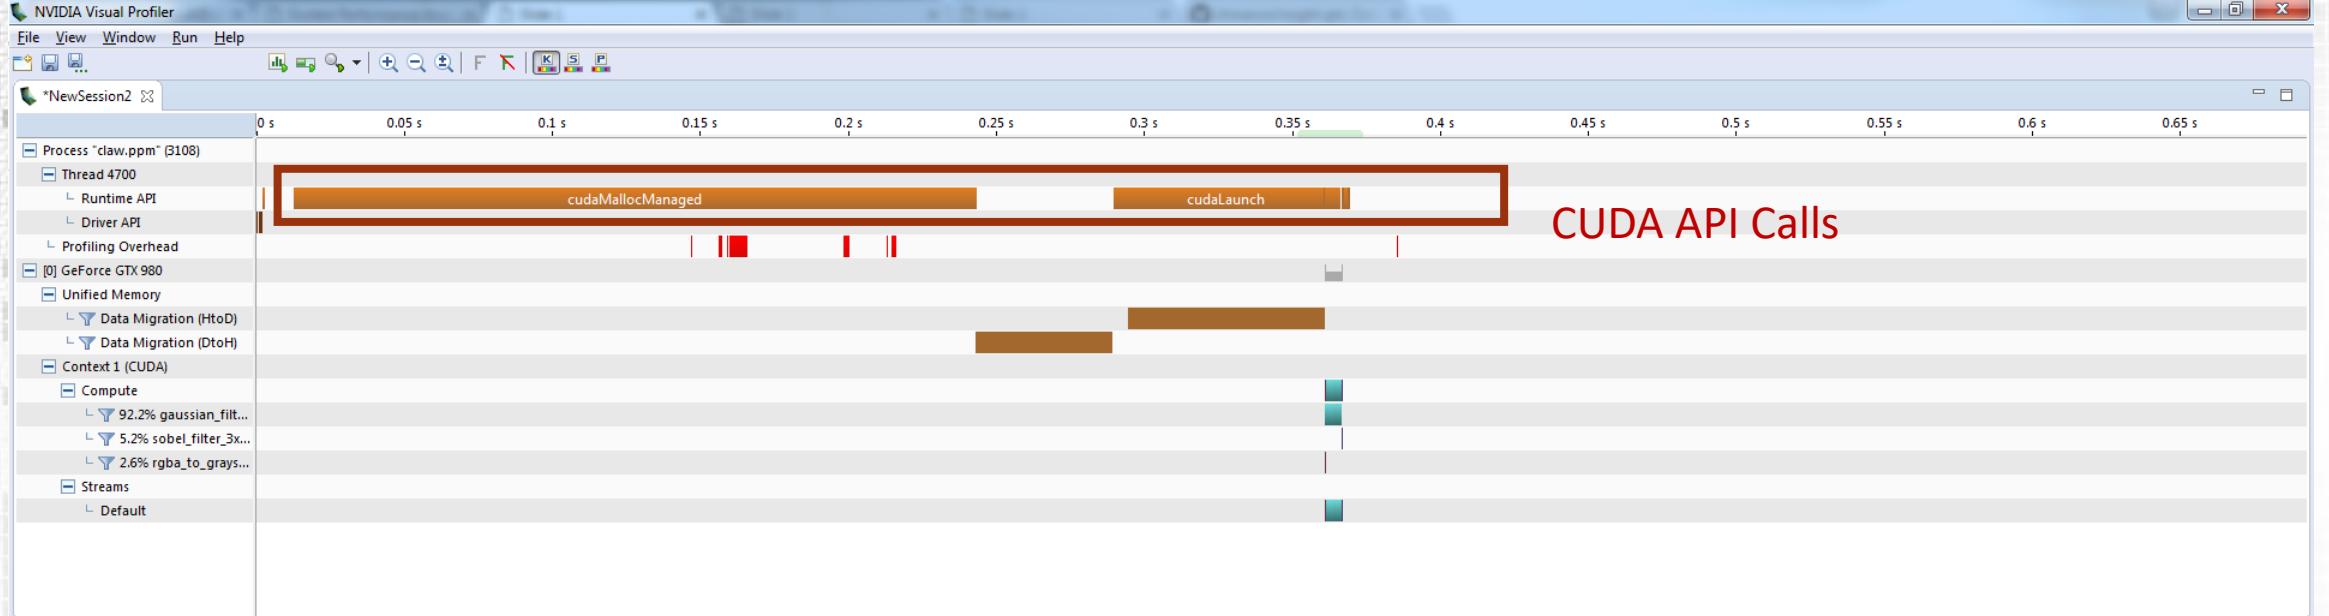
Low Compute Utilization [5.94 ms / 385.676 ms = 1.5%]
The multiprocessors of one or more GPUs are mostly idle.
[More...](#)

Compute Utilization
The device timeline shows an estimate of the amount of the total compute capacity being used by the kernels executing on the device.

Properties

Default

Duration	Session
385.676 n	



CUDA API Calls

Analysis x Details Console Settings

Export PDF Report

1. CUDA Application Analysis

2. Check Overall GPU Usage

The analysis results on the right indicate potential problems in how your application is taking advantage of the GPU's available compute and data movement capabilities. You should examine the information provided with each result to determine if you can make changes to your application to increase GPU utilization.

Examine Individual Kernels

You can also examine the performance of individual kernels to expose additional optimization opportunities.

Results

Low Kernel Concurrency [0 ns / 5.94 ms = 0%]
The percentage of time when two kernels are being executed in parallel is low.
[More...](#)

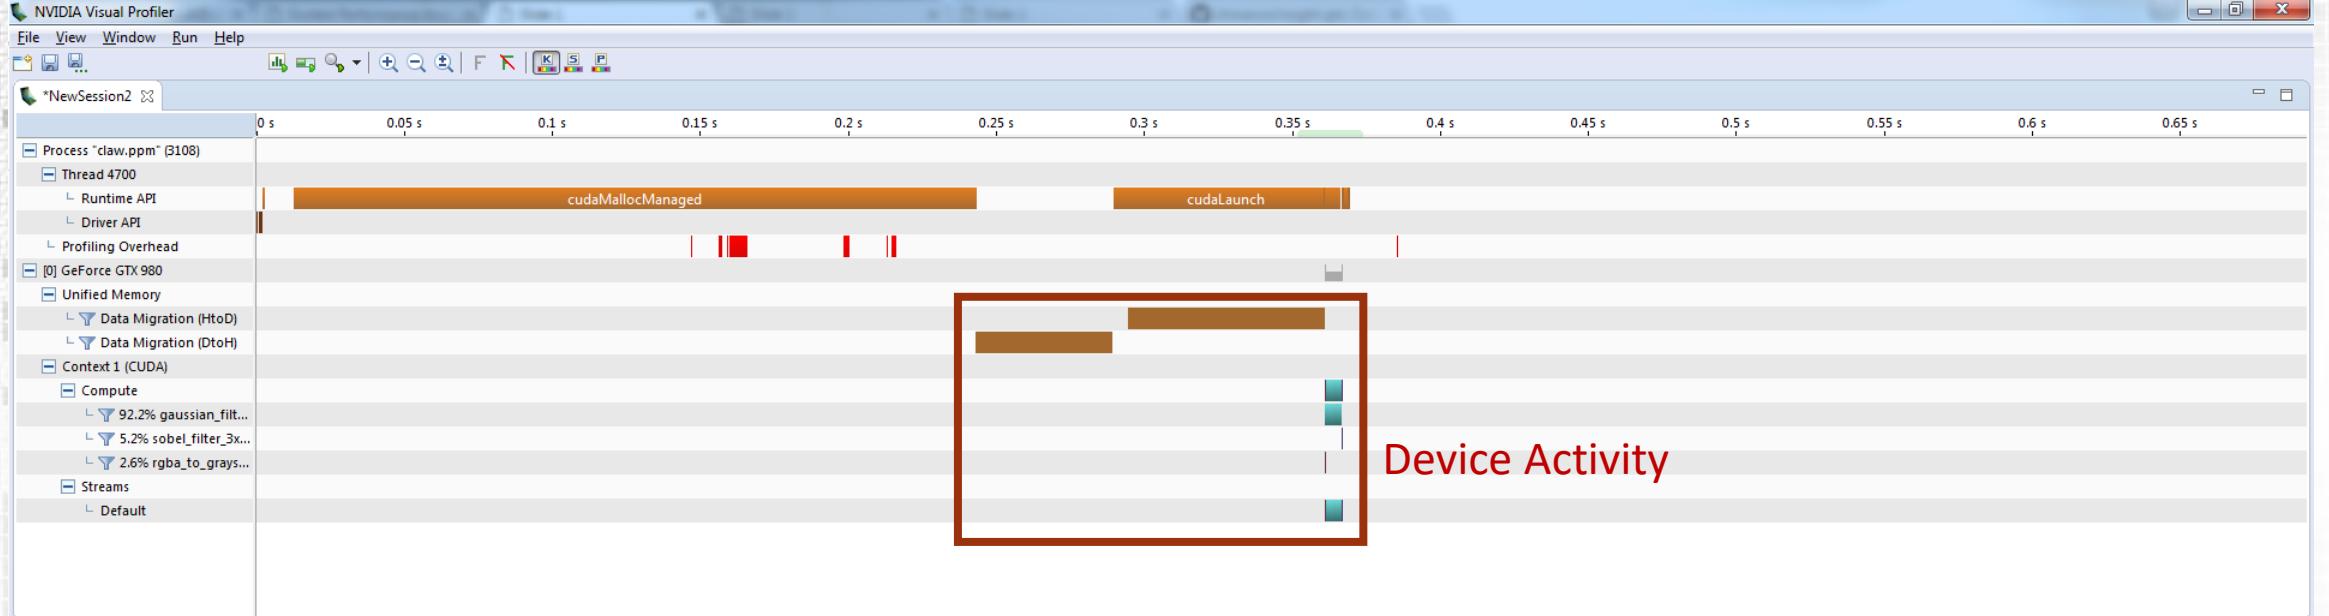
Low Compute Utilization [5.94 ms / 385.676 ms = 1.5%]
The multiprocessors of one or more GPUs are mostly idle.
[More...](#)

Compute Utilization
The device timeline shows an estimate of the amount of the total compute capacity being used by the kernels executing on the device.

Properties x

Default

Duration	Session
385.676 n	



Device Activity

Analysis Details Console Settings

Export PDF Report

1. CUDA Application Analysis

2. Check Overall GPU Usage

The analysis results on the right indicate potential problems in how your application is taking advantage of the GPU's available compute and data movement capabilities. You should examine the information provided with each result to determine if you can make changes to your application to increase GPU utilization.

Examine Individual Kernels

You can also examine the performance of individual kernels to expose additional optimization opportunities.

Results

Low Kernel Concurrency [0 ns / 5.94 ms = 0%]
The percentage of time when two kernels are being executed in parallel is low.
[More...](#)

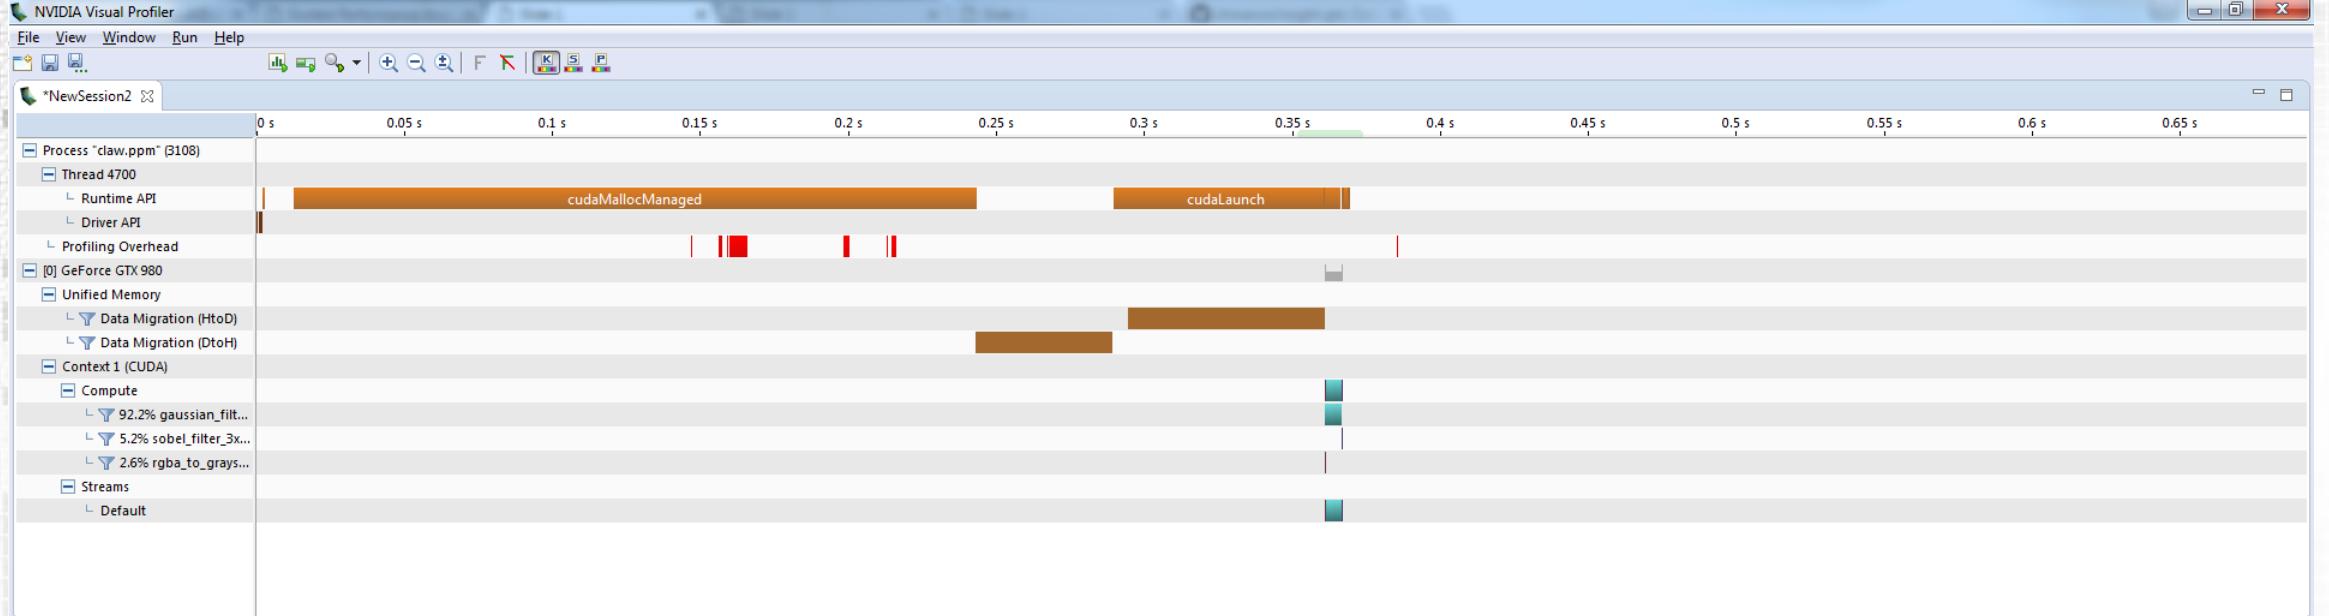
Low Compute Utilization [5.94 ms / 385.676 ms = 1.5%]
The multiprocessors of one or more GPUs are mostly idle.
[More...](#)

Compute Utilization
The device timeline shows an estimate of the amount of the total compute capacity being used by the kernels executing on the device.

Properties

Default

Duration	Session
385.676 n	



Analysis Details Console Settings

[Export PDF Report](#)

1. CUDA Application Analysis

2. Check Overall GPU Usage

The analysis results on the right indicate potential problems in how your application is taking advantage of the GPU's available compute and data movement capabilities. You should examine the information provided with each result to determine if you can make changes to your application to increase GPU utilization.

[Examine Individual Kernels](#)

You can also examine the performance of individual kernels to expose additional optimization opportunities.

Results

Low Kernel Concurrency [0 ns / 5.94 ms = 0%]
The percentage of time when two kernels are being executed in parallel is low.
[More...](#)

Low Compute Utilization [5.94 ms / 385.676 ms = 1.5%]
The multiprocessors of one or more GPUs are mostly idle.
[More...](#)

Compute Utilization
The device timeline shows an estimate of the amount of the total compute capacity being used by the kernels executing on the device.

Properties

Default

Duration	Session
385.676 n	

Hints

Results

⚠️ Low Kernel Concurrency [0 ns / 5.94 ms = 0%]

The percentage of time when two kernels are being executed in parallel is low.

[More...](#)

- We are using only a single stream
- Kernels have data dependencies so cant be executed in parallel

⚠️ Low Compute Utilization [5.94 ms / 385.676 ms = 1.5%]

The multiprocessors of one or more GPUs are mostly idle.

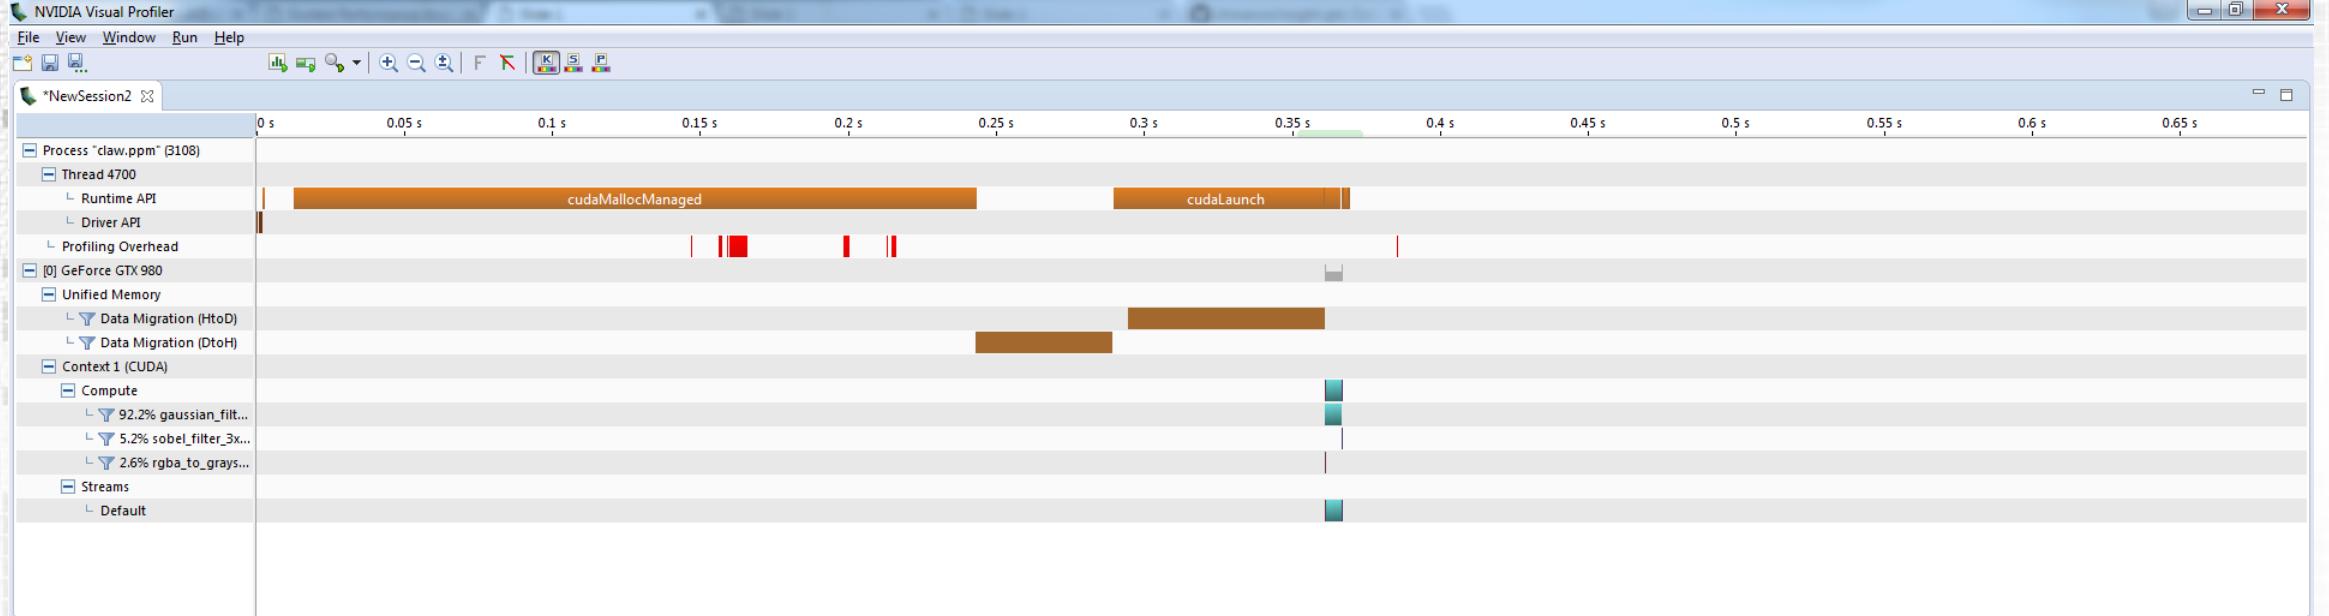
[More...](#)

- This is a problem
- The guided analysis will try and address this



The
University
Of
Sheffield.





Analysis Details Console Settings

Export PDF Report

Results

1. CUDA Application Analysis

2. Check Overall GPU Usage

The analysis results on the right indicate potential problems in how your application is taking advantage of the GPU's available compute and data movement capabilities. You should examine the information provided with each result to determine if you can make changes to your application to increase GPU utilization.

Examine Individual Kernels

You can also examine the performance of individual kernels to expose additional optimization opportunities.

Low Kernel Concurrency [0 ns / 5.94 ms = 0%]
The percentage of time when two kernels are being executed in parallel is low.
[More...](#)

Low Compute Utilization [5.94 ms / 385.676 ms = 1.5%]
The multiprocessors of one or more GPUs are mostly idle.
[More...](#)

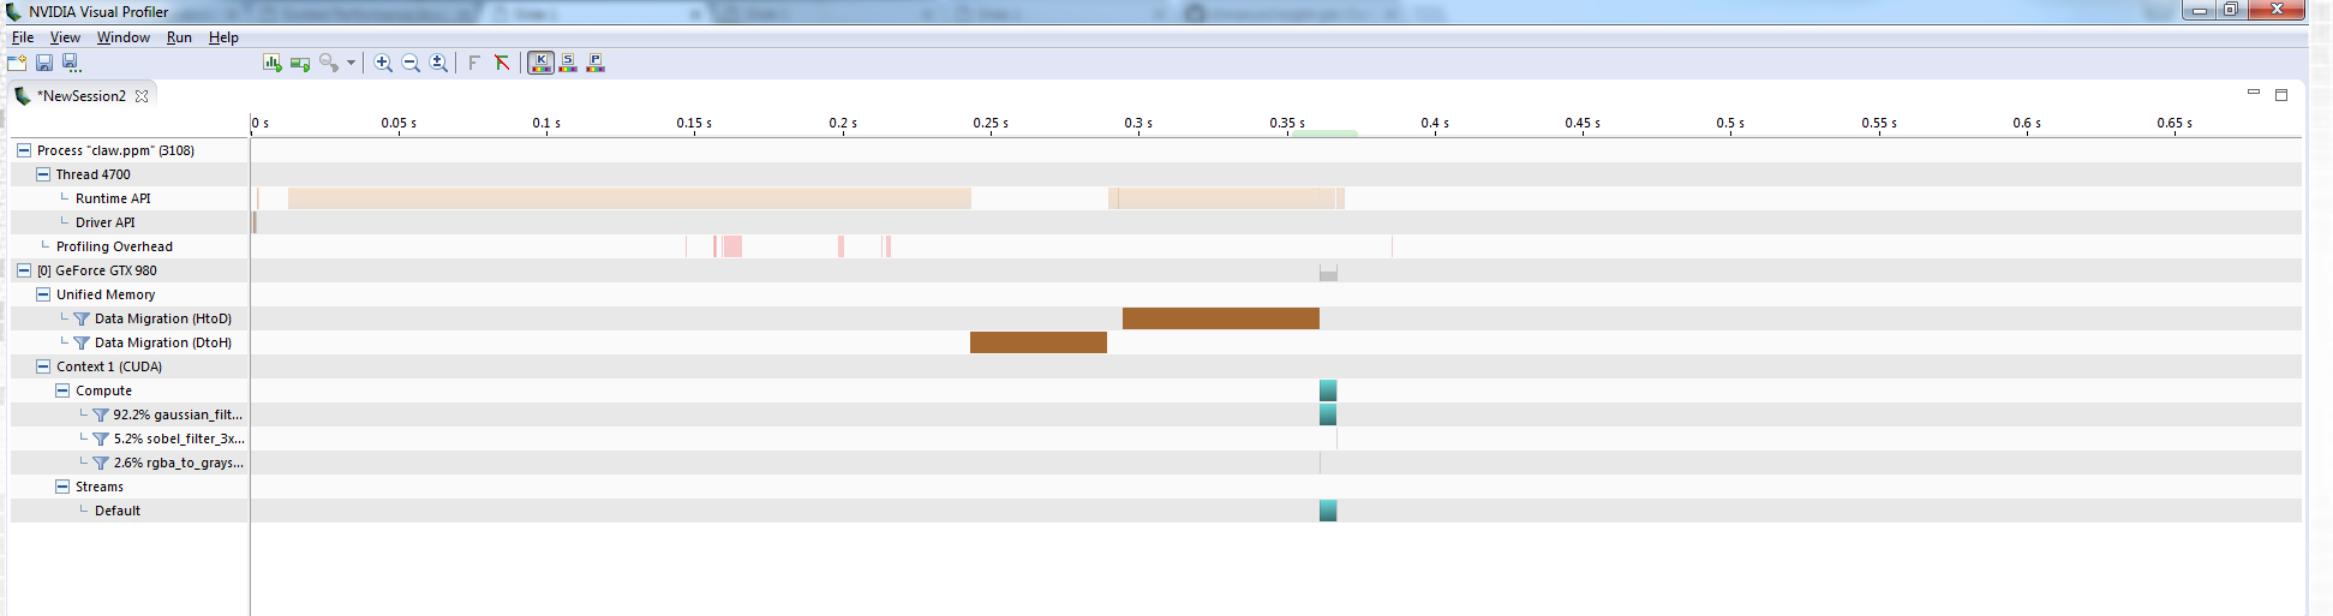
Compute Utilization
The device timeline shows an estimate of the amount of the total compute capacity being used by the kernels executing on the device.

Properties

Default

Duration	Session
385.676 n	

Guided Analysis



Analysis Details Console Settings

Export PDF Report

1. CUDA Application Analysis

2. Performance-Critical Kernels

The results on the right show your application's kernels ordered by potential for performance improvement. Starting with the kernels with the highest ranking, you should select an entry from the table and then perform kernel analysis to discover additional optimization opportunities.

Perform Kernel Analysis

Select a kernel from the table at right or from the timeline to enable kernel analysis. This analysis requires detailed profiling data, so your application will be run once to collect that data for the kernel if it is not already available.

Perform Additional Analysis

You can collect additional information to help identify kernels with potential performance problems. After running this analysis, select any of the new results at right to highlight the individual kernels for which the analysis applies.

Results

i Kernel Optimization Priorities

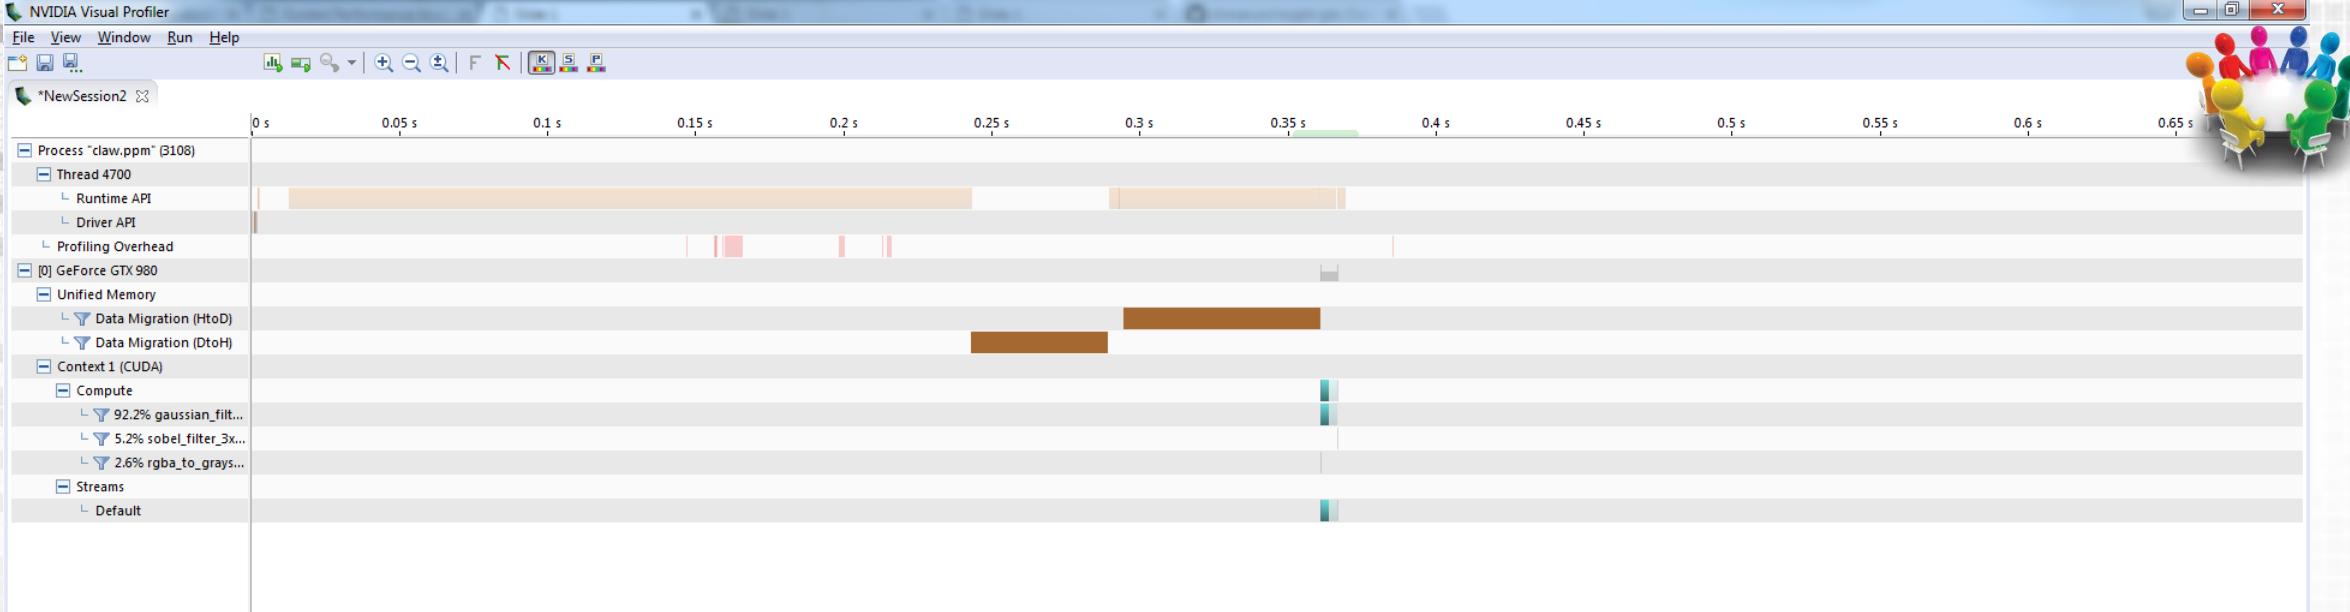
The following kernels are ordered by optimization importance based on execution time and achieved occupancy. Optimization of higher ranked kernels (those that appear first in the list) is more likely to improve performance compared to lower ranked kernels.

Rank	Description
100	[2 kernel instances] gaussian_filter_7x7_v0(int, int, unsigned char const *, unsigned char*)
5	[1 kernel instances] sobel_filter_3x3_v0(int, int, unsigned char const *, unsigned char*)
2	[1 kernel instances] rgba_to_grayscale_kernel_v0(int, int, uchar4 const *, unsigned char*)

guassian_filter_7x7_v0 kernel has highest rank

Properties

Select or highlight a single interval to see properties



Analysis Details Console Settings

Export PDF Report

Results

i Kernel Performance Is Bound By Memory Bandwidth

For device "GeForce GTX 980" the kernel's compute utilization is significantly lower than its memory utilization. These utilization levels indicate that the performance of the kernel is most likely being limited by the memory system. For this kernel the limiting factor in the memory system is the bandwidth of the L2 Cache memory.

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results at right indicate that the performance of kernel "gaussian_filter_7x7_v0" is most likely limited by memory bandwidth.

Perform Memory Bandwidth Analysis

The most likely bottleneck to performance for this kernel is memory bandwidth so you should first perform memory bandwidth analysis to determine how it is limiting performance.

Perform Compute Analysis

Perform Latency Analysis

Compute and instruction and memory latency are likely not the primary performance bottlenecks for this kernel, but you may still want to perform those analyses.

Rerun Analysis

If you modify the kernel you need to rerun your application to update this analysis.

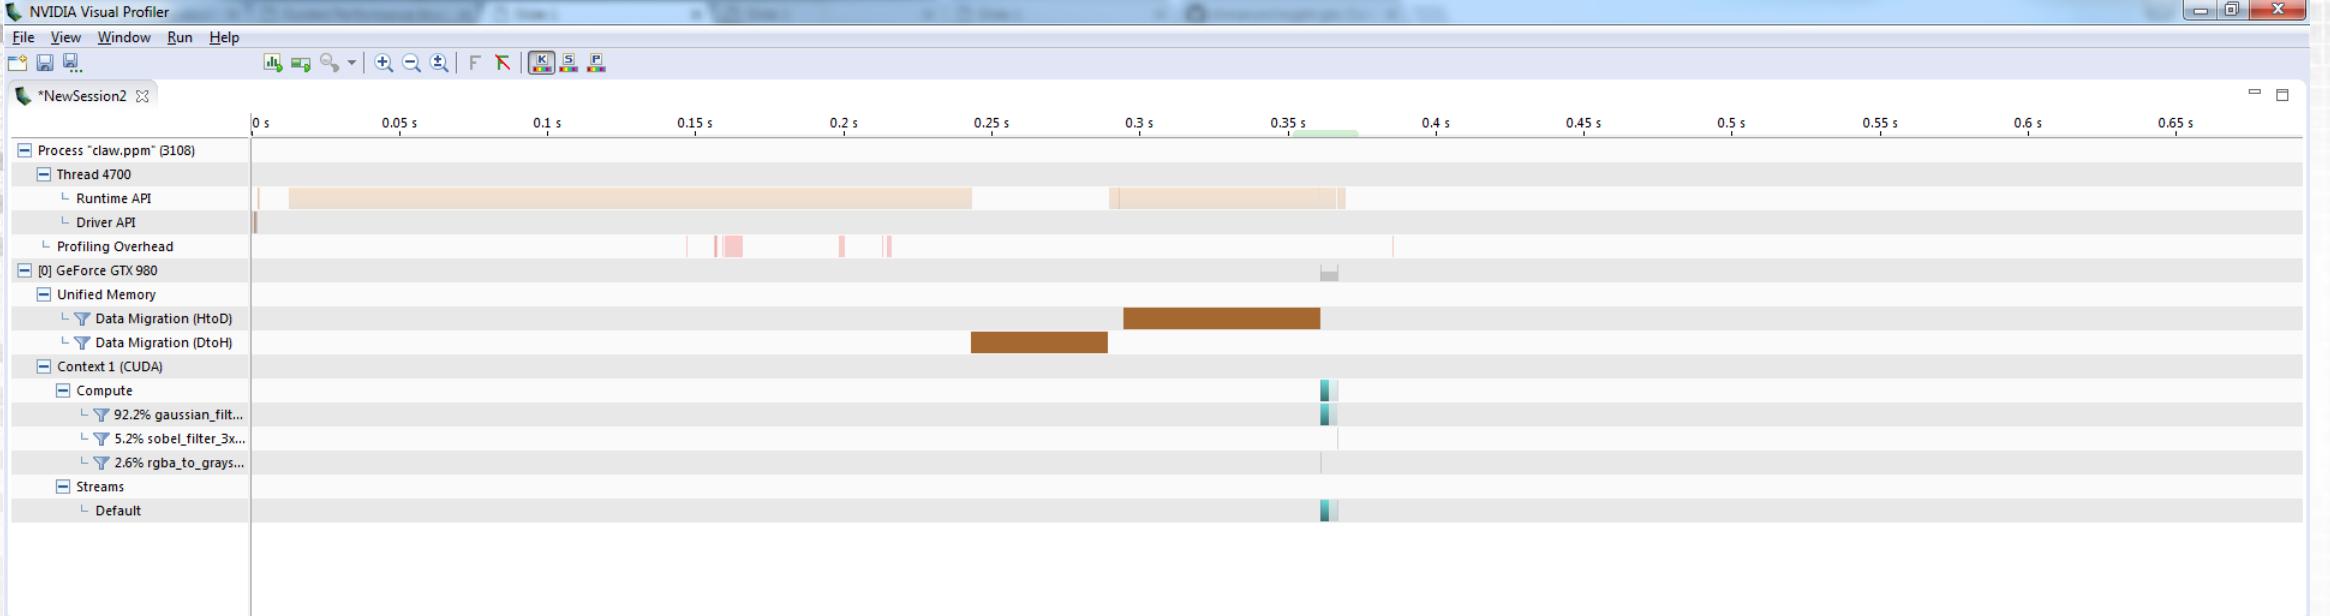
Utilization

Category	Utilization (%)
Compute	~25%
Memory (L2 Cache)	~75%

gaussian_filter_7x7_v0(int, int, unsigned char const *, unsigned char *)

Property	Value
Start	361.017 ms (361,017,459 ns)
End	363.759 ms (363,759,082 ns)
Duration	2.742 ms (2,741,623 ns)
Grid Size	[32,0,200,1]
Block Size	[8,8,1]
Registers/Thread	56
Shared Memory/Block	0 B
Occupancy	<ul style="list-style-type: none"> Achieved: 48.4% Theoretical: 56.2% Limiter: Registers
Shared Memory Configuration	<ul style="list-style-type: none"> Shared Memory Requested: 96 KiB Shared Memory Executed: 96 KiB Shared Memory Bank Size: 4 B
Global Cache Configuration	<ul style="list-style-type: none"> Global Cache Requested: off Global Cache Executed: off

What is this telling us about our code?



Analysis Details Console Settings

Export PDF Report

Results

i Kernel Performance Is Bound By Memory Bandwidth

For device "GeForce GTX 980" the kernel's compute utilization is significantly lower than its memory utilization. These utilization levels indicate that the performance of the kernel is most likely being limited by the memory system. For this kernel the limiting factor in the memory system is the bandwidth of the L2 Cache memory.

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results at right indicate that the performance of kernel "gaussian_filter_7x7_v0" is most likely limited by memory bandwidth.

Perform Memory Bandwidth Analysis

The most likely bottleneck to performance for this kernel is memory bandwidth so you should first perform memory bandwidth analysis to determine how it is limiting performance.

Perform Compute Analysis

Perform Latency Analysis

Compute and instruction and memory latency are likely not the primary performance bottlenecks for this kernel, but you may still want to perform those analyses.

Rerun Analysis

If you modify the kernel you need to rerun your application to update this analysis.

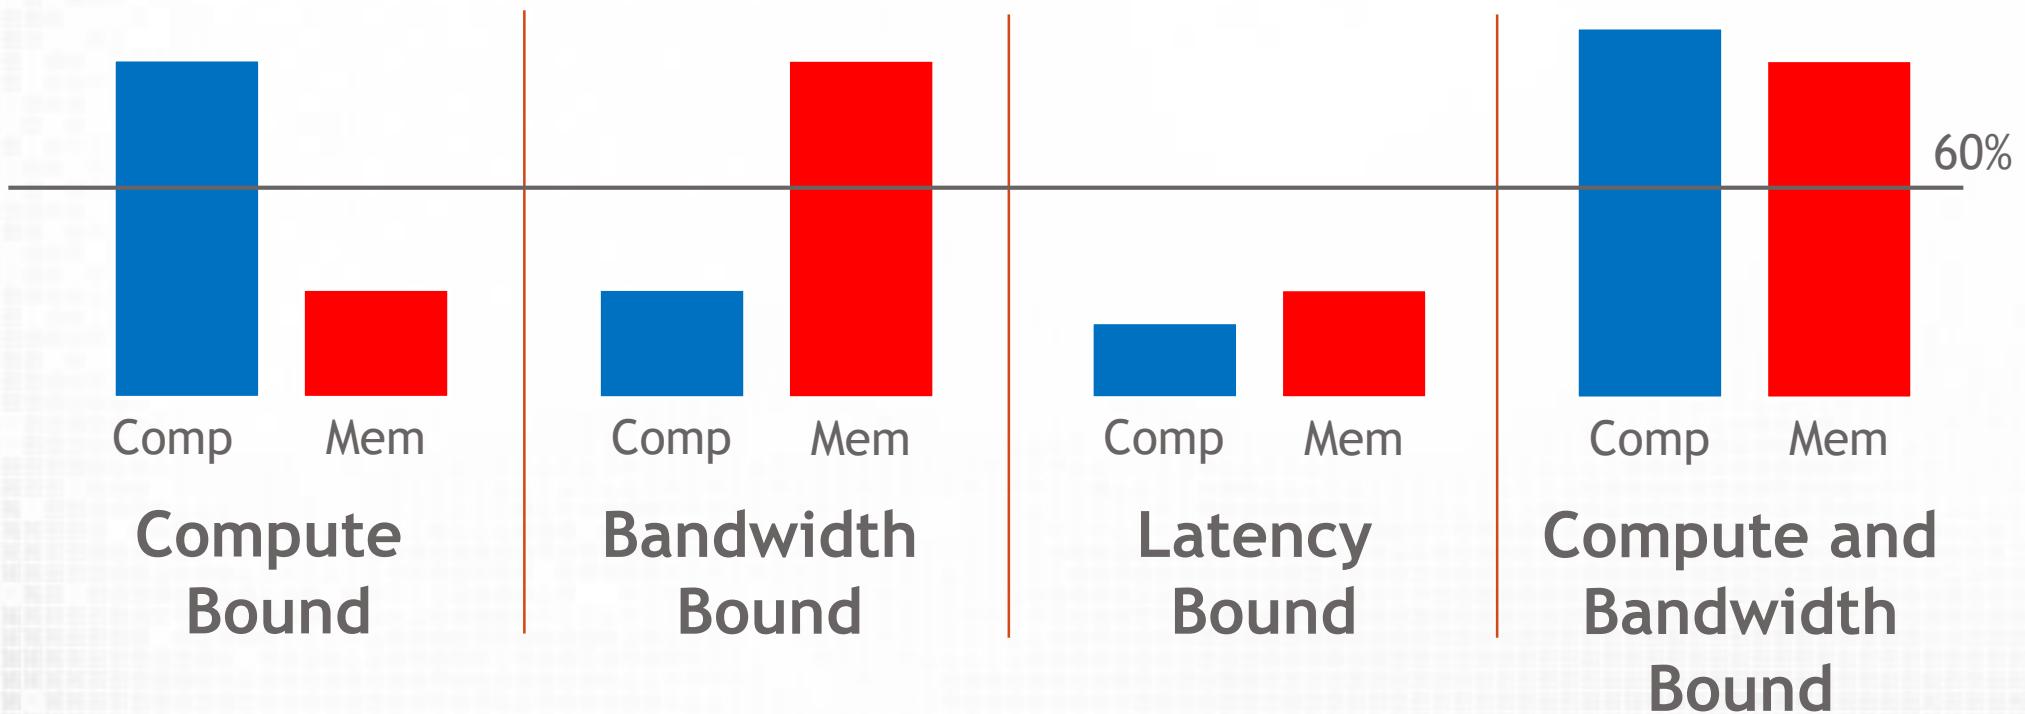
Memory Bound Problem!

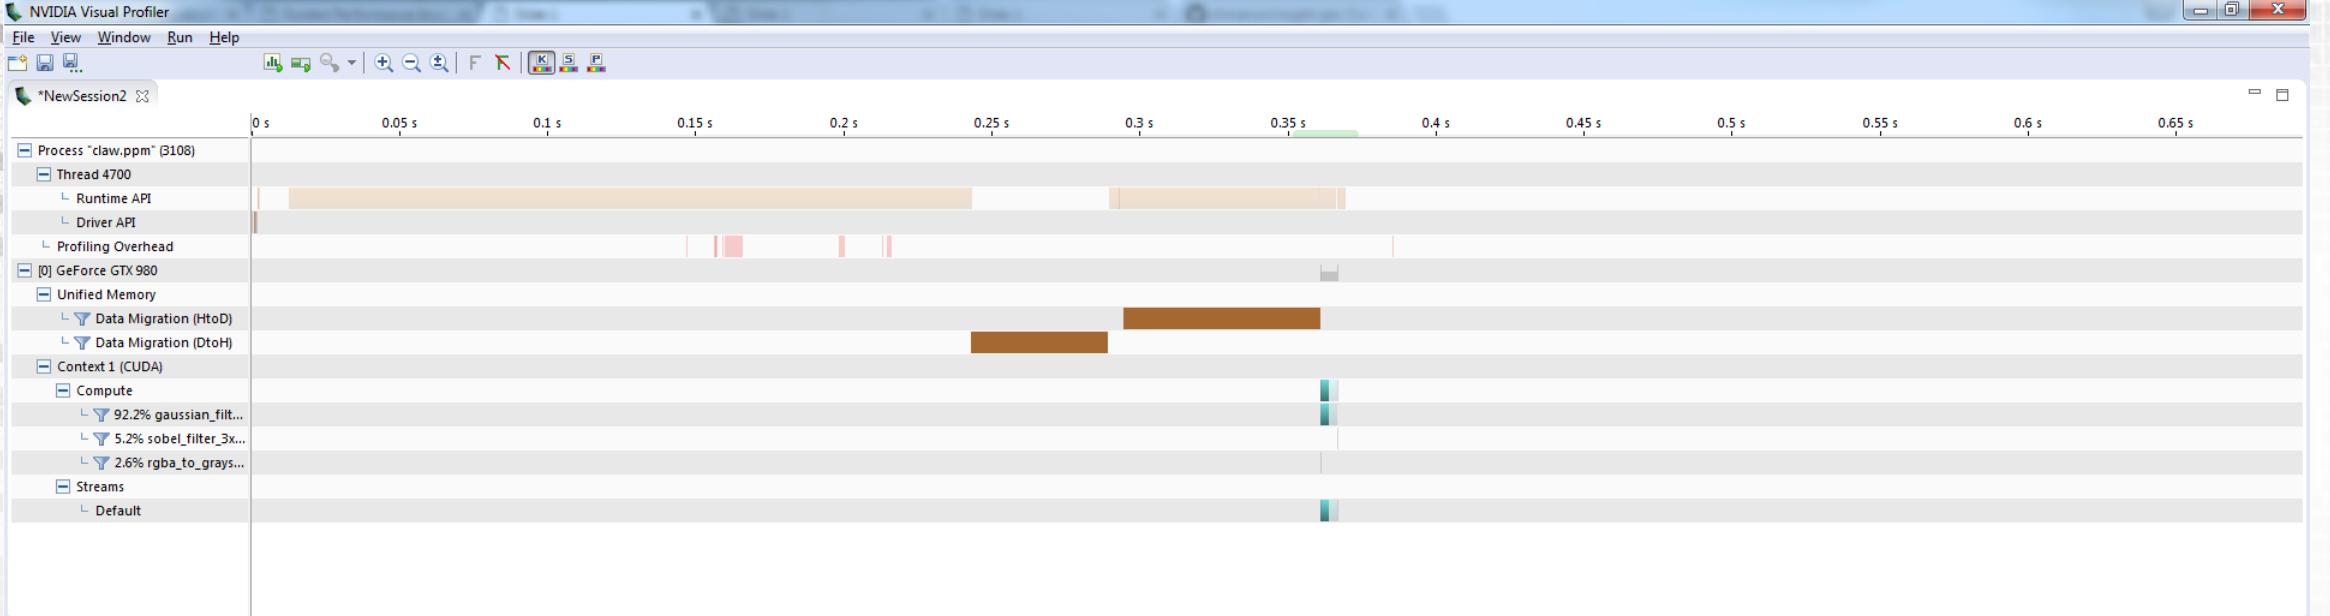
Properties

gaussian_filter_7x7_v0(int, int, unsigned char const *, unsigned char*)

Start	361.017 ms (361,017,459 ns)
End	363.759 ms (363,759,082 ns)
Duration	2.742 ms (2,741,623 ns)
Grid Size	[32,0,200,1]
Block Size	[8,8,1]
Registers/Thread	56
Shared Memory/Block	0 B
Occupancy	
Achieved	⚠ 48.4%
Theoretical	56.2%
Limiter	Registers
Shared Memory Configuration	
Shared Memory Requested	96 KiB
Shared Memory Executed	96 KiB
Shared Memory Bank Size	4 B
Global Cache Configuration	
Global Cache Requested	off
Global Cache Executed	off

Memory vs Compute vs Latency





Analysis Details Console Settings

Export PDF Report

Results

i Kernel Performance Is Bound By Memory Bandwidth

For device "GeForce GTX 980" the kernel's compute utilization is significantly lower than its memory utilization. These utilization levels indicate that the performance of the kernel is most likely being limited by the memory system. For this kernel the limiting factor in the memory system is the bandwidth of the L2 Cache memory.

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results at right indicate that the performance of kernel "gaussian_filter_7x7_v0" is most likely limited by memory bandwidth.

Perform Memory Bandwidth Analysis

The most likely bottleneck to performance for this kernel is memory bandwidth so you should first perform memory bandwidth analysis to determine how it is limiting performance.

Perform Compute Analysis

Perform Latency Analysis

Compute and instruction and memory latency are likely not the primary performance bottlenecks for this kernel, but you may still want to perform those analyses.

Rerun Analysis

If you modify the kernel you need to rerun your application to update this analysis.

Category	Utilization (%)
Compute	~25%
Memory (L2 Cache)	~75%

Properties

gaussian_filter_7x7_v0(int, int, unsigned char const *, unsigned char*)

Start	361.017 ms (361,017,459 ns)
End	363.759 ms (363,759,082 ns)
Duration	2.742 ms (2,741,623 ns)
Grid Size	[32,0,200,1]
Block Size	[8,8,1]
Registers/Thread	56
Shared Memory/Block	0 B
Occupancy	
Achieved	⚠ 48.4%
Theoretical	56.2%
Limiter	Registers
Shared Memory Configuration	
Shared Memory Requested	96 kB
Shared Memory Executed	96 kB
Shared Memory Bank Size	4 B
Global Cache Configuration	
Global Cache Requested	off
Global Cache Executed	off

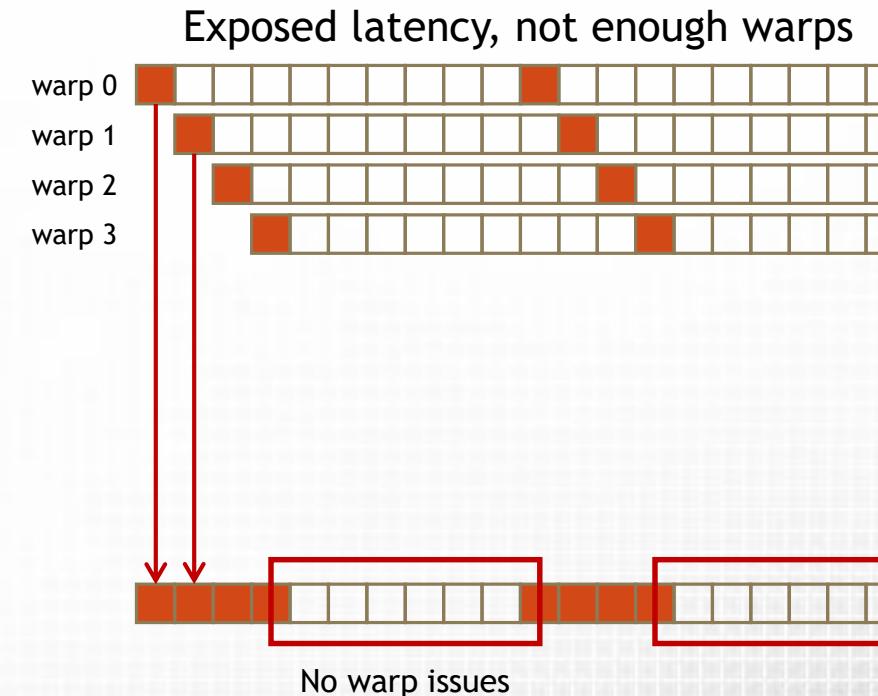
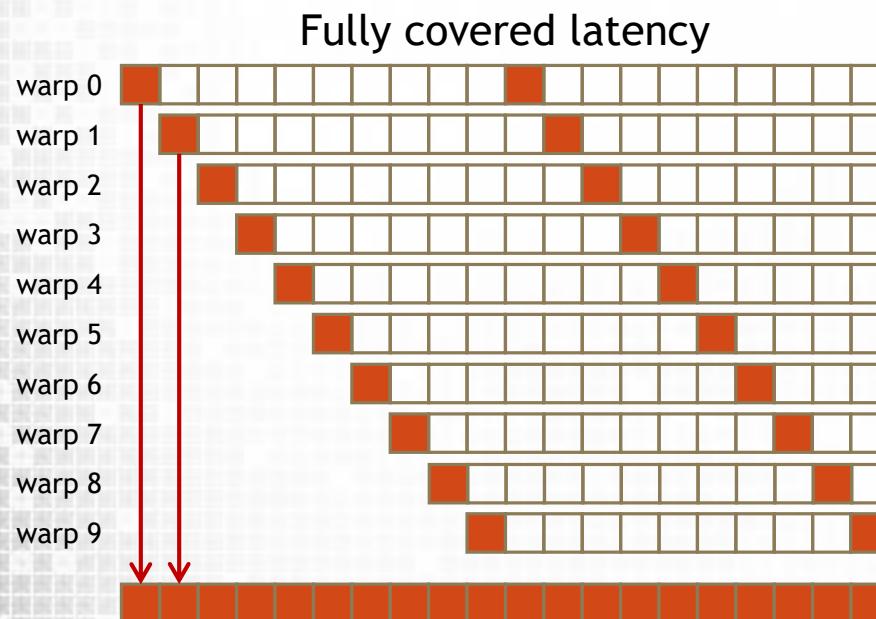
Better Occupancy might improve compute use

What about occupancy?

- ❑ Occupancy: “*number of active warps over max warps supported*”
- ❑ Increasing achieved occupancy can hide latency
 - ❑ More warps available for execution = more to hide latency

■ The warp issues

□ The warp waits (latency)



The University
Of
Sheffield.



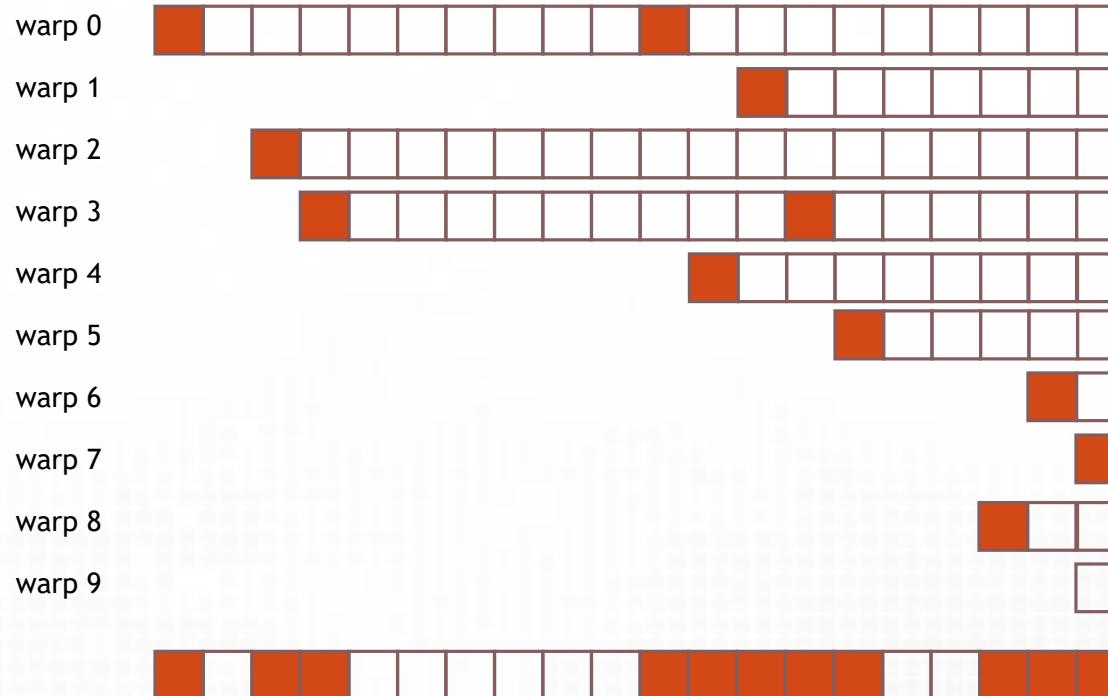
NVIDIA
GPU
RESEARCH
CENTER



Occupancy

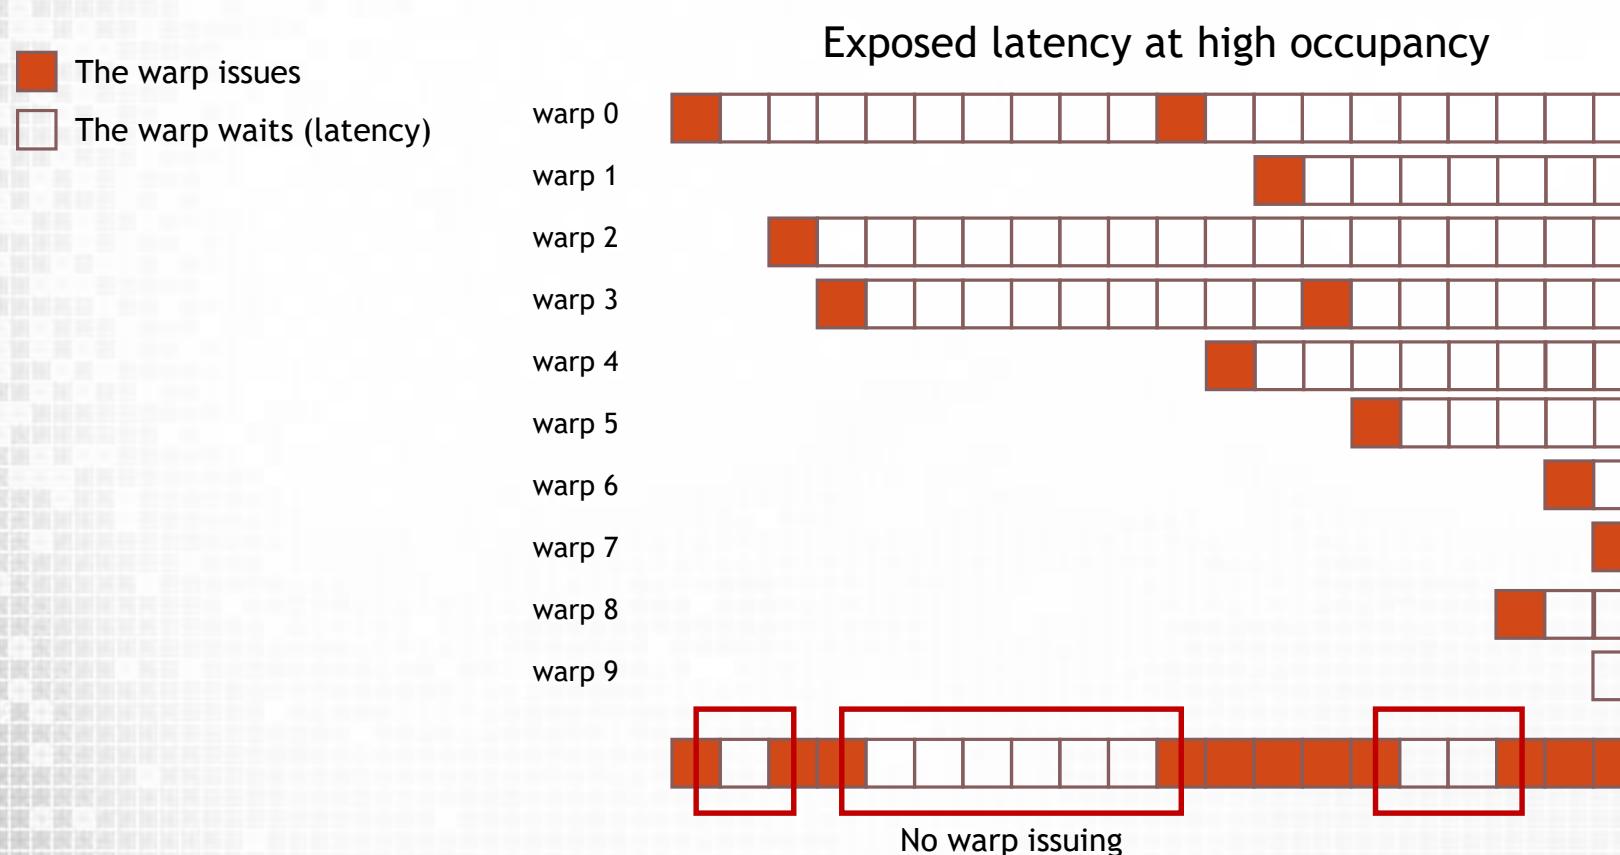
- ❑ In our case we are not achieving theoretical occupancy (we have latency)
What is the problem here?

█ The warp issues
█ The warp waits (latency)

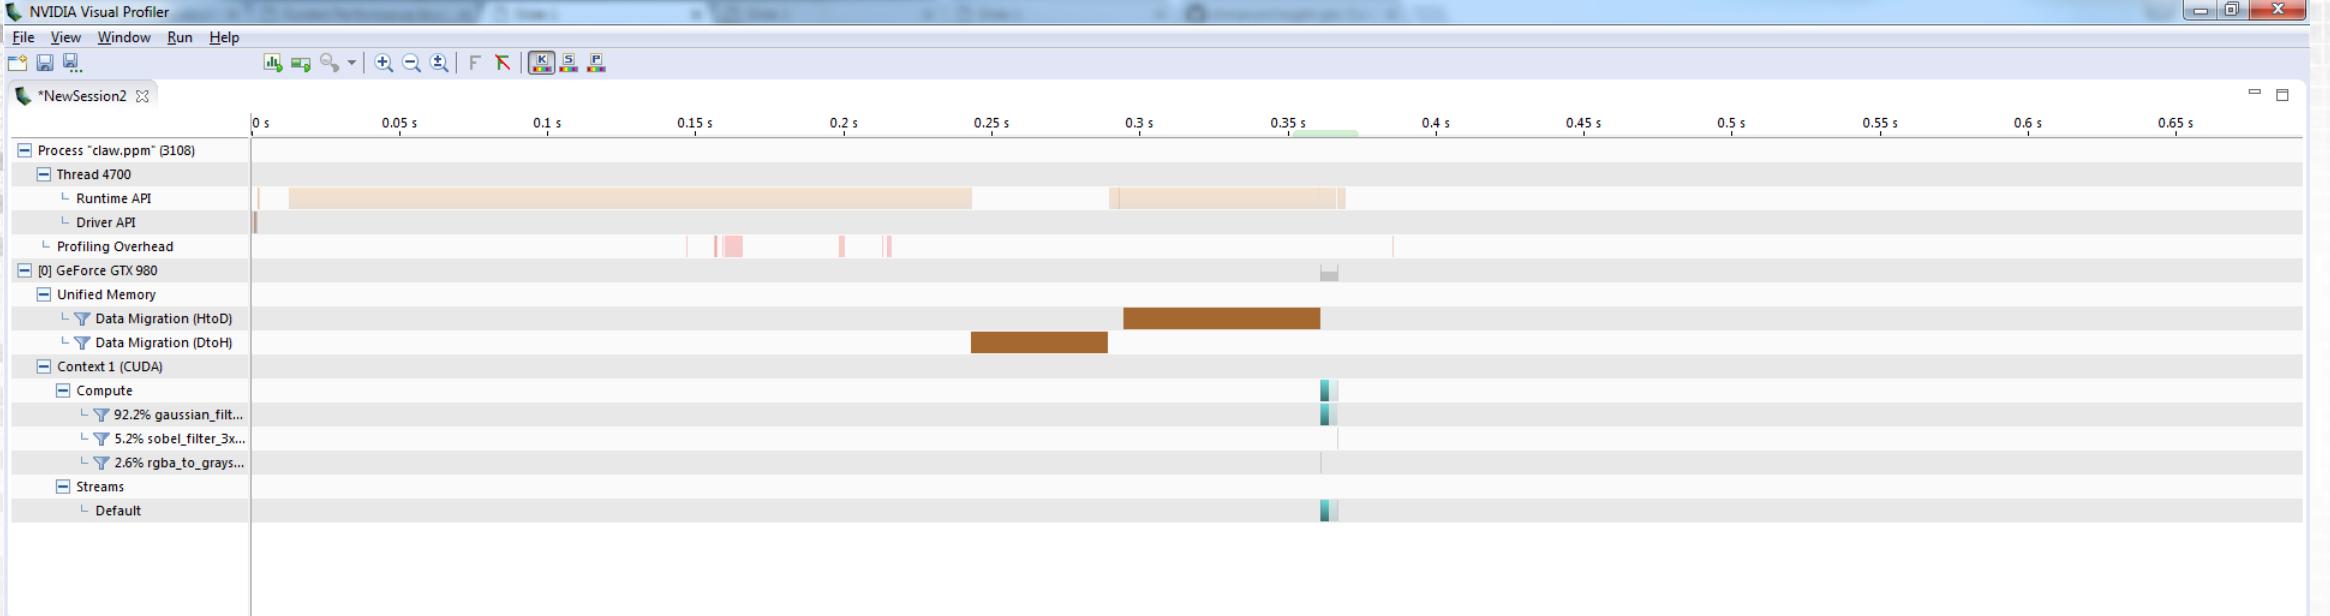


Occupancy

- ❑ In our case we have good occupancy but still high latency
 - ❑ Schedulers cant find eligible warps at every cycle



Warps are waiting for memory (transactions)



Analysis **Details** **Console** **Settings**

Export PDF Report

Results

i Kernel Performance Is Bound By Memory Bandwidth

For device "GeForce GTX 980" the kernel's compute utilization is significantly lower than its memory utilization. These utilization levels indicate that the performance of the kernel is most likely being limited by the memory system. For this kernel the limiting factor in the memory system is the bandwidth of the L2 Cache memory.

More information

Category	Utilization (%)
Compute	~25% (Memory operations) + ~75% (Arithmetic operations)
Memory (L2 Cache)	~100% (Memory operations)

1. CUDA Application Analysis

2. Performance-Critical Kernels

3. Compute, Bandwidth, or Latency Bound

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results at right indicate that the performance of kernel "gaussian_filter_7x7_v0" is most likely limited by memory bandwidth.

Perform Memory Bandwidth Analysis

The most likely bottleneck to performance for this kernel is memory bandwidth so you should first perform memory bandwidth analysis to determine how it is limiting performance.

Perform Compute Analysis

Perform Latency Analysis

Compute and instruction and memory latency are likely not the primary performance bottlenecks for this kernel, but you may still want to perform those analyses.

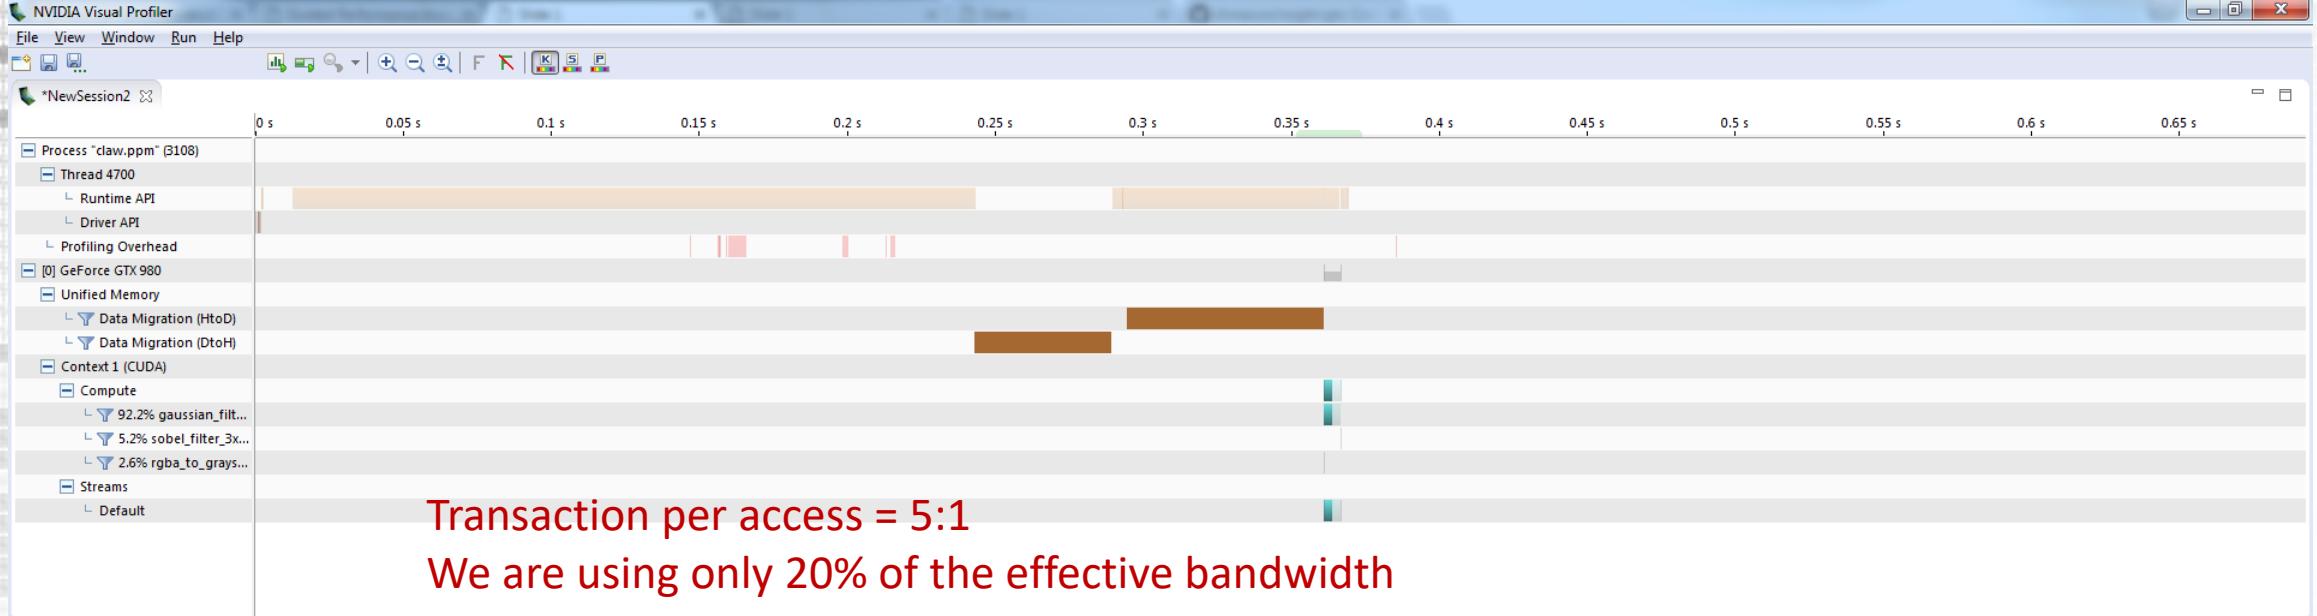
Rerun Analysis

If you modify the kernel you need to rerun your application to update this analysis.

Properties

gaussian_filter_7x7_v0(int, int, unsigned char const *, unsigned char*)

Start	361.017 ms (361,017,459 ns)
End	363.759 ms (363,759,082 ns)
Duration	2.742 ms (2,741,623 ns)
Grid Size	[32,0,200,1]
Block Size	[8,8,1]
Registers/Thread	56
Shared Memory/Block	0 B
▲ Occupancy	
Achieved	⚠ 48.4%
Theoretical	56.2%
Limiter	Registers
▲ Shared Memory Configuration	
Shared Memory Requested	96 KiB
Shared Memory Executed	96 KiB
Shared Memory Bank Size	4 B
▲ Global Cache Configuration	
Global Cache Requested	off
Global Cache Executed	off



Transaction per access = 5:1
We are using only 20% of the effective bandwidth

The screenshot shows the Nsight Compute interface with the following sections:

- Analysis**: The main menu bar.
- Details**: A panel showing the properties of the current kernel, `gaussian_filter_7x7_v0(int, int, unsigned char const *, unsigned char *)`. Properties include:
 - Start: 361.017 ms (361,017,459 ns)
 - End: 363.759 ms (363,759,082 ns)
 - Duration: 2.742 ms (2,741,623 ns)
 - Grid Size: [30,200,1]
 - Block Size: [8,8,1]
 - Registers/Thread: 56
 - Shared Memory/Block: 0 B
- Properties**: A table showing kernel properties for `gaussian_filter_7x7_v0(int, int, unsigned char const *, unsigned char *)`.
- Results**: The main analysis pane.
 - Global Memory Alignment and Access Pattern**: A warning message stating that memory bandwidth is used most efficiently when each global memory load and store has proper alignment and access pattern. It includes an optimization note: "Select each entry below to open the source code to a global load or store within the kernel with an inefficient alignment or access pattern. For each load or store improve the alignment and access pattern of the memory access." and a "More..." link.
 - A list of 245 global load transactions, each with an access count of 5, ideal transaction count of 1, and a range of L2 transactions for 128000 total executions.
- Actions**: A button labeled "Rerun Analysis".
- Message**: A message at the bottom left: "If you modify the kernel you need to rerun your application to update this analysis."

GPU Utilization Is Limited By Memory Bandwidth

The following table shows the memory bandwidth used by this kernel for the various types of memory on the device. The table also shows the utilization of each memory type relative to the maximum throughput supported by the memory. The results show that the kernel's performance is potentially limited by the bandwidth available from one or more of the memories on the device.

Optimization: Try the following optimizations for the memory with high bandwidth utilization.

Shared Memory - If possible use 64-bit accesses to shared memory and 8-byte bank mode to achieve 2x throughput.

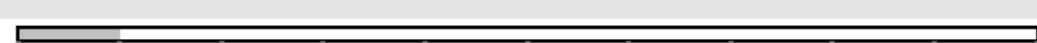
L2 Cache - Align and block kernel data to maximize L2 cache efficiency.

Unified Cache - Reallocate texture data to shared or global memory. Resolve alignment and access pattern issues for global loads and stores.

Device Memory - Resolve alignment and access pattern issues for global loads and stores.

System Memory (via PCIe) - Make sure performance critical data is placed in device or shared memory.

[More...](#)

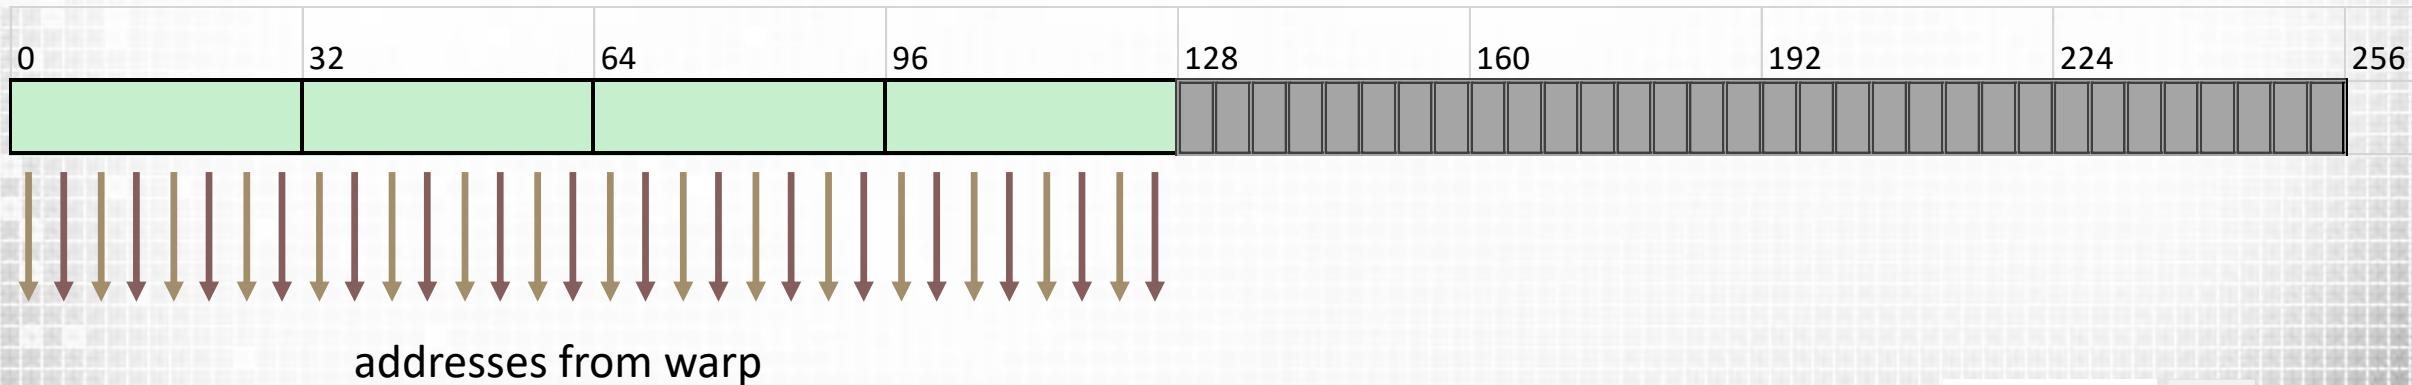
	Transactions	Bandwidth	Utilization
Shared Memory			
Shared Loads	0	0 B/s	
Shared Stores	0	0 B/s	
Shared Total	0	0 B/s	 Idle Low Medium High Max
L2 Cache			
Reads	30460192	392.435 GB/s	
Writes	1024006	13.193 GB/s	
Total	31484198	405.627 GB/s	 Idle Low Medium High Max
Unified Cache			
Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Global Loads	55425352	391.198 GB/s	
Global Stores	1024000	13.193 GB/s	
Texture Reads	25070080	322.991 GB/s	
Unified Total	81519432	727.382 GB/s	 Idle Low Medium High Max
Device Memory			
Reads	212615	2.739 GB/s	
Writes	127951	1.648 GB/s	
Total	340566	4.388 GB/s	 Idle Low Medium High Max
System Memory [PCIe configuration: Gen2 x16, 5 Gbit/s]			
Reads	0	0 B/s	 Idle Low Medium High Max
Writes	5	64.417 kB/s	 Idle Low Medium High Max

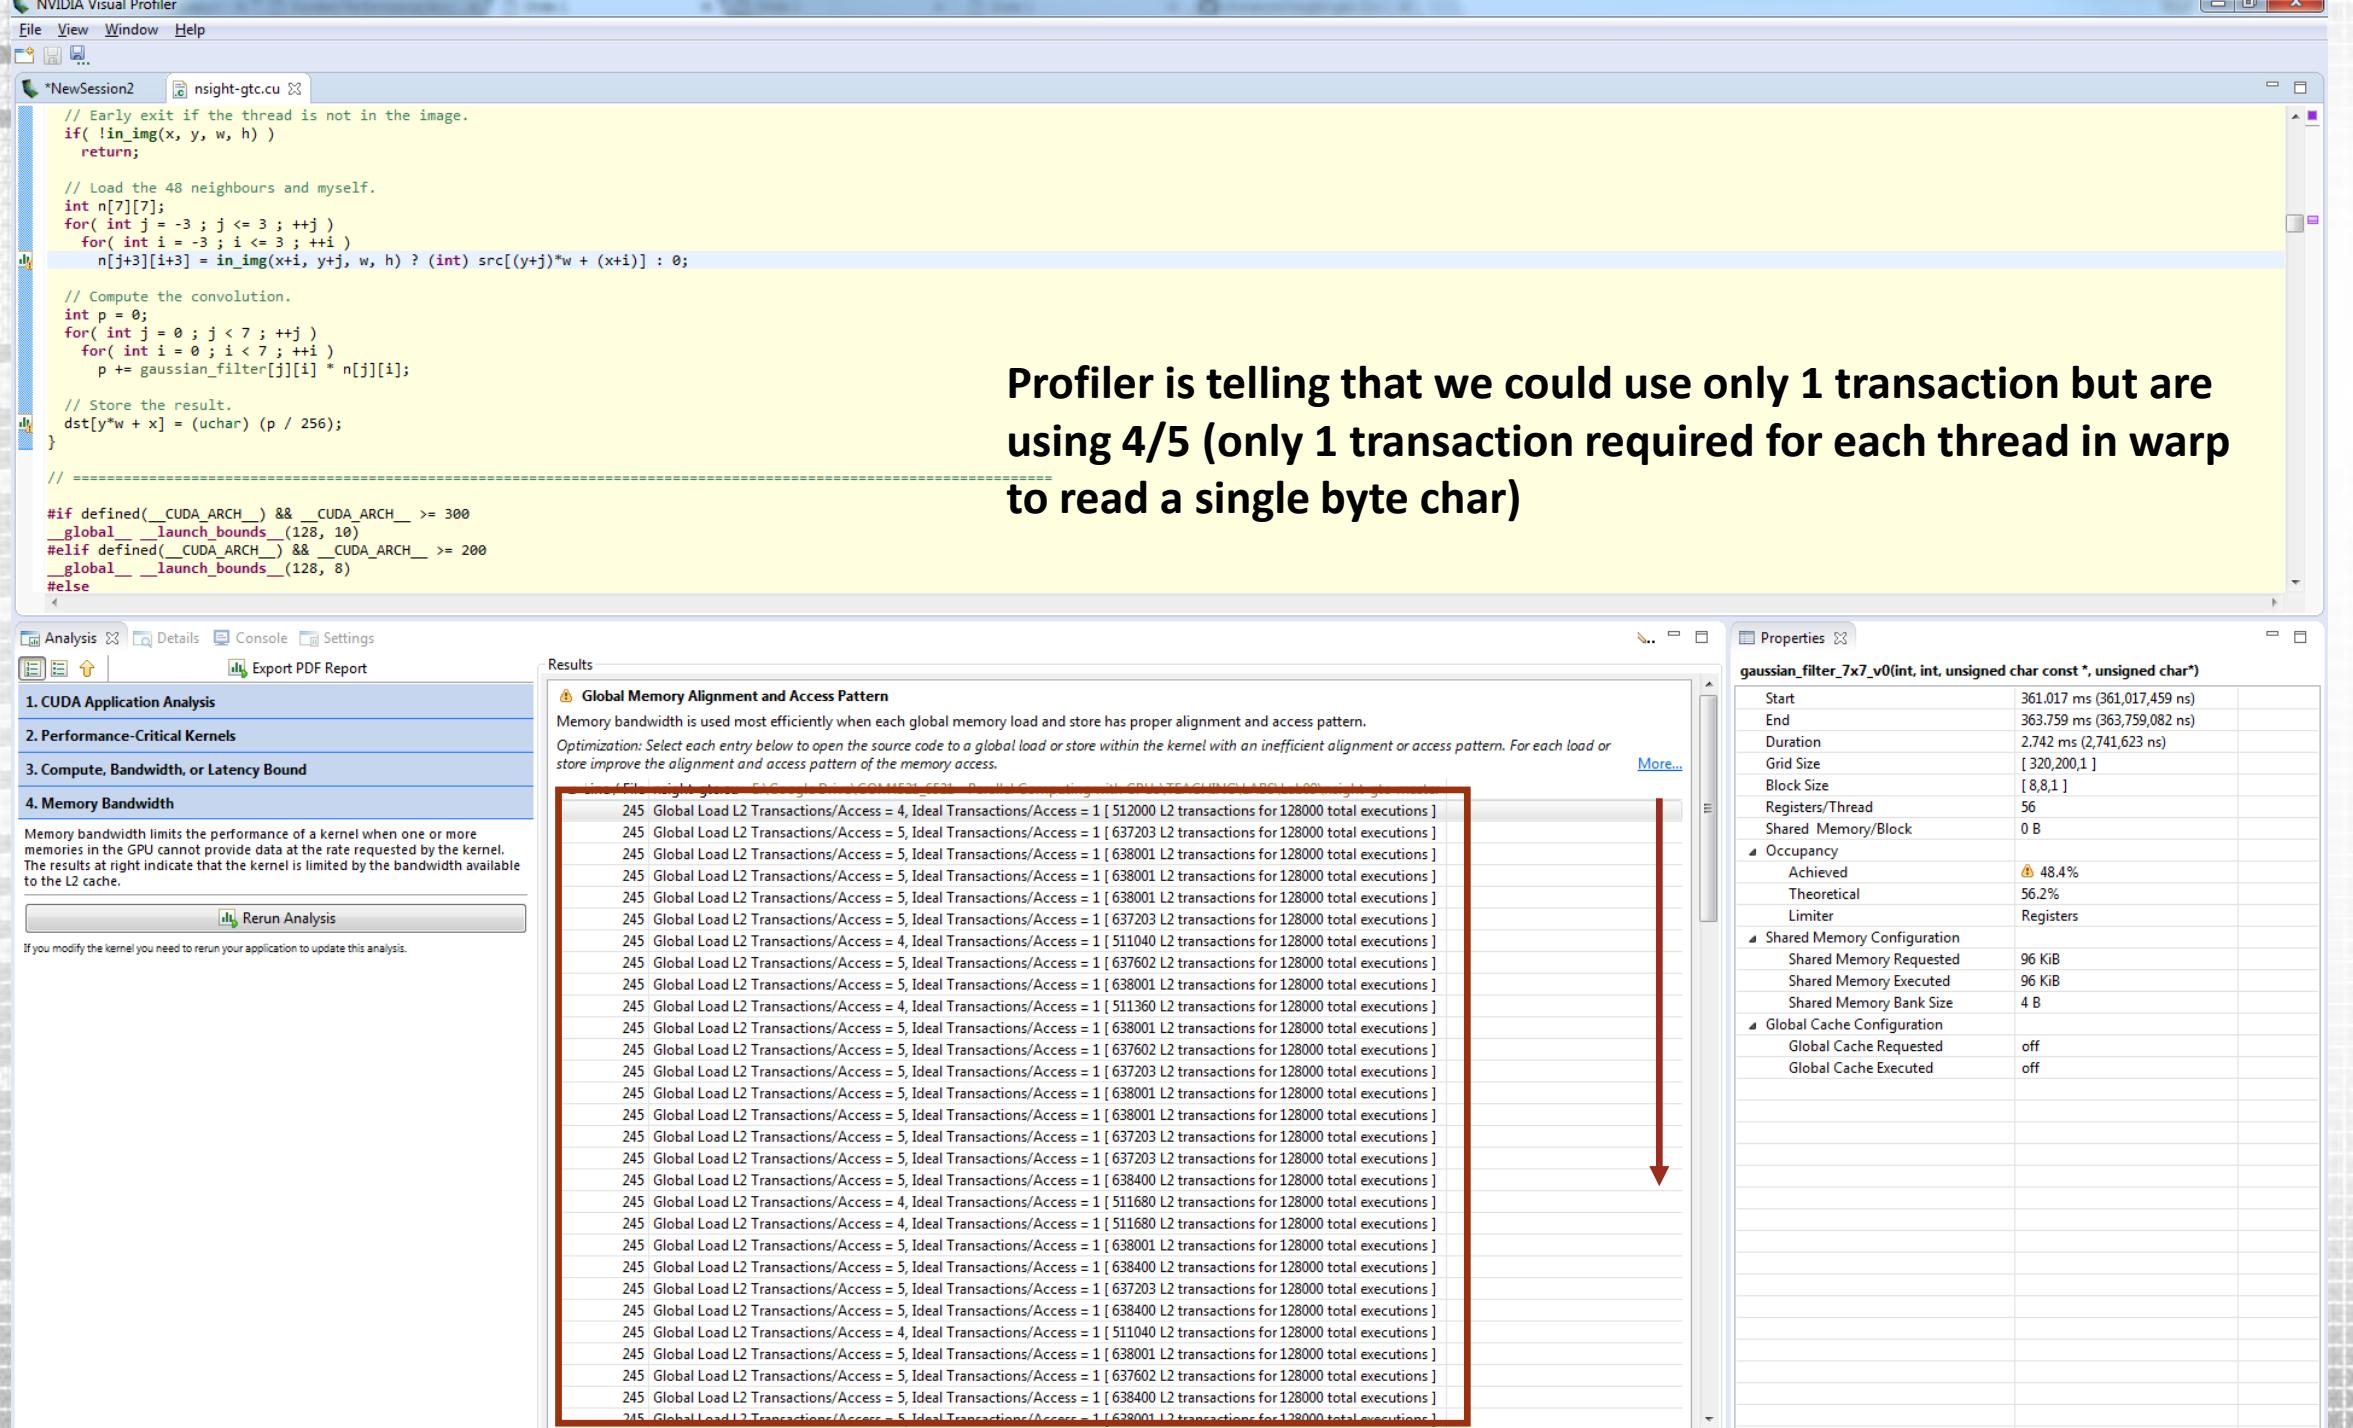


Transactions per access?

- ❑ Think back to Lecture 11
 - ❑ To get 100% efficiency **our threads need to access consecutive 4 byte values**
 - ❑ 32 Threads in warp accessing 4B each
 - ❑ 128B total via 4 L2 cache lines

```
__global__ void copy(float *odata, float* idata)
    int xid = blockIdx.x * blockDim.x + threadIdx.x;
    odata[xid] = idata[xid];
}
```





NVIDIA Visual Profiler

File View Window Help

*NewSession2 nsight-gtc.cu

```
// Early exit if the thread is not in the image.
if( !in_img(x, y, w, h) )
    return;

// Load the 48 neighbours and myself.
int n[7][7];
for( int j = -3 ; j <= 3 ; ++j )
    for( int i = -3 ; i <= 3 ; ++i )
        n[j+3][i+3] = in_img(x+i, y+j, w, h) ? (int) src[(y+j)*w + (x+i)] : 0;

// Compute the convolution.
int p = 0;
for( int j = 0 ; j < 7 ; ++j )
    for( int i = 0 ; i < 7 ; ++i )
        p += gaussian_filter[j][i] * n[j][i];

// Store the result.
dst[y*w + x] = (uchar) (p / 256);

// =====

#if defined(__CUDA_ARCH__) && __CUDA_ARCH__ >= 300
__global__ __launch_bounds__(128, 10)
#elif defined(__CUDA_ARCH__) && __CUDA_ARCH__ >= 200
__global__ __launch_bounds__(128, 8)
#else

```

Memory is indexed based on $x==\text{threadIdx.x}$: Suggests access is coalesced.
Cause not clear.....

Analysis Details Console Settings

Export PDF Report

Results

1. CUDA Application Analysis

2. Performance-Critical Kernels

3. Compute, Bandwidth, or Latency Bound

4. Memory Bandwidth

Memory bandwidth limits the performance of a kernel when one or more memories in the GPU cannot provide data at the rate requested by the kernel. The results at right indicate that the kernel is limited by the bandwidth available to the L2 cache.

Rerun Analysis

If you modify the kernel you need to rerun your application to update this analysis.

Line / File nsight-gtc.cu - E:\Google Drive\COM4521_6521 - Parallel Computing with GPUs\TEACHING\LABS\Lab09\nsight-gtc-master

Global Load L2 Transactions/Access	Ideal Transactions/Access	Start	End	Duration
4	= 1 [512000 L2 transactions for 128000 total executions]	361,017 ms	(361,017,459 ns)	361,017 ms (361,017,459 ns)
5	= 1 [637203 L2 transactions for 128000 total executions]	363,759 ms	(363,759,082 ns)	363,759 ms (363,759,082 ns)
5	= 1 [638001 L2 transactions for 128000 total executions]	2,742 ms	(2,741,623 ns)	2,742 ms (2,741,623 ns)
5	= 1 [638001 L2 transactions for 128000 total executions]	[32,200,1]		[32,200,1]
5	= 1 [638001 L2 transactions for 128000 total executions]	56		56
5	= 1 [638001 L2 transactions for 128000 total executions]	0 B		0 B
5	= 1 [638001 L2 transactions for 128000 total executions]	Registers		Registers
5	= 1 [638001 L2 transactions for 128000 total executions]	Achieved	48.4%	Achieved
5	= 1 [638001 L2 transactions for 128000 total executions]	Theoretical	56.2%	Theoretical
5	= 1 [638001 L2 transactions for 128000 total executions]	Limiter	Registers	Limiter
5	= 1 [638001 L2 transactions for 128000 total executions]	Shared Memory Requested	96 KiB	Shared Memory Configuration
5	= 1 [638001 L2 transactions for 128000 total executions]	Shared Memory Executed	96 KiB	Shared Memory Executed
5	= 1 [638001 L2 transactions for 128000 total executions]	Shared Memory Bank Size	4 B	Shared Memory Bank Size
5	= 1 [638001 L2 transactions for 128000 total executions]	Global Cache Requested	off	Global Cache Configuration
5	= 1 [638001 L2 transactions for 128000 total executions]	Global Cache Executed	off	Global Cache Executed

Properties

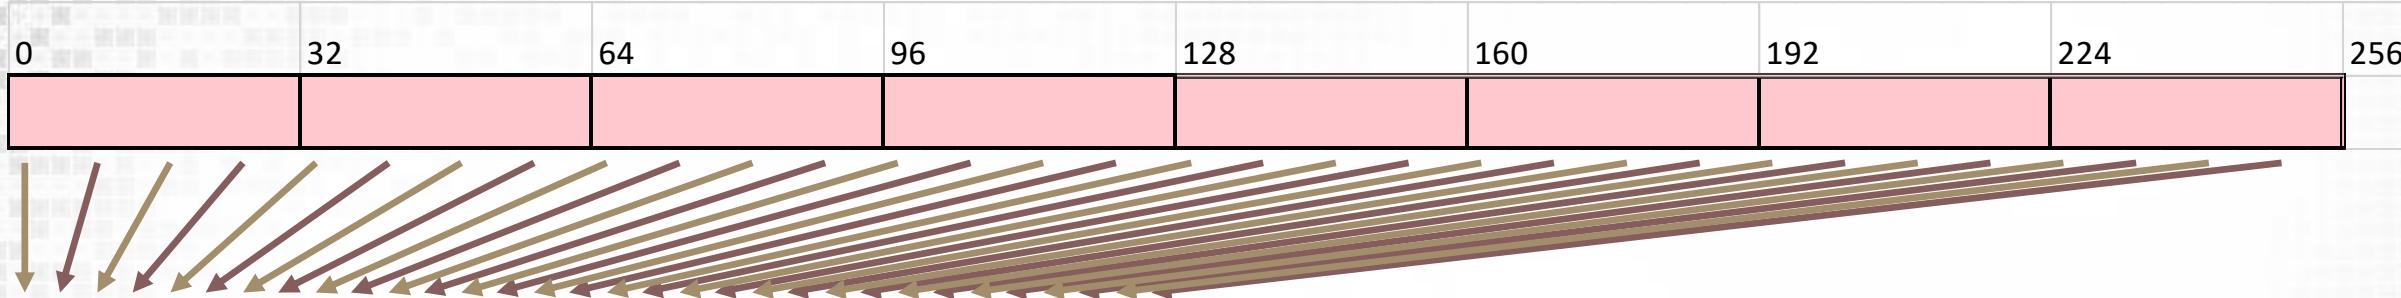
gaussian_filter_7x7_v0(int, int, unsigned char const *, unsigned char *)

Start	361,017 ms (361,017,459 ns)
End	363,759 ms (363,759,082 ns)
Duration	2,742 ms (2,741,623 ns)
Grid Size	[32,200,1]
Block Size	[8,8,1]
Registers/Thread	56
Shared Memory/Block	0 B
Occupancy	
Achieved	48.4%
Theoretical	56.2%
Limiter	Registers
Shared Memory Configuration	
Shared Memory Requested	96 KiB
Shared Memory Executed	96 KiB
Shared Memory Bank Size	4 B
Global Cache Configuration	
Global Cache Requested	off
Global Cache Executed	off

Analysis

- ❑ The limiting factor of our code is L2 Throughput
 - ❑ There is nothing wrong with having high throughput
 - ❑ Except: There is not enough compute to hide this
 - ❑ We can't increase occupancy any further to hide this
- ❑ Solution: We need to reduce the time it takes to get data to the device to do compute on it. Either by
 - ❑ Moving data closer to the SMPs
 - ❑ **Making our L2 reads/writes more efficient**
 - ❑ Currently ~4-5 Transactions/Access
 - ❑ Our L2 cache lines are being used ineffectively

Causes of Transaction per access: Striding?

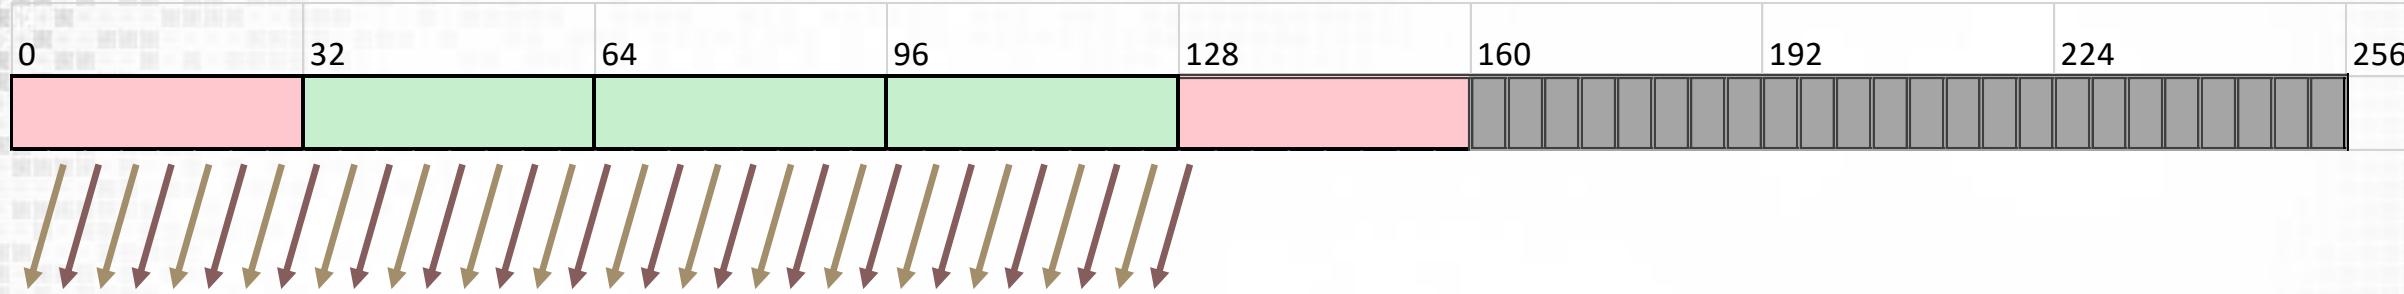


```
__global__ void copy(float *odata, float* idata)
    int xid = (blockIdx.x * blockDim.x + threadIdx.x)* 2;
    odata[xid] = idata[xid];
}
```

❑ Lecture 11 example

- ❑ Strides (like above) cause poor transactions per access
- ❑ In the above case 8 transactions where we could have used 4

Causes of Transaction per access: Offset?



```
__global__ void copy(float *odata, float* idata)
    int xid = blockIdx.x * blockDim.x + threadIdx.x + 1;
    odata[xid] = idata[xid];
}
```

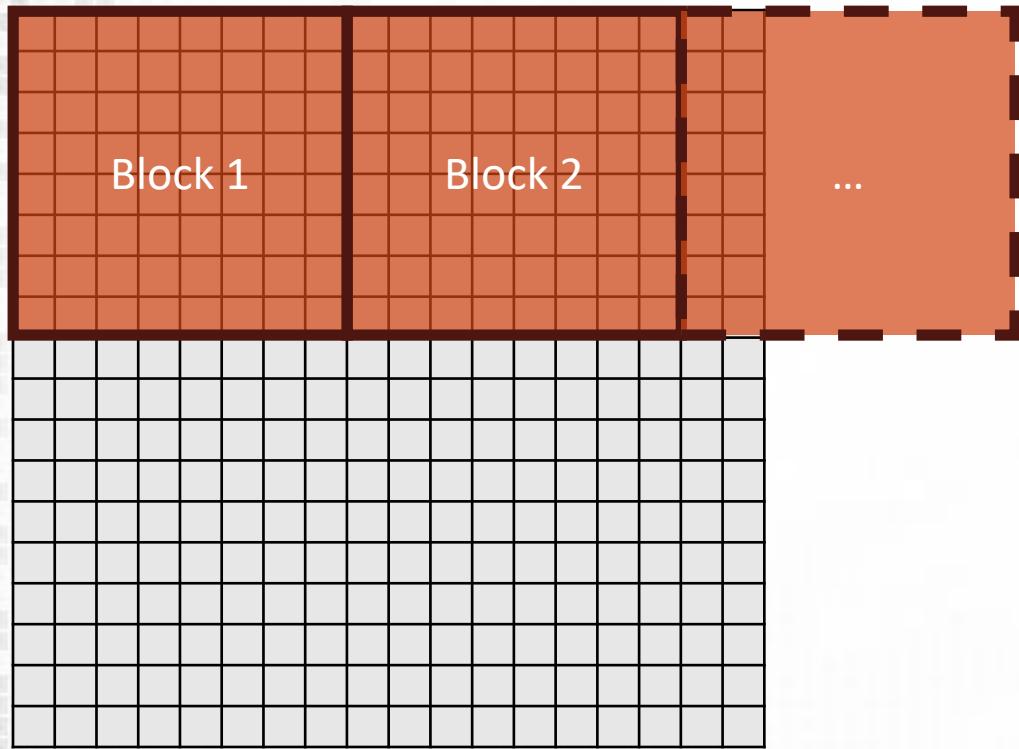
❑ Lecture 11 Example:

- ❑ If memory accesses are offset then parts of the cache line will be unused (shown in red) e.g.
- ❑ Use thread blocks sizes of multiples of 32!



What is our current data layout?

Blocks are 8x8



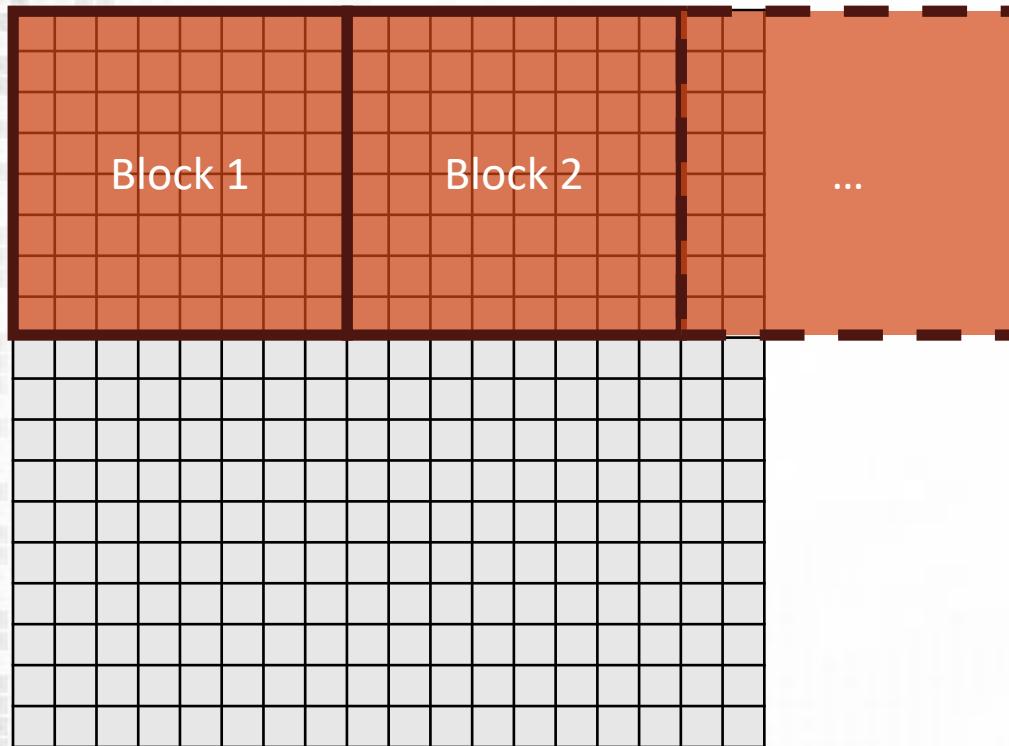
Warp are 8x4



Why might this be a problem

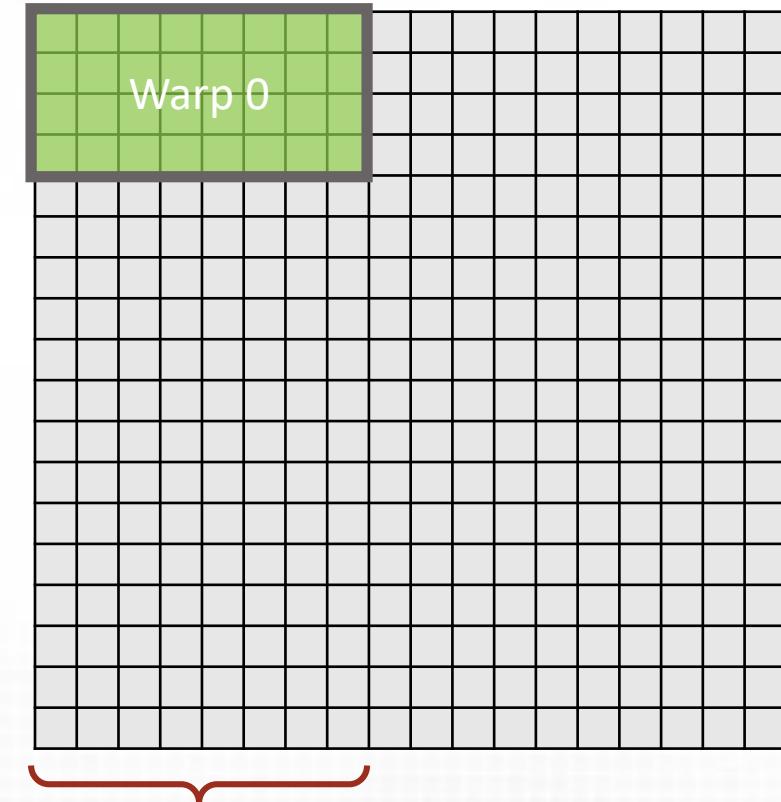
What is our current data layout?

Blocks are 8x8



Warp 0

Warp s are 8x4



threadIdx.x not consecutive within the warp

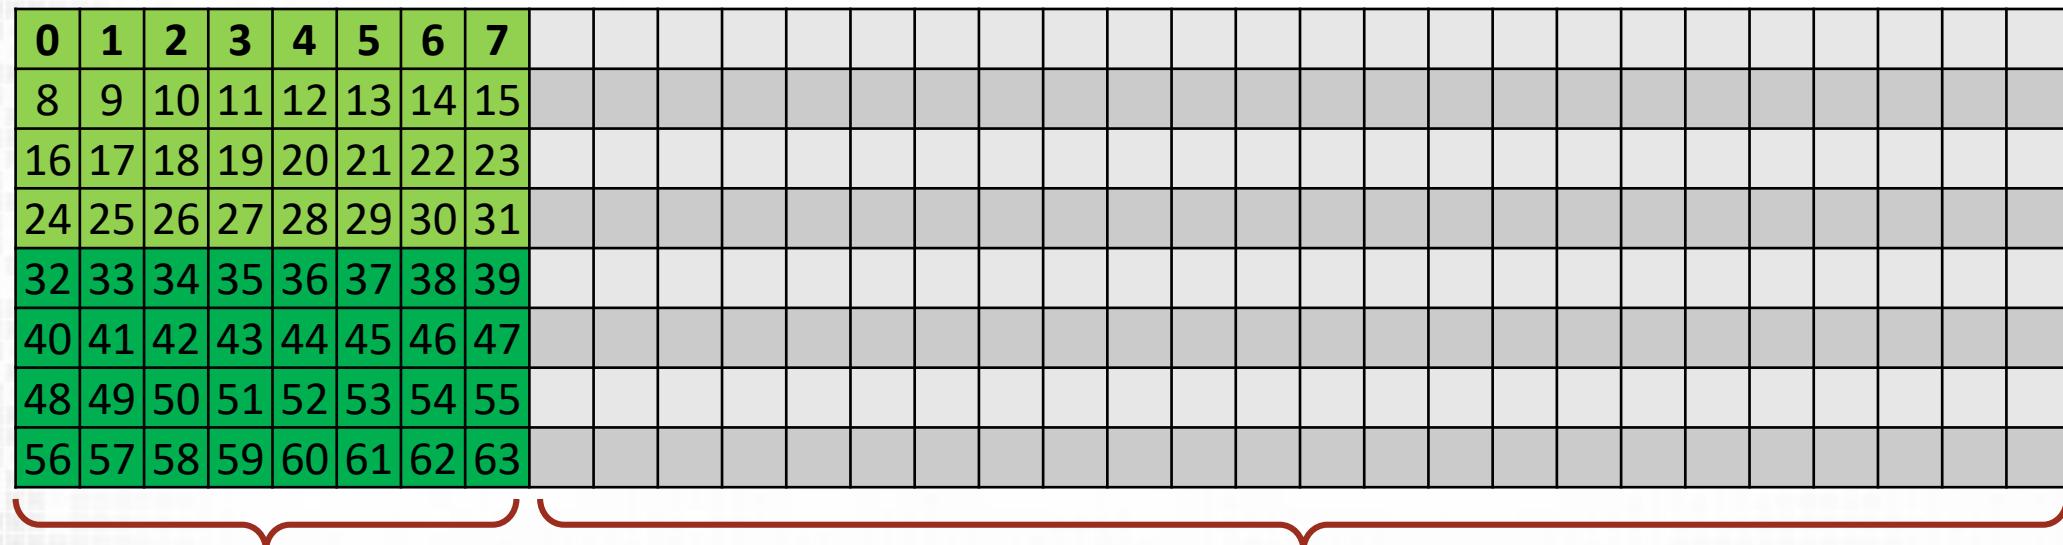
Overfetch from L2 Cache

Line 245:

src[(y+j)*w + (x+i)]

Line 245 for i=0, j=0: **src[x]**

//threads 0-7 only



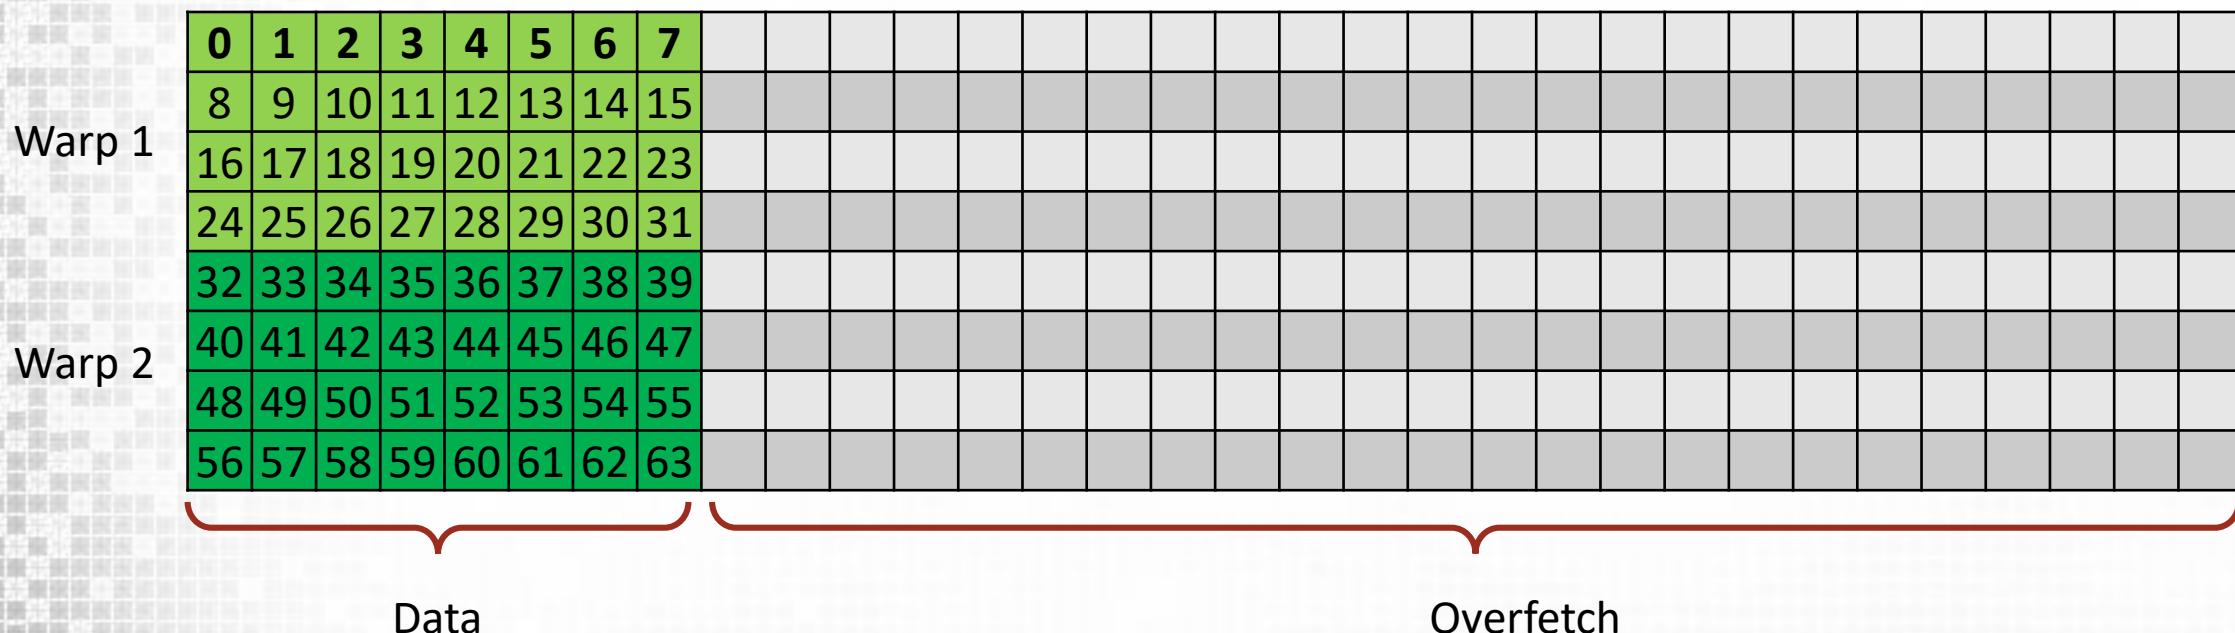
Cache line (always aligned by 32B boundaries)



Overfetch with L1 Caching

Line 245: **src[(y+j)*w + (x+i)]**

Line 245 for i=0, j=0: **src[y*w + x]**

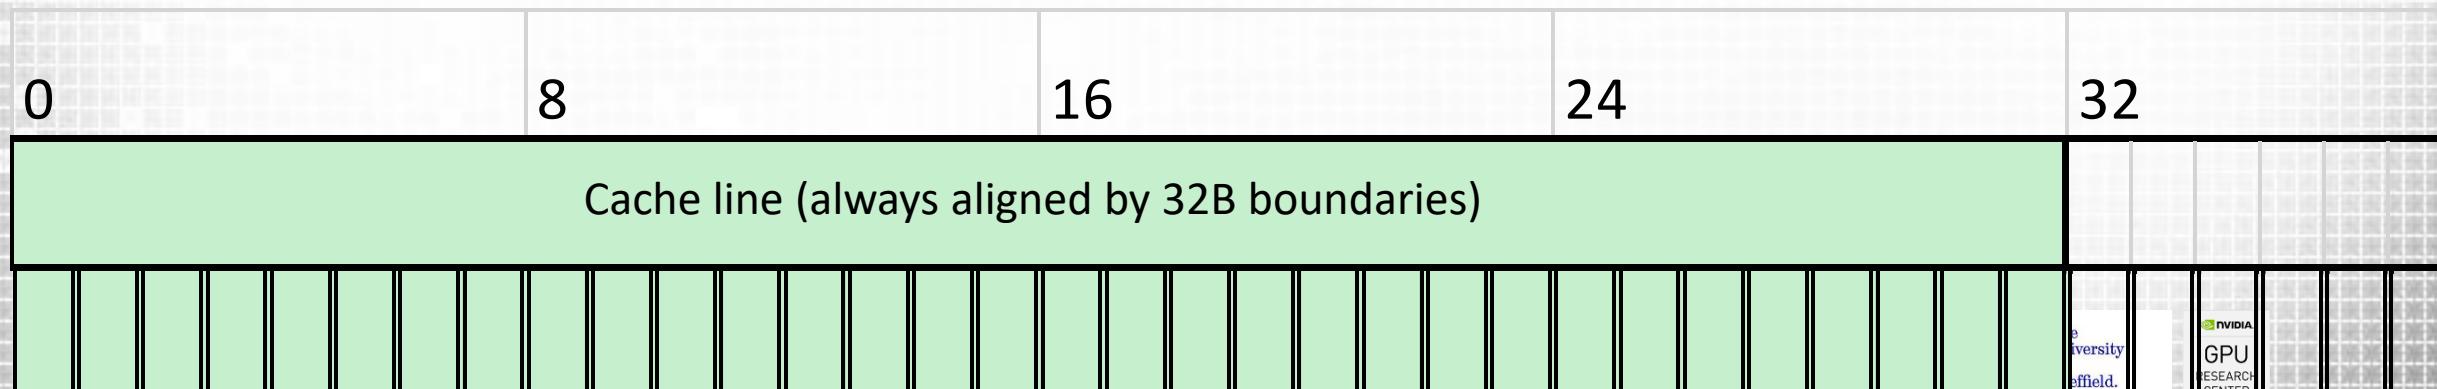


Any Ideas for improving this?

Optimisation: Improved Memory layout

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47	48	49	50	51	52	53	54	55	56	57	58	59	60	61	62	63

- ❑ Minimum block width should be 32 (each thread requires only 1 byte)
- ❑ Use Layout of 32x2



Deploy: Improved Memory layout

Kernel	Time (ms)	Speedup	Rel. Speedup
Gaussian_filter (Step 0)	5.49	1.00x	-
Gaussian_filter (Step 1a)	1.00	5.49x	5.49x



The
University
Of
Sheffield.



NVIDIA
GPU
RESEARCH
CENTER

Break

- ❑ What do we expect the analysis to look like next?
- ❑ Any ideas for what else may be required?



The
University
Of
Sheffield.



Half time summary

- ❑ The guided profiler will help us optimise the right thing
- ❑ Hotspot tells us the most appropriate place to optimise
- ❑ Performance Limiter tells us what to focus on to improve
- ❑ Code may be Memory, Compute or Latency Bound

- ❑ Improvements so far
 - ❑ Changed the access pattern (by changing block size)
 - ❑ Reduced memory dependencies?

- ❑ Profiling Introduction
- ❑ The Problem
- ❑ Visual Profiler Guided Analysis
 - ❑ Iteration 1
 - ❑ Iteration 2



The
University
Of
Sheffield.



Identify the hotspot

- ❑ Examine GPU Usage in Visual Profiler
- ❑ Examine Individual Kernels
 - ❑ Gaussian filter kernel still the highest rank

i Kernel Optimization Priorities

The following kernels are ordered by optimization importance based on execution time and achieved occupancy. Optimization of higher ranked kernels (those that appear first in the list) is more likely to improve performance compared to lower ranked kernels.

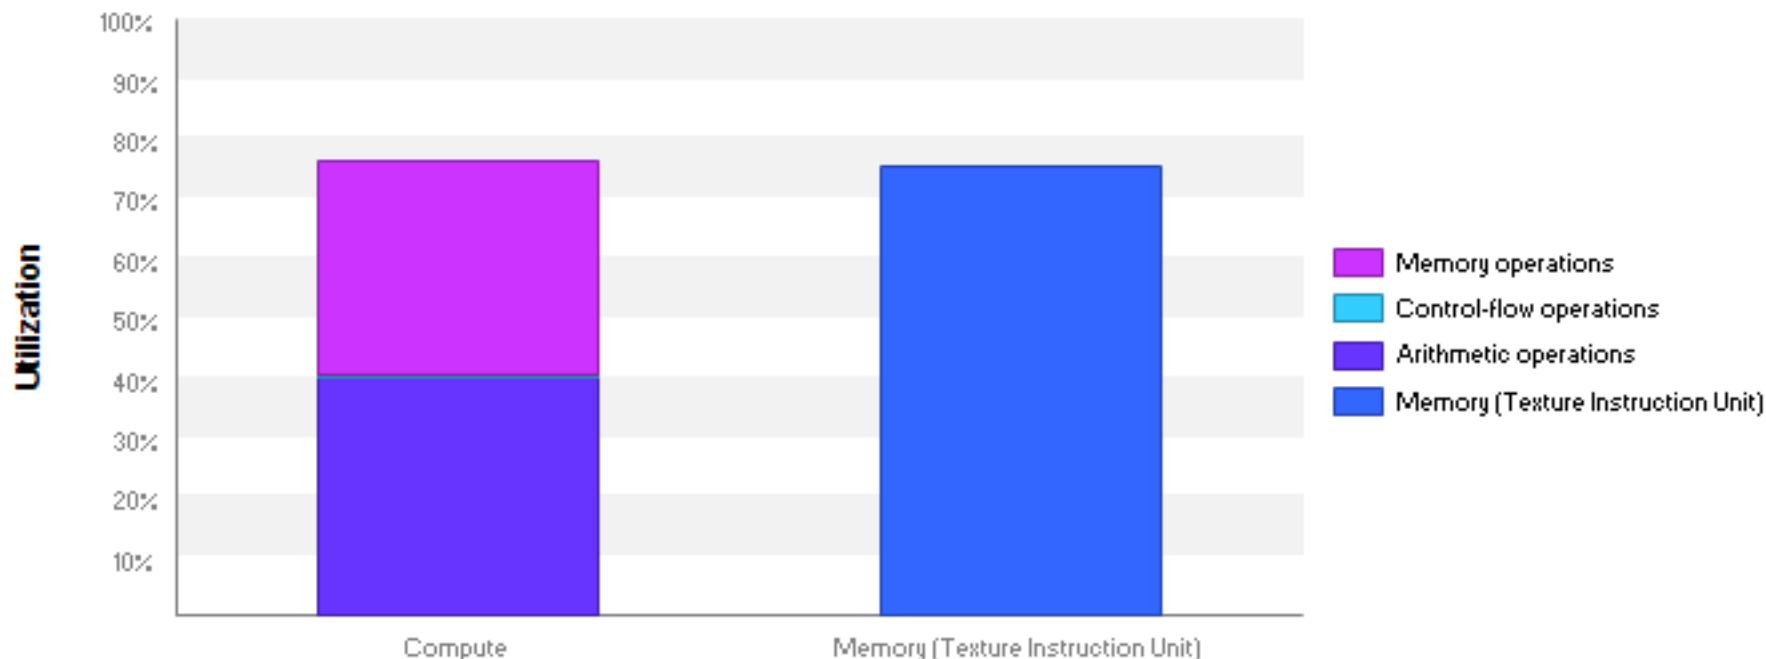
Rank	Description
100	[1 kernel instances] gaussian_filter_7x7_v0(int, int, unsigned char const *, unsigned char*)
29	[1 kernel instances] sobel_filter_3x3_v0(int, int, unsigned char const *, unsigned char*)
14	[1 kernel instances] rgba_to_grayscale_kernel_v0(int, int, uchar4 const *, unsigned char*)



Performance Limiter

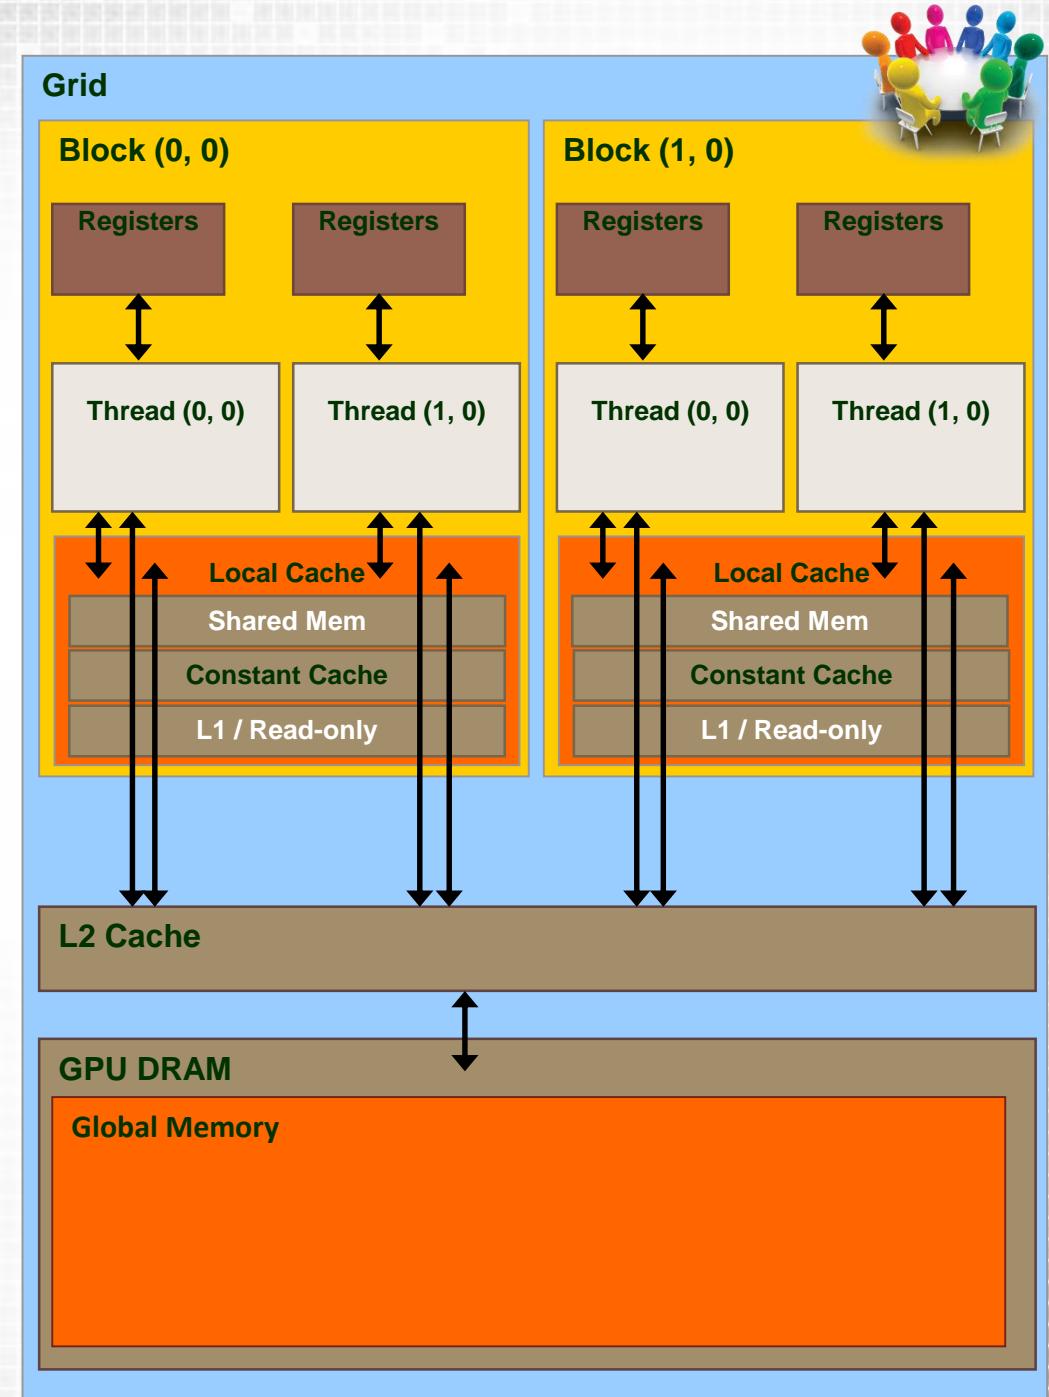
i Kernel Performance Is Bound By Memory Bandwidth

For device "GeForce GTX 980" the kernel's compute utilization is significantly lower than its memory utilization. These utilization levels indicate that the performance of the kernel is most likely being limited by the memory system. For this kernel the limiting factor in the memory system is the bandwidth of the texture instruction units within the multiprocessors.



Tex Instruction Units?

- ❑ What are texture instruction units and why might our code be using them?

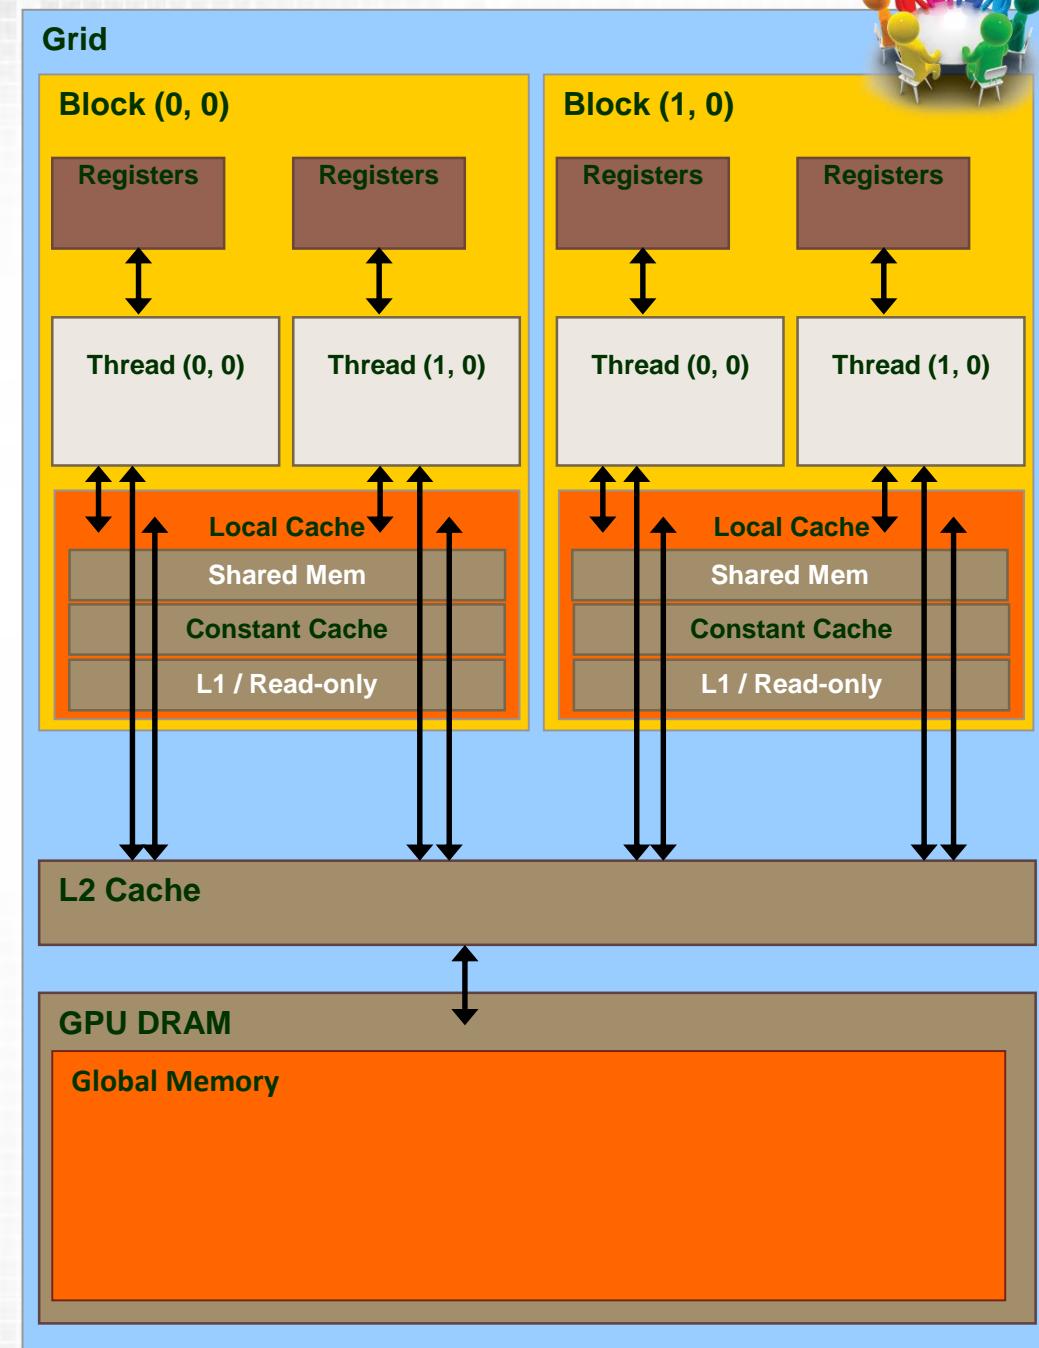


Tex Instruction Units?

- ❑ What are texture instruction units and why might our code be using them?

- ❑ Hint:

```
void gaussian_filter_7x7_v1(int w,  
                           int h,  
                           const uchar *src,  
                           uchar *dst)
```

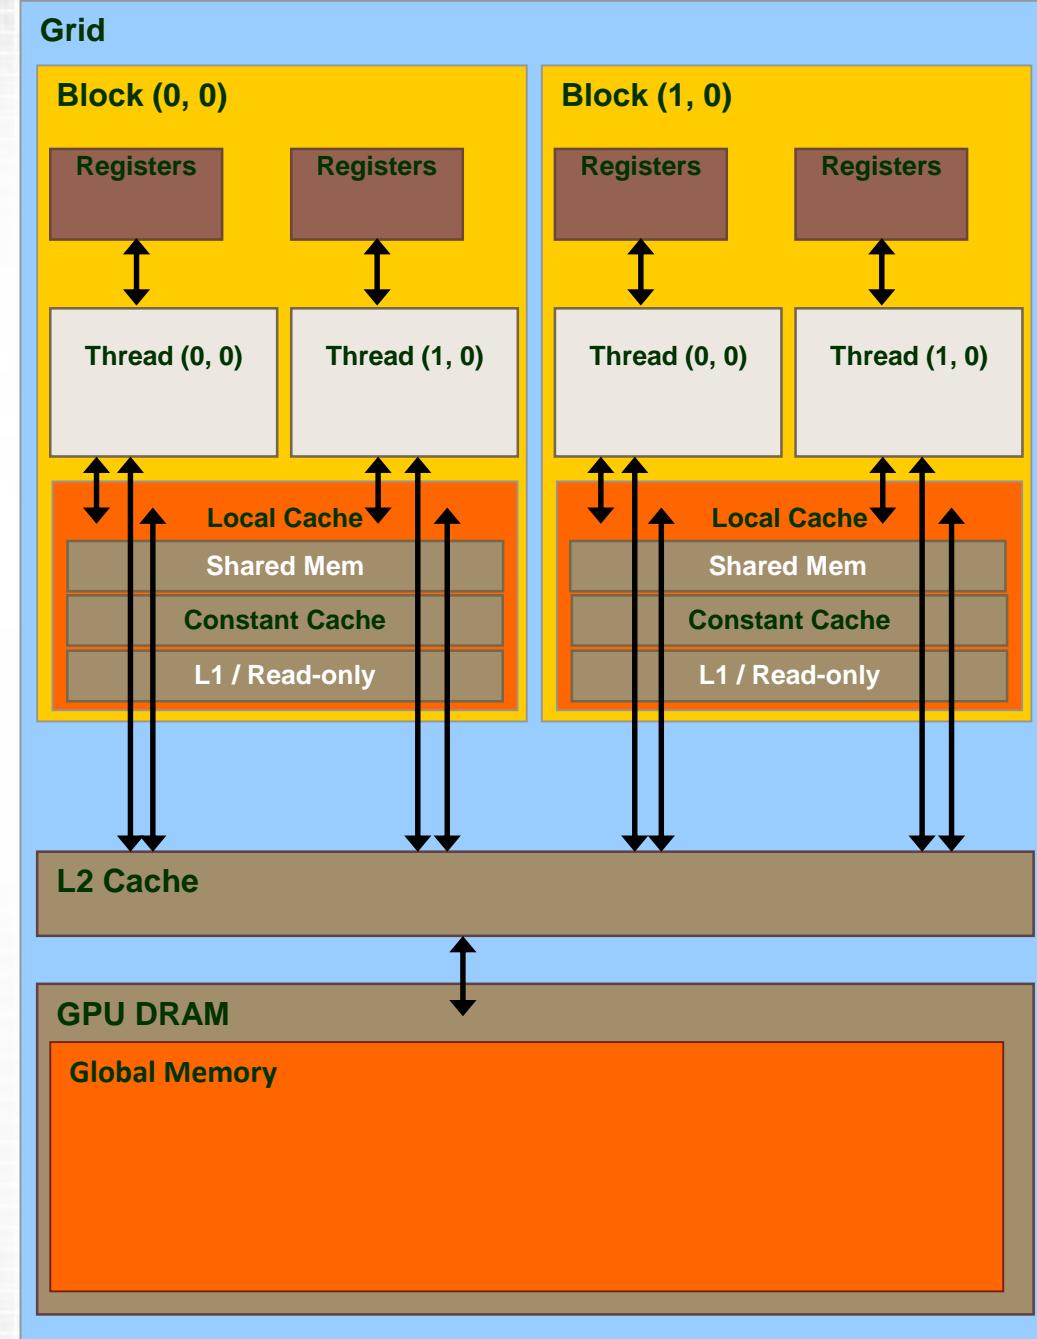


Tex Instruction Units?

- ❑ What are texture instruction units and why might our code be using them?

```
void gaussian_filter_7x7_v1(int w,  
                           int h,  
                           const uchar *src,  
                           uchar *dst)
```

- ❑ Compiler is reading `src` as read-only through Unified L1/Read-Only (texture cache)



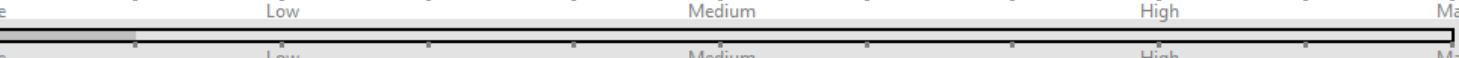
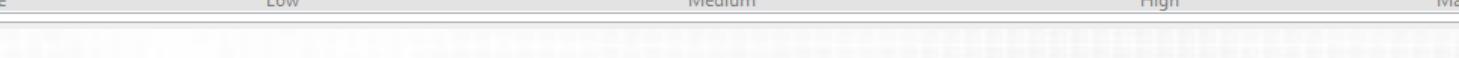
Guided Bandwidth Analysis

⚠ GPU Utilization Is Limited By Memory Instruction Execution

The kernel's performance is potentially limited by the texture instruction units within the multiprocessors. These units are responsible for executing the instructions that result in accesses to memory. The table below shows the memory bandwidth used by this kernel for the various types of memory on the device.

Optimization: Examine the compute analysis results for this kernel to determine how to reduce utilization and improve efficiency of the texture instruction units.

[More...](#)

	Transactions	Bandwidth	Utilization
Shared Memory			
Shared Loads	0	0 B/s	
Shared Stores	0	0 B/s	
Shared Total	0	0 B/s	 Idle Low Medium High Max
L2 Cache			
Reads	11597012	415.583 GB/s	
Writes	128006	4.587 GB/s	
Total	11725018	420.17 GB/s	 Idle Low Medium High Max
Unified Cache			
Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Global Loads	30364232	414.557 GB/s	
Global Stores	128000	4.587 GB/s	
Texture Reads	25061120	898.074 GB/s	
Unified Total	55553352	1,317.218 GB/s	 Idle Low Medium High Max
Device Memory			
Reads	157383	5.64 GB/s	
Writes	127574	4.572 GB/s	
Total	284957	10.212 GB/s	 Idle Low Medium High Max
System Memory [PCIe configuration: Gen2 x16, 5 Gbit/s]			
Reads	0	0 B/s	 Idle Low Medium High Max
Writes	5	179.176 kB/s	 Idle Low Medium High Max

- ❑ We are doing lots of reading/writing through unified cache

Guided Bandwidth Analysis

⚠ Global Memory Alignment and Access Pattern

Memory bandwidth is used most efficiently when each global memory load and store has proper alignment and access pattern.

Optimization: Select each entry below to open the source code to a global load or store within the kernel with an inefficient alignment or access pattern. For each load or store improve the alignment and access pattern of the memory access.

[More...](#)

Line / File	ninspect-gtc.cu - E:\Google Drive\COM4521_6521 - Parallel Computing with GPUs\TEACHING\LABS\Lab09\ninspect-gtc-master
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254082 L2 transactions for 127840 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254082 L2 transactions for 127840 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254082 L2 transactions for 127840 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254241 L2 transactions for 127920 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254082 L2 transactions for 127840 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254241 L2 transactions for 127920 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254082 L2 transactions for 127840 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254241 L2 transactions for 127920 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254082 L2 transactions for 127840 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254241 L2 transactions for 127920 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254082 L2 transactions for 127840 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254241 L2 transactions for 127920 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254082 L2 transactions for 127840 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254241 L2 transactions for 127920 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254400 L2 transactions for 128000 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254082 L2 transactions for 127840 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254241 L2 transactions for 127920 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254241 L2 transactions for 127920 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254241 L2 transactions for 127920 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254082 L2 transactions for 127840 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254400 L2 transactions for 128000 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [253923 L2 transactions for 127760 total executions]
245	Global Load L2 Transactions/Access = 2, Ideal Transactions/Access = 1 [254241 L2 transactions for 127920 total executions]

☐ Still parts of the code reporting 2 transactions per access?



Transaction per request

Line 245:

src[(y+j)*w + (x+i)]

Line 245 for i=1, j=0: **src[x+1]**

What is wrong with this access pattern?



The
University
Of
Sheffield.



NVIDIA
GPU
RESEARCH
CENTER



Transaction per request

Line 245:

src[(y+j)*w + (x+i)]

Line 245 for i=1, j=0: **src[x+1]**

What is wrong with this access pattern?

Hint: Cache Lines are aligned by 32B boundaries



The
University
Of
Sheffield.

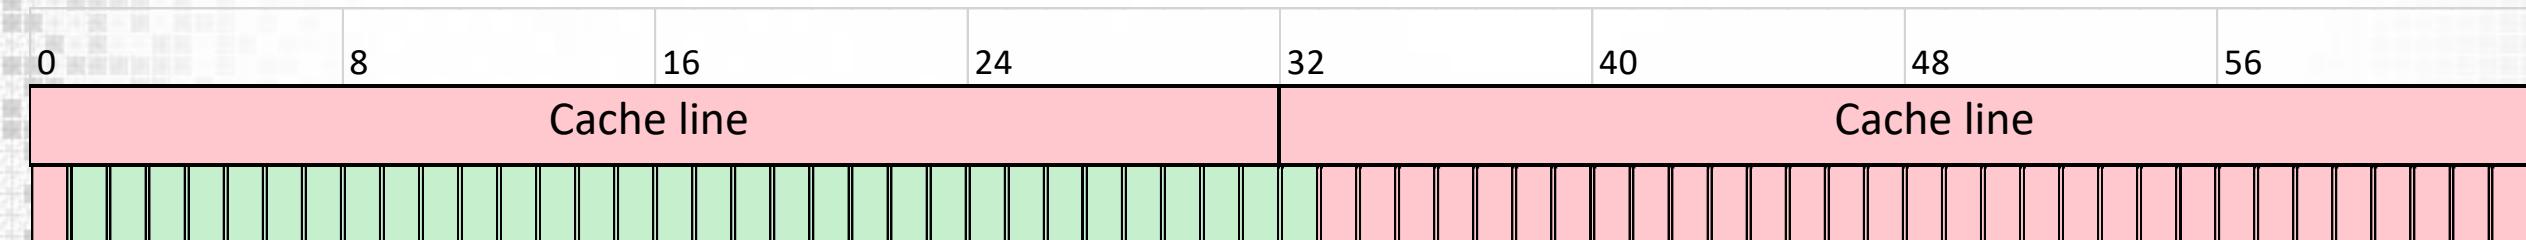


NVIDIA
GPU
RESEARCH
CENTER

Transaction per request

Line 245: **src[(y+j)*w + (x+i)]**

Line 245 for i=1, j=0: **src[x+1]**



We have an offset access pattern

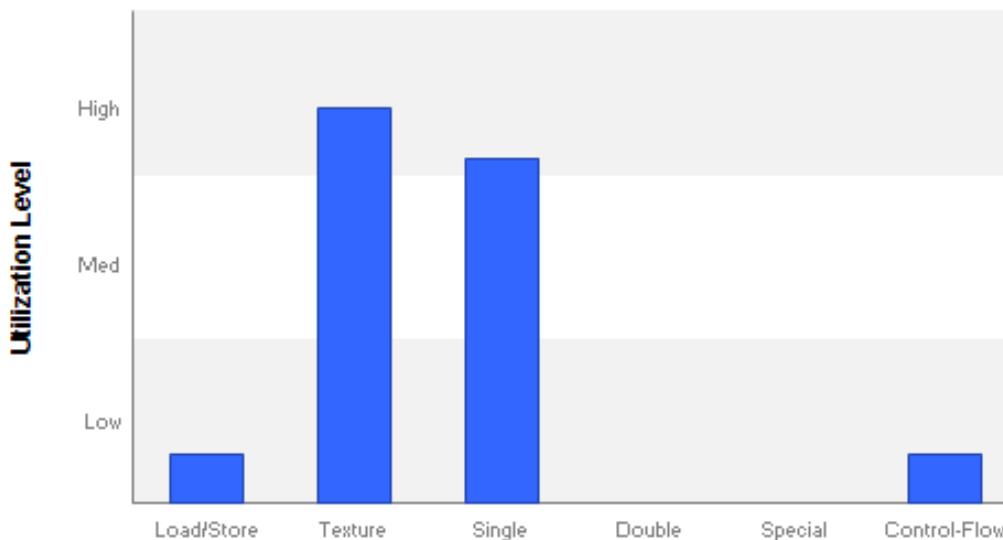
Guided Compute Analysis

- ❑ The guided analysis suggests that lots of our compute cycles are spent issuing texture load/stores

Function Unit Utilization

Different types of instructions are executed on different function units within each SM. Performance can be limited if a function unit is over-used by the instructions executed by the kernel. The following results show that the kernel's performance is not limited by overuse of any function unit.

Load/Store - Load and store instructions for shared and constant memory.
Texture - Load and store instructions for local, global, and texture memory.
Single - Single-precision integer and floating-point arithmetic instructions.
Double - Double-precision floating-point arithmetic instructions.
Special - Special arithmetic instructions such as sin, cos, popc, etc.
Control-Flow - Direct and indirect branches, jumps, and calls.



Guided Latency Analysis: Occupancy

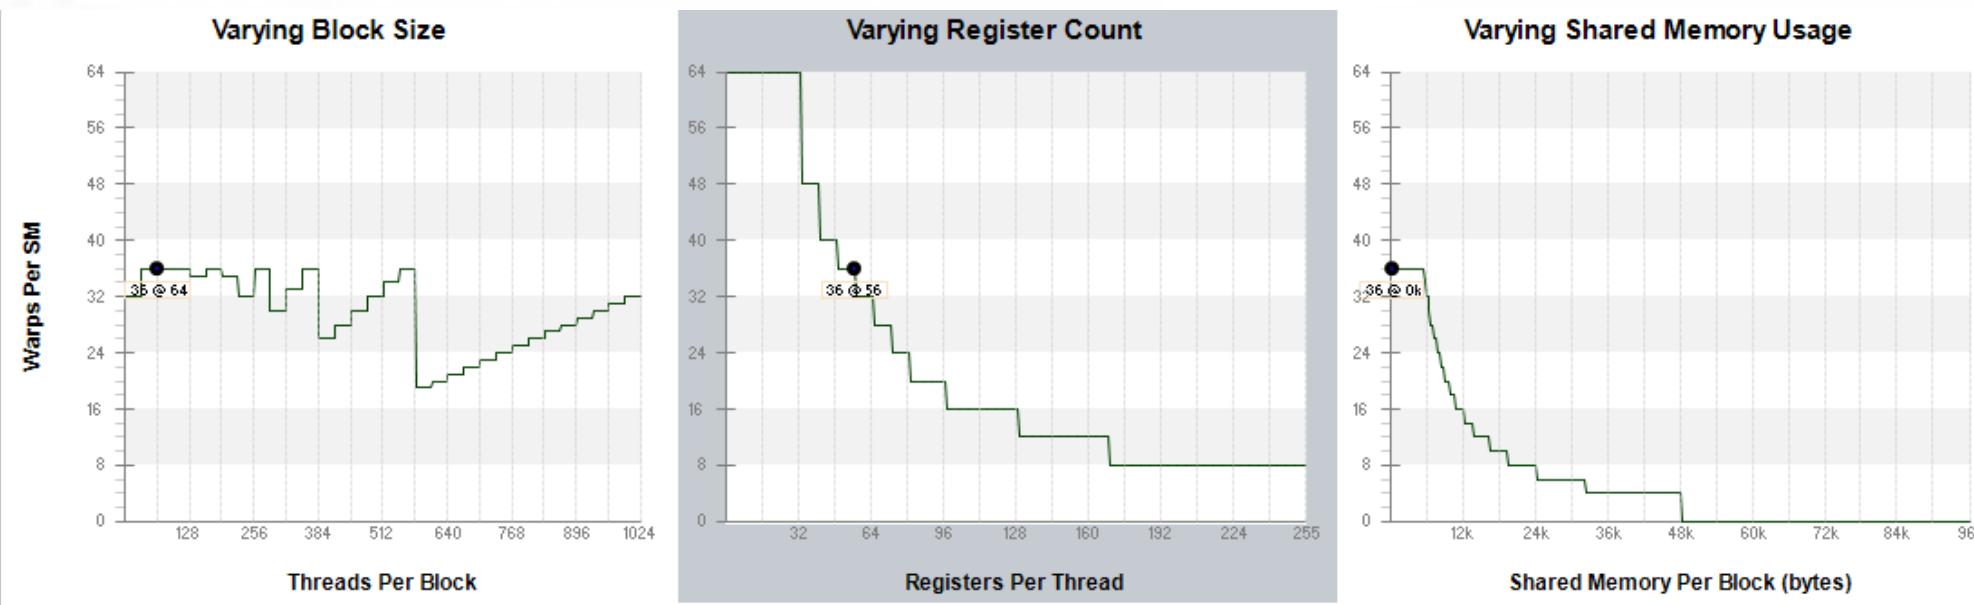
⚠ GPU Utilization May Be Limited By Register Usage

Theoretical occupancy is less than 100% but is large enough that increasing occupancy may not improve performance. You can attempt the following optimization to increase the number of warps on each SM but it may not lead to increased performance.

The kernel uses 56 registers for each thread (3584 registers for each block). This register usage is likely preventing the kernel from fully utilizing the GPU. Device "GeForce GTX 980" provides up to 65536 registers for each block. Because the kernel uses 3584 registers for each block each SM is limited to simultaneously executing 18 blocks (36 warps). Chart "Varying Register Count" below shows how changing register usage will change the number of blocks that can execute on each SM.

Optimization: Use the `-maxregcount` flag or the `_launch_bounds_` qualifier to decrease the number of registers used by each thread. This will increase the number of blocks that can execute on each SM. On devices with Compute Capability 5.2 turning global cache off can increase the occupancy limited by register usage.

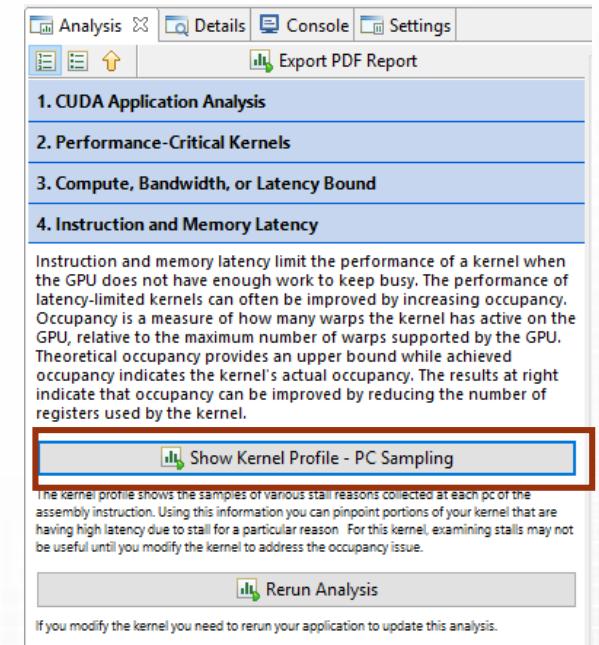
[More...](#)



Guided Latency Analysis: Occupancy

- ❑ Register usage is very high
- ❑ Occupancy currently limited by register usage

- ❑ Increasing occupancy might not help us however as we are dominated by texture load stores
 - ❑ More work per SMP will just mean even more texture load stores!
 - ❑ We can confirm this by looking at the unguided analysis: Kernel Latency



PC Sampling

i Instruction Latencies

Instruction stall reasons indicate the condition that prevents warps from executing on any given cycle. The following chart shows the break-down of stalls reasons averaged over the entire execution of the kernel. The kernel has low theoretical or achieved occupancy. Therefore, it is likely that the instruction stall reasons described below are not the primary limiters of performance and so should not be considered until any occupancy issues are resolved.

Constant - A constant load is blocked due to a miss in the constants cache.

Memory Dependency - A load/store cannot be made because the required resources are not available or are fully utilized, or too many requests of a given type are outstanding. Data request stalls can potentially be reduced by optimizing memory alignment and access patterns.

Pipeline Busy - The compute resource(s) required by the instruction is not yet available.

Texture - The texture sub-system is fully utilized or has too many outstanding requests.

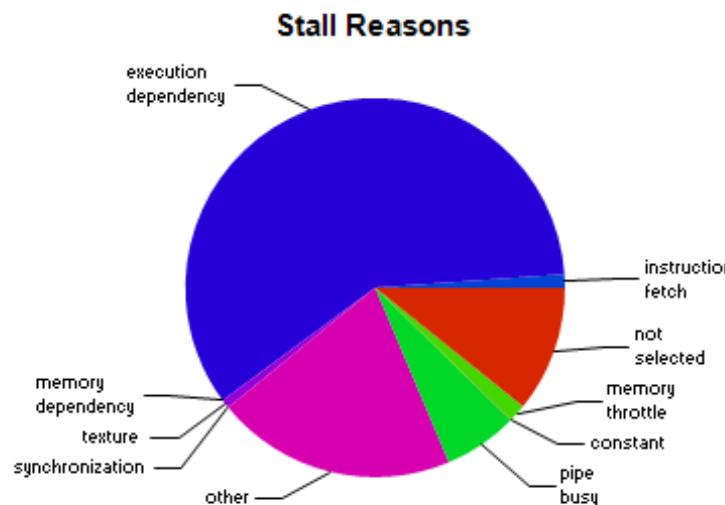
Execution Dependency - An input required by the instruction is not yet available. Execution dependency stalls can potentially be reduced by increasing instruction-level parallelism.

Synchronization - The warp is blocked at a `_syncthreads()` call.

Not Selected - Warp was ready to issue, but some other warp issued instead. You may be able to sacrifice occupancy without impacting latency hiding and doing so may help improve cache hit rates.

Memory Throttle - Large number of pending memory operations prevent further forward progress. These can be reduced by combining several memory transactions into one.

Instruction Fetch - The next assembly instruction has not yet been fetched.





Execution/Memory Dependency

- ❑ Rank these are best to worst
- ❑ Which have instruction and memory dependencies?

```
int a = b + c;  
  
int d = a + e;  
  
//b, c and e are local ints
```

```
int a = b[i];  
  
int d = a + e;  
  
//b is global memory  
//i and e are local ints
```

```
int a = b + c;  
  
int d = e + f;  
  
//b, c, e and f are local ints
```



The
University
Of
Sheffield.



NVIDIA
GPU
RESEARCH
CENTER

Instruction/Memory Dependency

- ❑ Rank these are best to worst
- ❑ Which have instruction and memory dependencies?



```
int a = b + c;  
          ↑  
int d = a + e;  
  
//b, c and e are local ints
```

```
int a = b[i];  
          ↑  
int d = a + e;  
  
//b is global memory  
//i and e are local ints
```

```
int a = b + c;  
  
int d = e + f;  
  
//b, c, e and f are local ints
```

- ❑ Instruction Dependency
- ❑ Second add must wait for first

- ❑ Memory Dependency
- ❑ Second add must wait for memory request

- ❑ No dependencies
- ❑ Independent Adds



Analysis

- ❑ Our compute engine is dominated by load/store instructions for the texture cache
 - ❑ Our texture bandwidth is good BUT
- ❑ Our warps are stalling as instructions are waiting to issue texture fetch instructions
- ❑ We still have poorly aligned access pattern within our inner loops

- ❑ Solution: Reduce dependencies on texture loads
 - ❑ Move data closer to the SMP
 - ❑ Only read from global memory with nicely aligned cache lines
- ❑ How?

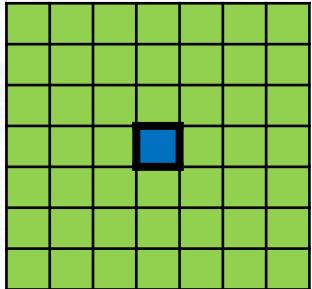
Analysis

- ❑ Our compute engine is dominated by load/store instructions for the texture cache
 - ❑ Our texture bandwidth is good BUT
- ❑ Our warps are stalling as instructions are waiting to issue texture fetch instructions
- ❑ We still have poorly aligned access pattern within our inner loops

- ❑ Solution: Reduce dependencies on texture loads
 - ❑ Move data closer to the SMP
 - ❑ Only read from global memory with nicely aligned cache lines
 - ❑ **Shared Memory**

Shared Memory

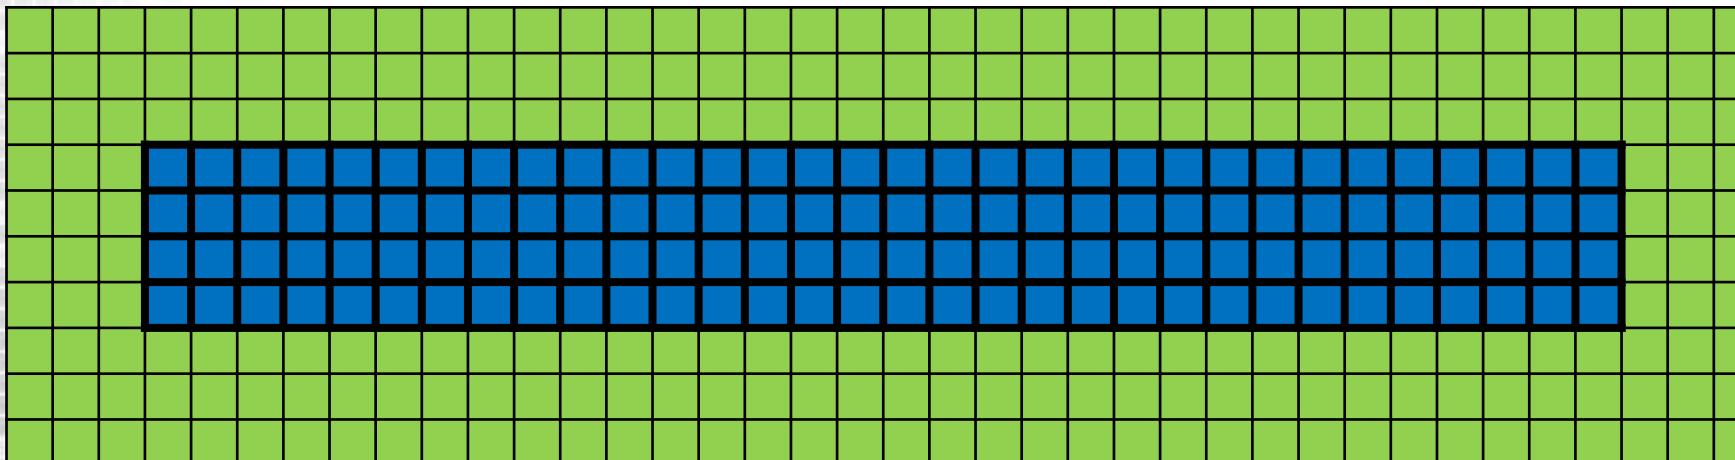
Single thread uses $7 \times 7 = 49$ values



Use shared memory to store all pixels for the block

What important factor should we be considering?

Single block (32x4) uses $32 \times 10 = 320$ values



Also increased
Block size

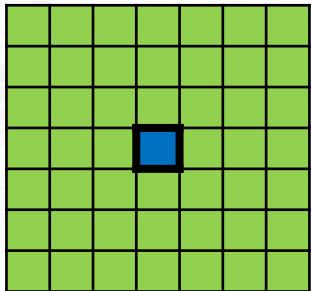


The
University
Of
Sheffield.

NVIDIA
GPU
RESEARCH
CENTER

Shared Memory

Single thread uses $7 \times 7 = 49$ values

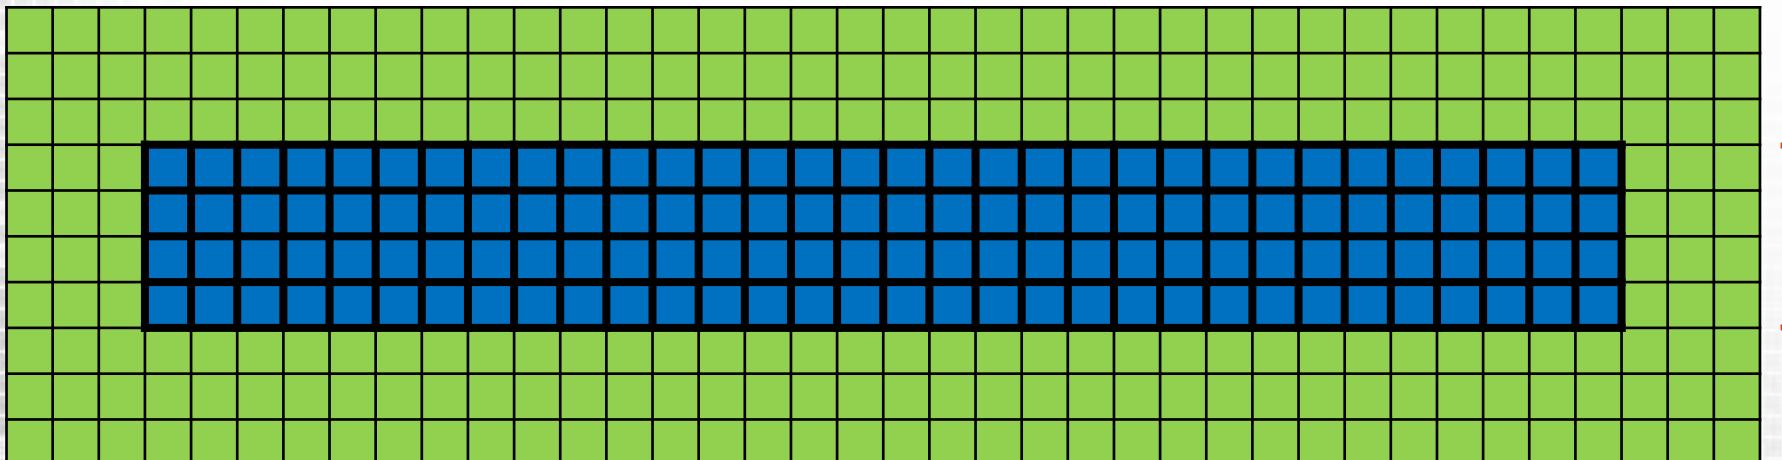


Use shared memory to store all pixels for the block

```
__shared__ unsigned char smem_pixels[10][64]
```

SM bank conflicts

Single block (32x4) uses $38 \times 10 = 680$ values



Also increased
Block size



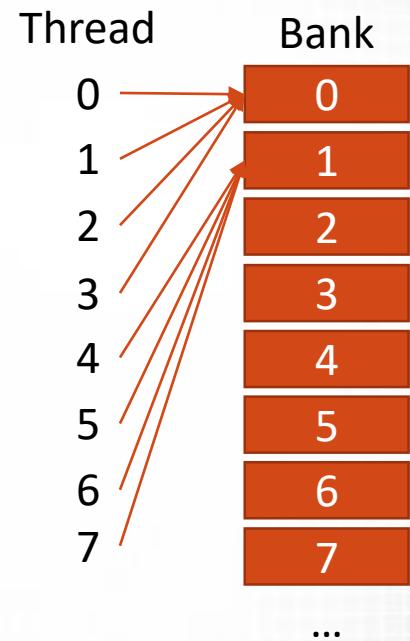
The
University
Of
Sheffield.



BUT WAIT ! ! ! ! ! ! ! ! ! ! ! !

- ❑ Wouldn't aligned char access have 4 way bank conflicts?
 - ❑ NOT for Compute Mode 2.0+...

*"A shared memory request for a warp does not generate a bank conflict between two threads that access any address within the same 32-bit word (even though the two addresses fall in the same bank): In that case, for read accesses, the word is **broadcast** to the requesting threads (multiple words can be broadcast in a single transaction) ..."*



I.e. A Stride of less than 1 (4B word) can be read conflict free if threads access aligned data

Improvement

❑ Significant

Kernel	Time (ms)	Speedup	Rel. Speedup
Gaussian_filter (Step 0)	5.49	1.00x	-
Gaussian_filter (Step 1a)	1.00	5.49x	5.49x
Gaussian_filter (Step 40)	0.49	11.20x	2.04x



The
University
Of
Sheffield.



- ❑ Profiling Introduction
- ❑ The Problem
- ❑ Visual Profiler Guided Analysis
 - ❑ Iteration 1
 - ❑ Iteration 2
 - ❑ Iteration 3



The
University
Of
Sheffield.



Identify the hotspot

- ❑ Examine GPU Usage in Visual Profiler
- ❑ Examine Individual Kernels
 - ❑ Gaussian filter kernel still the highest rank
 - ❑ Getting much closer though

i Kernel Optimization Priorities

The following kernels are ordered by optimization importance based on execution time and achieved occupancy. Optimization of higher ranked kernels (those that appear first in the list) is more likely to improve performance compared to lower ranked kernels.

Rank	Description
100	[1 kernel instances] gaussian_filter_7x7_v2(int, int, unsigned char const *, unsigned char*)
60	[1 kernel instances] sobel_filter_3x3_v0(int, int, unsigned char const *, unsigned char*)
29	[1 kernel instances] rgba_to_grayscale_kernel_v0(int, int, uchar4 const *, unsigned char*)



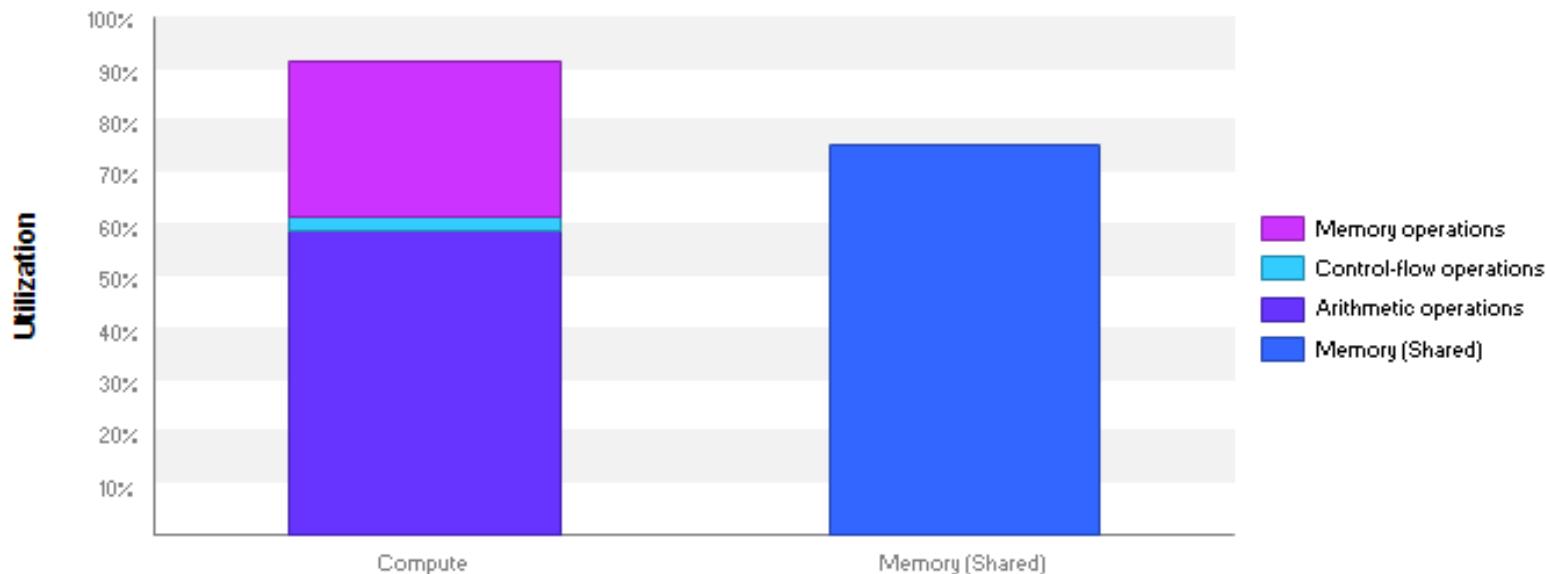
The
University
Of
Sheffield.



Performance Limiter

i Kernel Performance Is Bound By Memory Bandwidth

For device "GeForce GTX 980" the kernel's compute utilization is significantly lower than its memory utilization. These utilization levels indicate that the performance of the kernel is most likely being limited by the memory system. For this kernel the limiting factor in the memory system is the bandwidth of the Shared memory.



- Actually very close to magical 60% of compute
- Lets examine
 - 1) The compute analysis
 - 2) The latency analysis

Guided Bandwidth Analysis

⚠ GPU Utilization Is Limited By Memory Bandwidth

The following table shows the memory bandwidth used by this kernel for the various types of memory on the device. The table also shows the utilization of each memory type relative to the maximum throughput supported by the memory. The results show that the kernel's performance is potentially limited by the bandwidth available from one or more of the memories on the device.

Optimization: Try the following optimizations for the memory with high bandwidth utilization.

Shared Memory - If possible use 64-bit accesses to shared memory and 8-byte bank mode to achieve 2x throughput.

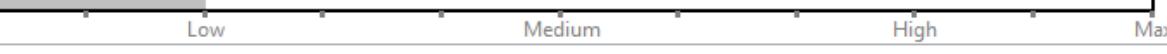
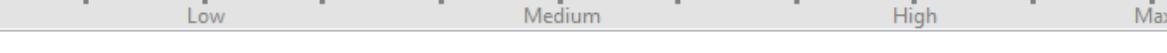
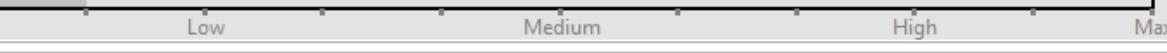
L2 Cache - Align and block kernel data to maximize L2 cache efficiency.

Unified Cache - Reallocate texture data to shared or global memory. Resolve alignment and access pattern issues for global loads and stores.

Device Memory - Resolve alignment and access pattern issues for global loads and stores.

System Memory (via PCIe) - Make sure performance critical data is placed in device or shared memory.

[More...](#)

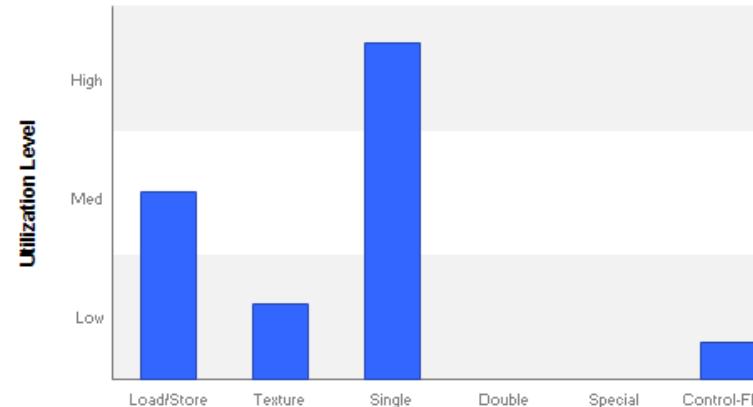
	Transactions	Bandwidth	Utilization
Shared Memory			
Shared Loads	6272000	1,826.774 GB/s	
Shared Stores	638720	186.033 GB/s	
Shared Total	6910720	2,012.807 GB/s	 High
L2 Cache			
Reads	1318182	95.983 GB/s	
Writes	128006	9.321 GB/s	
Total	1446188	105.304 GB/s	 Medium
Unified Cache			
Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Global Loads	3171255	92.366 GB/s	
Global Stores	128000	9.32 GB/s	
Texture Reads	2540993	185.022 GB/s	
Unified Total	5840248	286.707 GB/s	 High
Device Memory			
Reads	128446	9.353 GB/s	
Writes	128008	9.321 GB/s	
Total	256454	18.674 GB/s	 Medium
System Memory [PCIe configuration: Gen2 x16, 5 Gbit/s]			
Reads	0	0 B/s	 Idle
Writes	5	364.073 kB/s	 Low



i Function Unit Utilization

Different types of instructions are executed on different function units within each SM. Performance can be limited if a function unit is over-used by the instructions executed by the kernel. The following results show that the kernel's performance is not limited by overuse of any function unit.

- Load/Store - Load and store instructions for shared and constant memory.
- Texture - Load and store instructions for local, global, and texture memory.
- Single - Single-precision integer and floating-point arithmetic instructions.
- Double - Double-precision floating-point arithmetic instructions.
- Special - Special arithmetic instructions such as sin, cos, popc, etc.
- Control-Flow - Direct and indirect branches, jumps, and calls.



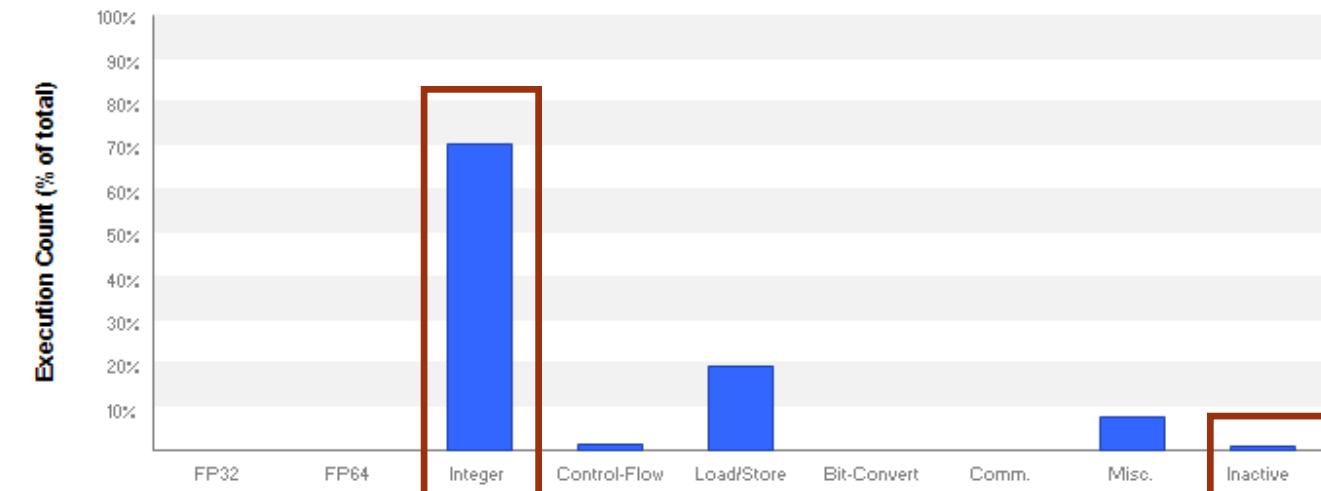
Compute Analysis

- We are simply doing lots of compute
- Additional floating point operations graph shows no activity i.e. all of our instructions are Integer

i Instruction Execution Counts

The following chart shows the mix of instructions executed by the kernel. The instructions are grouped into classes and for each class the chart shows the percentage of thread execution cycles that were devoted to executing instructions in that class. The "Inactive" result shows the thread executions that did not execute any instruction because the thread was predicated or inactive due to divergence.

Screenshot of the CUDA Application Analysis interface. The main menu includes "Export PDF Report", "1. CUDA Application Analysis", "2. Performance-Critical Kernels", "3. Compute, Bandwidth, or Latency Bound", and "4. Compute Resources". A note below states: "GPU compute resources limit the performance of a kernel when those resources are insufficient or poorly utilized." A red box highlights the "Show Kernel Profile - Instruction Execution" button, which is described in the tooltip as showing the execution count, inactive threads, and predicated threads for each source and assembly line of the kernel. Another red box highlights the "Rerun Analysis" button at the bottom.



Compute Analysis by Line

- ❑ Selecting the CUDA function from compute analysis results allows a line by line breakdown
 - ❑ This will switch to unguided analysis

Line	Exec Count	File - /E:/Google Drive/COM4521_6521 - Parallel Computing with GPUs/TEACHING/LABS/Lab09/nsight-gtc-master/nsight-gtc.cu
304	512000	uchar *smem_img_ptr = &smem_img[threadIdx.y][threadIdx.x];
305	2620560	for(int iy = y-3 ; iy <= blockIdx.y*blockDim.y+6 ; iy += 4, smem_img_ptr += 4*64) {
306		smem_img_ptr[0] = in_img(x-3, iy, w, h) ? src[iy*w + (x-3)] : 0;
307	5365280	smem_img_ptr[32] = in_img(x+29, iy, w, h) ? src[iy*w + (x+29)] : 0; // 29 = 32-3.
308	1916240	}
309		__syncthreads();
310	128000	
311		// Load the 48 neighbours and myself.
312		int n[7][7];
313		for(int j = 0 ; j <= 6 ; ++j)
314		for(int i = 0 ; i <= 6 ; ++i)
315		n[j][i] = smem_img[threadIdx.y+j][threadIdx.x+i];
316	6272000	
317		// Compute the convolution.
318		int p = 0;
319		for(int j = 0 ; j < 7 ; ++j)
320		for(int i = 0 ; i < 7 ; ++i)
321		p += gaussian_filter[j][i] * n[j][i];
322	18816000	
323		// Store the result.
324		if(in_img(x, y, w, h))
325	640000	dst[y*w + x] = (uchar)(p / 256);
326	1536000	
327	128000	}
328		// =====
329		
330		
331		_global_ void gaussian_filter_7x7_v3(int w, int h, const uchar * __restrict src, uchar *dst)
332		{

Also PTX instruction breakdown provided

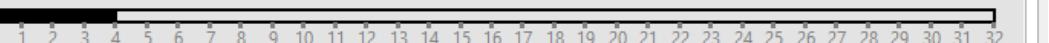
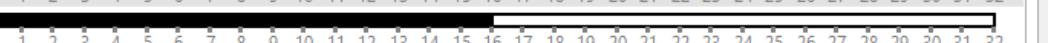
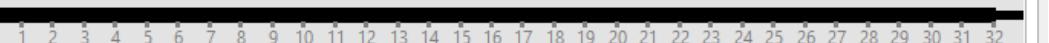


Guided Latency Analysis

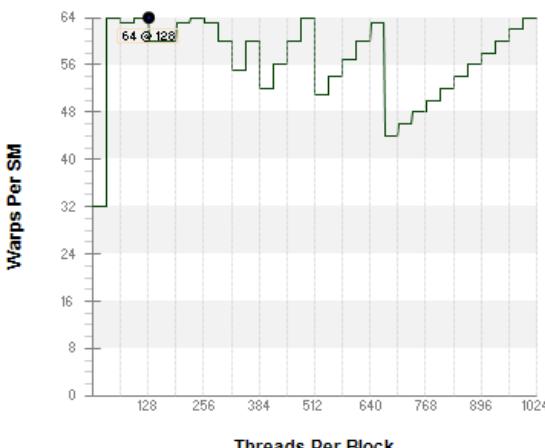
i Occupancy Is Not Limiting Kernel Performance

The kernel's block size, register usage, and shared memory usage allow it to fully utilize all warps on the GPU.

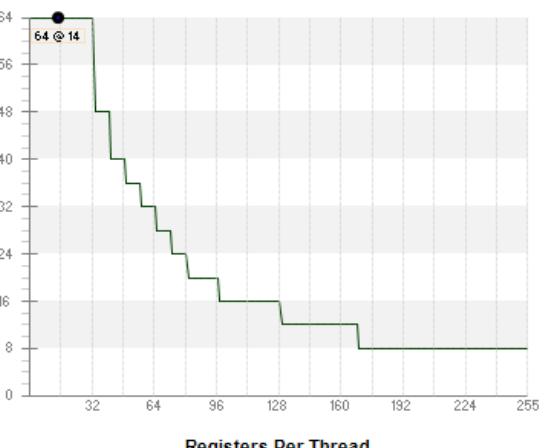
[More...](#)

Variable	Achieved	Theoretical	Device Limit	Grid Size: [80,400,1] (32000 blocks) Block Size: [32,4,1] (128 threads)
Occupancy Per SM				
Active Blocks		16	32	
Active Warps	60.99	64	64	
Active Threads		2048	2048	
Occupancy	95.3%	100%	100%	
Warp Metrics				
Threads/Block		128	1024	
Warp/Block		4	32	
Block Limit		16	32	
Registers				
Registers/Thread		14	255	
Registers/Block		2048	65536	
Block Limit		32	32	
Shared Memory				
Shared Memory/Block		640	98304	
Block Limit		128	32	

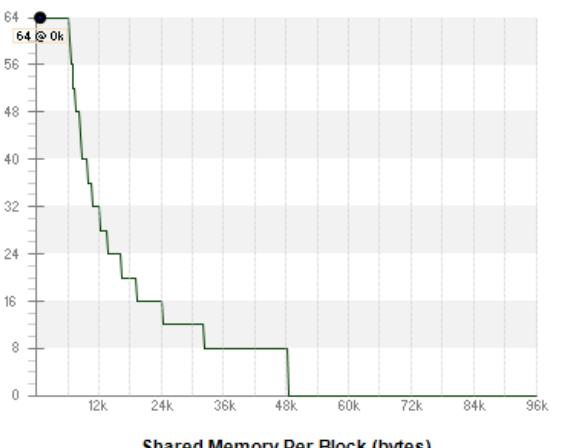
Varying Block Size



Varying Register Count



Varying Shared Memory Usage



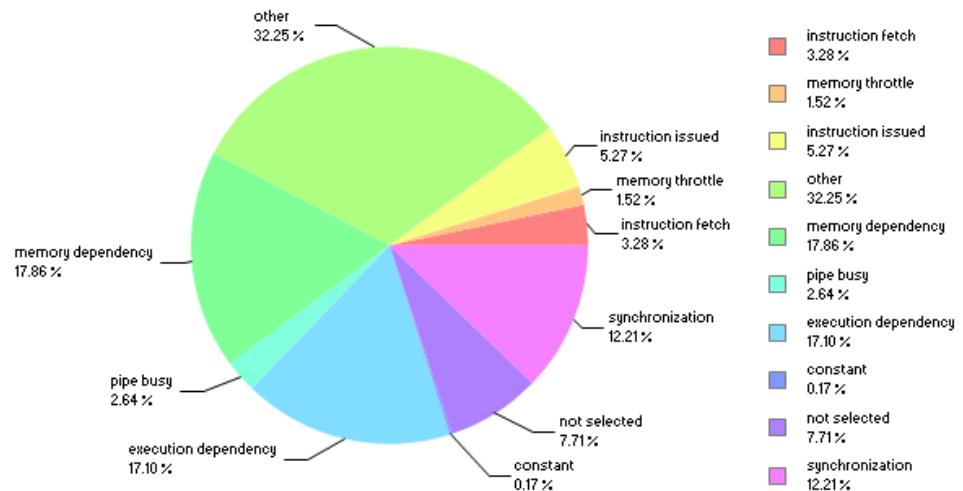
Would changing the block size, register usage or amount of shared memory per block improve occupancy?

Guided Latency Analysis

Source files :

file:///E:/Google%20Drive/COM4521_6521-%20-%20Parallel%20Computing%20with%20GPUs/TEACHING/LABS/Lab09/nsight-gtc-master/nsight-gtc.cu

Sample distribution



Line by Line Breakdown

The Kernel Profile - PC Sampling gives the number of samples for each source and assembly line with various stall reasons. Using this information you can pinpoint portions of your kernel that are introducing latencies and the reason for the latency. Samples are taken in round robin order for all active warps at a fixed number of cycles regardless of whether the warp is issuing an instruction or not.

Instruction Issued - Warp was issued

Instruction Fetch - The next assembly instruction has not yet been fetched.

Execution Dependency - An input required by the instruction is not yet available. Execution dependency stalls can potentially be reduced by increasing instruction-level parallelism.

Memory Dependency - A load/store cannot be made because the required resources are not available or are fully utilized, or too many requests of a given type are outstanding. Data request stalls can potentially be reduced by optimizing memory alignment and access patterns.

Texture - The texture sub-system is fully utilized or has too many outstanding requests.

Synchronization - The warp is blocked at a `_syncthreads()` call.

Constant - A constant load is blocked due to a miss in the constants cache.

Pipe Busy - The compute resource(s) required by the instruction is not yet available.

Memory Throttle - Large number of pending memory operations prevent further forward progress. These can be reduced by combining several memory transactions into one.

Not Selected - Warp was ready to issue, but some other warp issued instead. You may be able to sacrifice occupancy without impacting latency hiding and doing so may help improve cache hit rates.

Other - The warp is blocked for an uncommon reason.

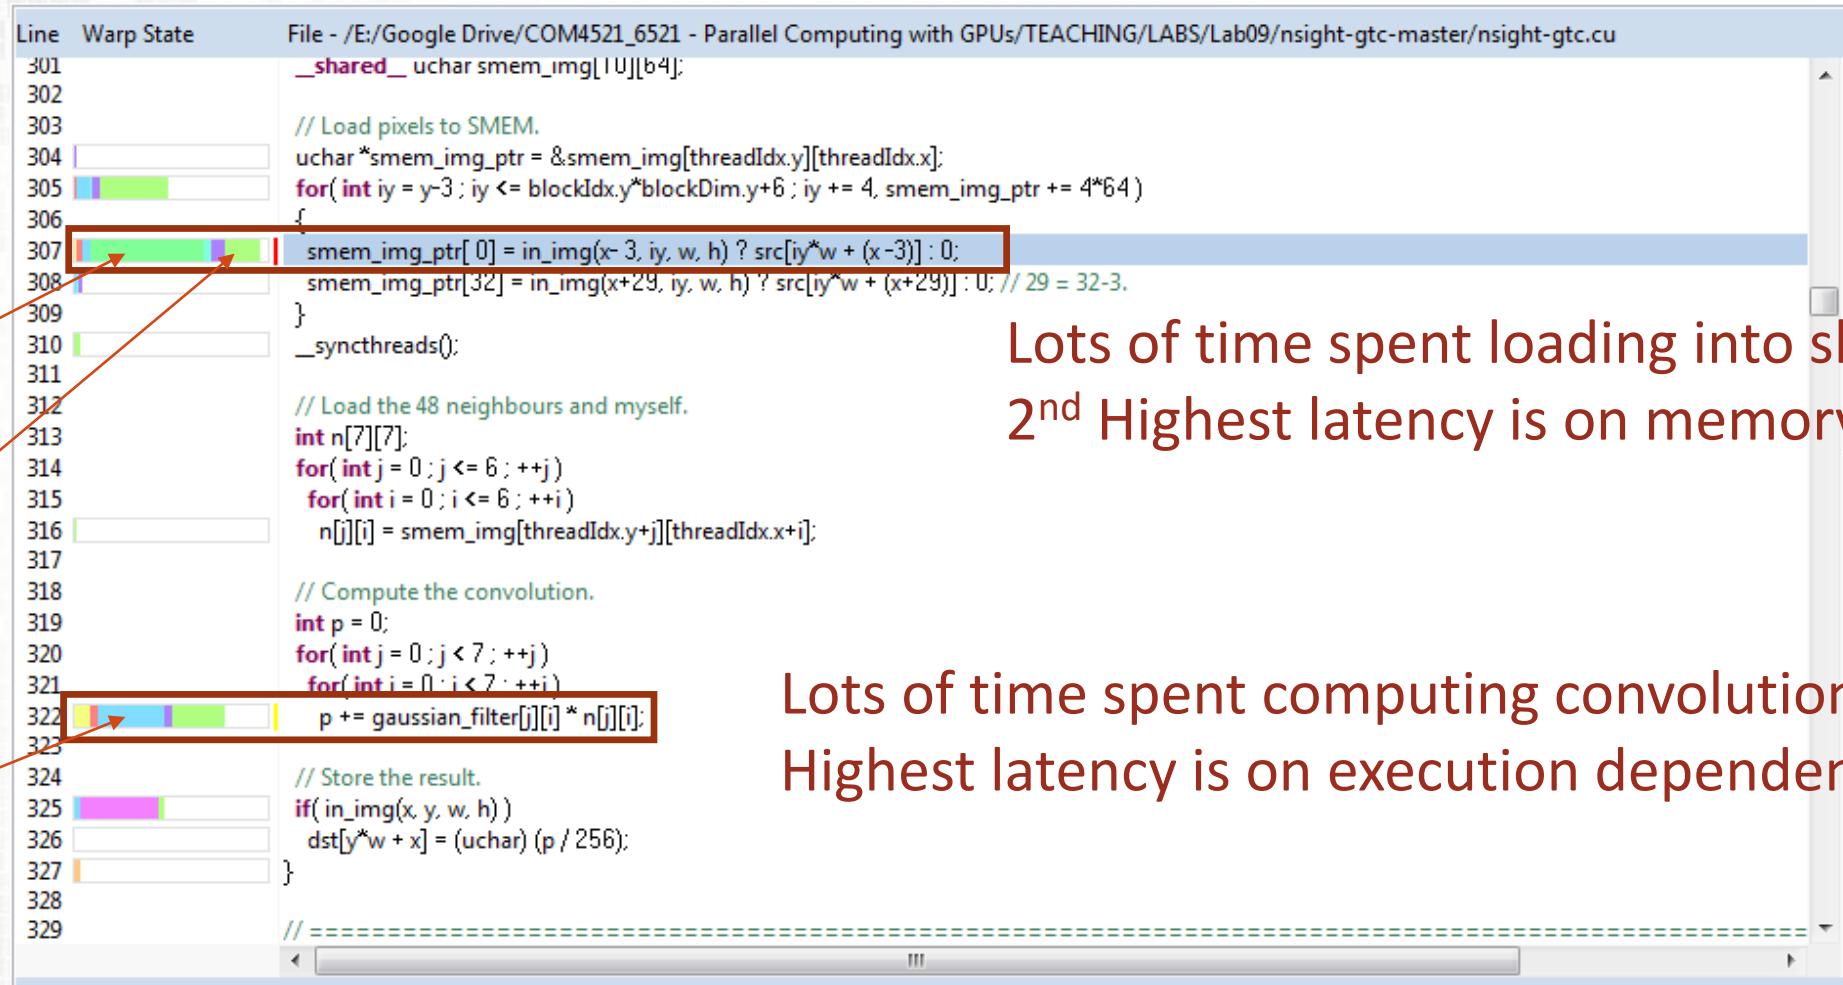
[More...](#)

Latency Overview: Other 32.25%

- ❑ Stall reason other generally means that there is no obvious action to improve performance
- ❑ Other stall reasons may indicate either;
 1. Execution unit is busy
 - ❑ Solution: Potentially reduce use of low throughput integer operations if possible
 2. Register bank conflicts : a compiler issue that can sometimes be made worst by heavy use of vector data types
 - ❑ Solution: None
 3. Too few warps per scheduler
 - ❑ Solution: Increase occupancy, decrease latency

Guided Latency Analysis: Line by Line

memory
other
execution



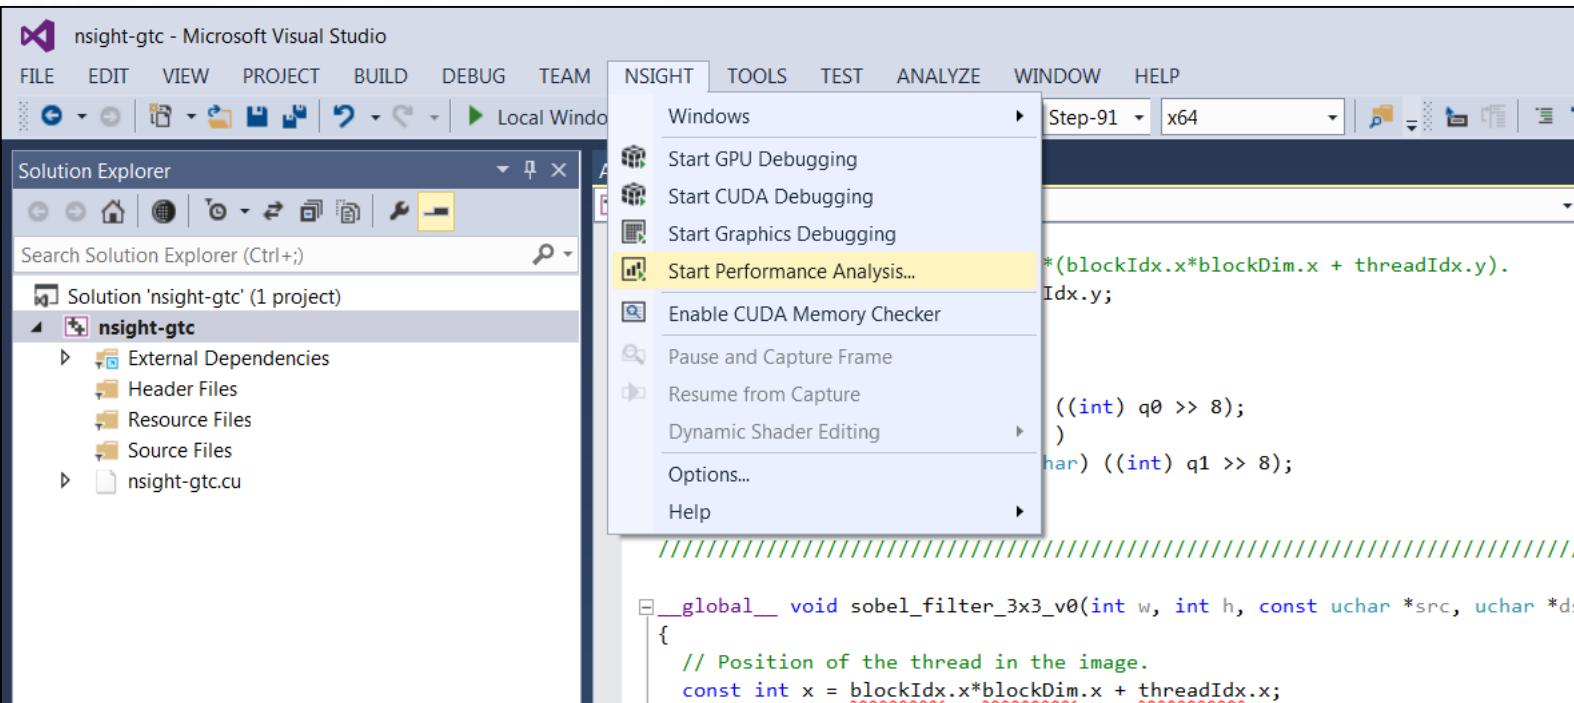
Lots of time spent loading into shared memory
2nd Highest latency is on memory

Lots of time spent computing convolution
Highest latency is on execution dependency

1st Analysis

- ❑ We have a reasonably well balanced use of the Compute and Memory pipes.
- ❑ There is some latency in loading data to and from shared memory
- ❑ Our compute cycles are dominated by Integer operations
 - ❑ What operations are they?
 - ❑ We can either examine the code and PTX instructions (from Compute or Latency Analysis) or run additional analysis via Nsight within Visual Studio
 - ❑ More detailed analysis
 - ❑ Not guided like the visual profiler

Start profiling



Kernel Analysis

The screenshot shows the 'Nsight Visual Studio Edition' interface for CUDA kernel analysis. It is divided into three main sections:

- Select Profile CUDA Application:** A box highlights the 'Activity Type' section, which contains four options:
 - Trace Application:** Collects events from the target application. The analysis session and data collection are stopped when the launched application exits.
 - Trace Process Tree:** Collects events from the target application and all native child processes of the target application. The analysis session and data collection are not stopped when the launched application exits. The session and data collection must be stopped manually.
 - Profile CUDA Application (selected):** Collects counters, statistics and derived values for given CUDA kernel launches.
 - Profile CUDA Process Tree:** Collects counters, statistics and derived values for given CUDA kernel launches from the target application and all native child processes of the target application. The analysis session and data collection are not stopped when the launched application exits. The session and data collection must be stopped manually.
- Select the Kernel (optional, will profile all kernels otherwise):** A box highlights the 'Experiment Settings' section, specifically the 'Kernel Selection' area where the kernel 'gaussian_filter_7x7_v0' is listed.
- Select the Experiments (All):** A box highlights the 'Experiment Configuration' section, where 'Experiments to Run' is set to 'All (Kernel-Level)'.

At the bottom right of the interface, there is a 'Launch' button with a red arrow pointing to it, indicating the final step to start the analysis session.

CUDA Launches View

Performance Indicators

Viewing: 3 / 3

Function Name	Grid Dimensions	Block Dimensions	Start Time (μs)	Duration (μs)	Occupancy	Registers per Thread	Static Shared Memory per Block (bytes)	Dynamic Shared Memory per Block (bytes)	Cache Configuration Executed	
1 gaussian_filter_7x7_v2	{80, 400, 1}	{32, 4, 1}	2,393,353.624	438.400	100.00 %	14	640	0	PREFER_SHARED	C
2 sobel_filter_3x3_v0	{80, 200, 1}	{32, 8, 1}	3,110,031.704	267.232	100.00 %	18	0	0	PREFER_SHARED	C
3 rgba_to_grayscale_kernel_v0	{80, 200, 1}	{32, 8, 1}	1,384,434.679	133.696	100.00 %	8	0	0	PREFER_SHARED	C

Kernel: gaussian_filter_7x7_v2 Grid Dim: {80, 400, 1} 32000 Block Dim: {32, 4, 1} 128
Device: GeForce GTX 980 Compute Capability: 5.2 Dyn Shm/Block: 0 Stat Shm/Block: 640

Variable Achieved Theoretical Device Limit

Occupancy Per SM

Active Blocks		16	32	
Active Warps	60.68	64	64	
Active Threads		2048	2048	
Occupancy	94.81 %	100.00 %	100.00 %	

Warp Statistics

Threads/Block		128	1024	
Warps/Block		4	32	
Block Limit		16	32	

Registers

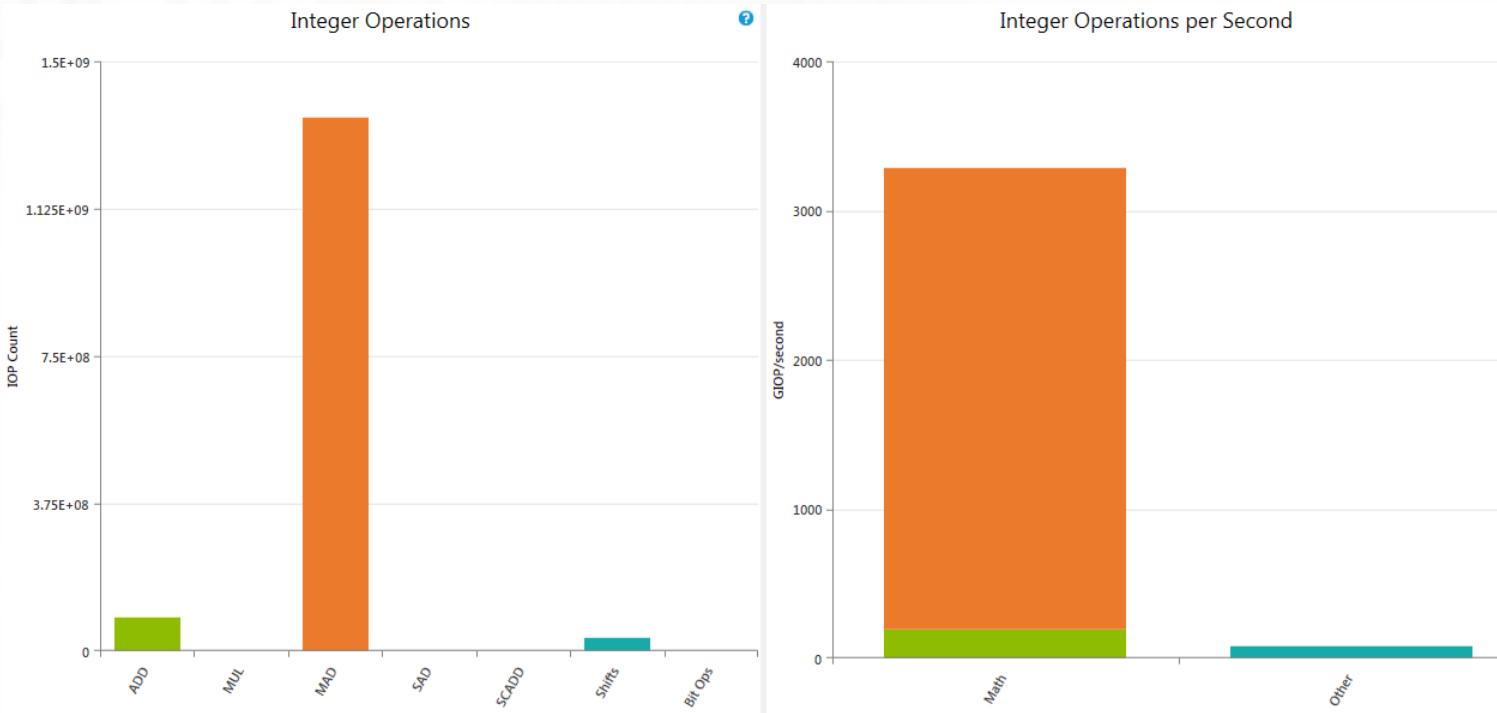
Registers		14	14	
-----------	--	----	----	--

Occupancy Data Occupancy Graphs

The left sidebar shows a tree view of performance metrics for the gaussian_filter_7x7_v2 kernel:

- gaussian_filter_7x7_v2 << <<32000,128>> [CUDA Launch]
 - Device Launches
 - Call Graph
- gaussian_filter_7x7_v2 [CUDA Kernel]
- Experiment Results
 - Occupancy
 - Instruction Statistics
 - Branch Statistics
 - Issue Efficiency
 - Achieved FLOPS
 - Achieved IOPS
 - Pipe Utilization
 - Memory Statistics
- Source Profiler
 - Instruction Count
 - Divergent Branch
 - Memory Transactions

Achieved IOPS



❑ No surprises...

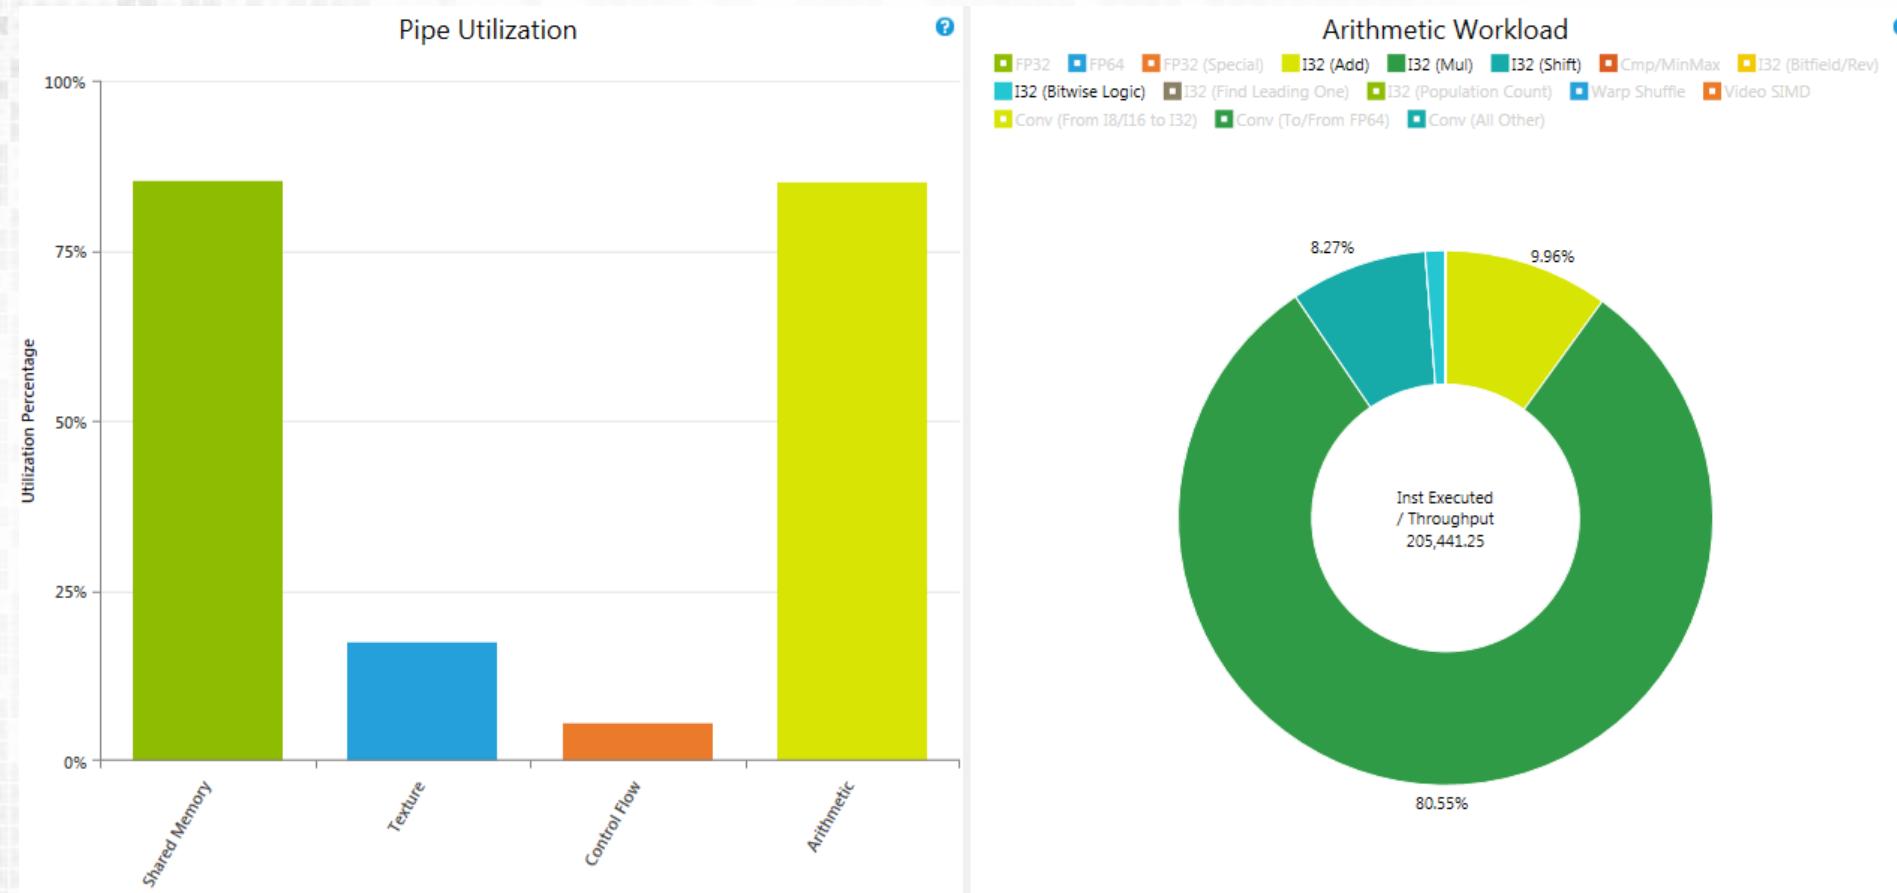
```
int p = 0;
for( int j = 0 ; j < 7 ; ++j )
    for( int i = 0 ; i < 7 ; ++i )
        p += gaussian_filter[j][i] * n[j][i];
```



The University
Of
Sheffield.

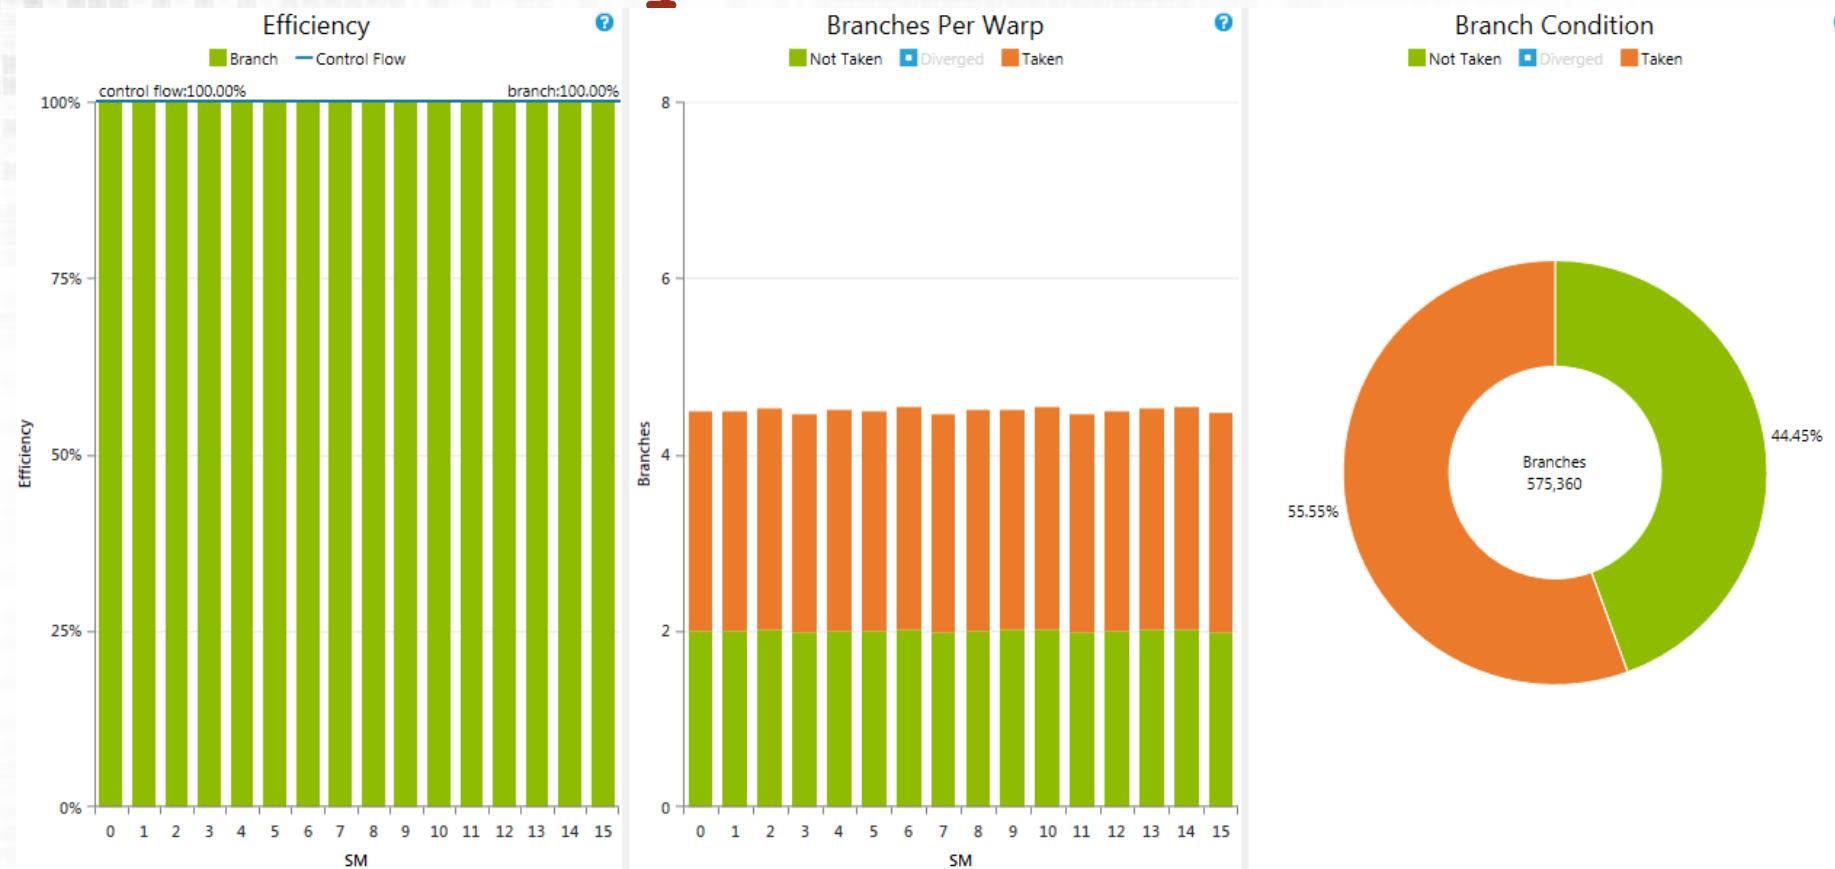
NVIDIA
GPU
RESEARCH
CENTER

Pipe Utilisation



- ❑ More detailed confirmation
- ❑ Integer operations dominate

Issue Efficiency



- This is good
- We have no divergent code

2nd Analysis

- ❑ We have a reasonably well balanced use of the Compute and Memory pipes.
- ❑ There is some latency in loading data to shared memory and on executions to read it back
- ❑ Our compute cycles are dominated by Integer operations

There is some latency in loading data to shared memory and on executions to read it back



- ❑ Consider a simplified problem
- ❑ Each thread needs to load an r, g, b, a value into shared memory
 - ❑ Which has fewer shared memory load instructions?

```
__shared__ char sm[TPB*4];  
  
char r,g,b,a;  
  
r = sm[threadidx.x];  
g = sm[threadidx.x+1];  
b = sm[threadidx.x+2];  
a = sm[threadidx.x+3];
```

```
__shared__ char4 sm[TPB];  
  
char r,g,b,a;  
char4 rgba;  
rgba = sm[threadidx.x];  
r = rgba.r;  
g = rgba.g;  
b = rgba.b;  
a = rgba.a;
```

There is some latency in loading data to shared memory and on executions to read it back

- Consider a simplified problem
- Each thread needs to load an r, g, b, a value into shared memory
 - Which has fewer shared memory load instructions?

```
__shared__ char sm[TPB*4];  
  
char r,g,b,a;  
  
r = sm[threadidx.x];  
g = sm[threadidx.x+1];  
b = sm[threadidx.x+2];  
a = sm[threadidx.x+3];
```

```
__shared__ char4 sm[TPB];  
  
char r,g,b,a;  
char4 rgba;  
rgba = sm[threadidx.x];  
r = rgba.r;  
g = rgba.g;  
b = rgba.b;  
a = rgba.a;
```



The
University
Of
Sheffield.



Our compute cycles are dominated by Integer operations



```
int p = 0;  
for( int j = 0 ; j < 7 ; ++j )  
    for( int i = 0 ; i < 7 ; ++i )  
        p += gaussian_filter[j][i] * n[j][i];
```

- Which of the following is faster?

```
int a, b, c;  
a = sm_a[i]; b = sm_b[i];  
  
c += a * b;
```

```
float a, b, c;  
a = sm_a[i]; b = sm_b[i];  
  
c += a * b;
```

Our compute cycles are dominated by Integer operations

```
int p = 0;  
for( int j = 0 ; j < 7 ; ++j )  
    for( int i = 0 ; i < 7 ; ++i )  
        p += gaussian_filter[j][i] * n[j][i];
```

- Which of the following is faster?

```
int a, b, c;  
a = sm_a[i]; b = sm_b[i];  
  
c += a * b;
```

```
float a, b, c;  
a = sm_a[i]; b = sm_b[i];  
  
c += a * b;
```

Integer multiply add is 16 cycles

Float combined multiply add is 4 cycles

Analysis

- ❑ We have a reasonably well balanced use of the from Compute and Memory pipes.
- ❑ There is some latency in loading data to shared memory and on executions to read it back
 - ❑ **Solution 1:** Reduce SM Load Stores dependencies by using wider requests.
i.e. 4B values rather than 1B (chars)
 - ❑ I.e. Store shared memory values as 4B minimum
- ❑ Our compute cycles are dominated by Integer operations
 - ❑ Almost all MAD operations
 - ❑ **Solution:** Change slower Integer MAD instructions to faster floating point FMAD instructions
 - ❑ I.e. Use floating point multiply and cast result to uchar at end

Improvement

❑ Significant

Kernel	Time (ms)	Speedup	Rel. Speedup
Gaussian_filter (Step 0)	5.49	1.00x	-
Gaussian_filter (Step 1a)	1.00	5.49x	5.49x
Gaussian_filter (Step 40)	0.49	11.20x	2.04x
Gaussian_filter (Step 5a)	0.28	19.60x	1.75x



The
University
Of
Sheffield.



NVIDIA
GPU
RESEARCH
CENTER

- ❑ Profiling Introduction
- ❑ The Problem
- ❑ Visual Profiler Guided Analysis
 - ❑ Iteration 1
 - ❑ Iteration 2
 - ❑ Iteration 3
 - ❑ Iteration 4



The
University
Of
Sheffield.





Identify the hotspot

- ❑ Examine GPU Usage in Visual Profiler
- ❑ What should be our next step?

i Kernel Optimization Priorities

The following kernels are ordered by optimization importance based on execution time and achieved occupancy. Optimization of higher ranked kernels (those that appear first in the list) is more likely to improve performance compared to lower ranked kernels.

Rank	Description
100	[1 kernel instances] sobel_filter_3x3_v0(int, int, unsigned char const *, unsigned char*)
93	[1 kernel instances] gaussian_filter_7x7_v3_bis(int, int, unsigned char const *, unsigned char*)
49	[1 kernel instances] rgba_to_grayscale_kernel_v0(int, int, uchar4 const *, unsigned char*)



Identify the hotspot

- ❑ Examine GPU Usage in Visual Profiler
- ❑ Examine Individual Kernels
 - ❑ Gaussian filter kernel no longer highest rank!
 - ❑ We can now optimise the sobel_filter kernel

i Kernel Optimization Priorities

The following kernels are ordered by optimization importance based on execution time and achieved occupancy. Optimization of higher ranked kernels (those that appear first in the list) is more likely to improve performance compared to lower ranked kernels.

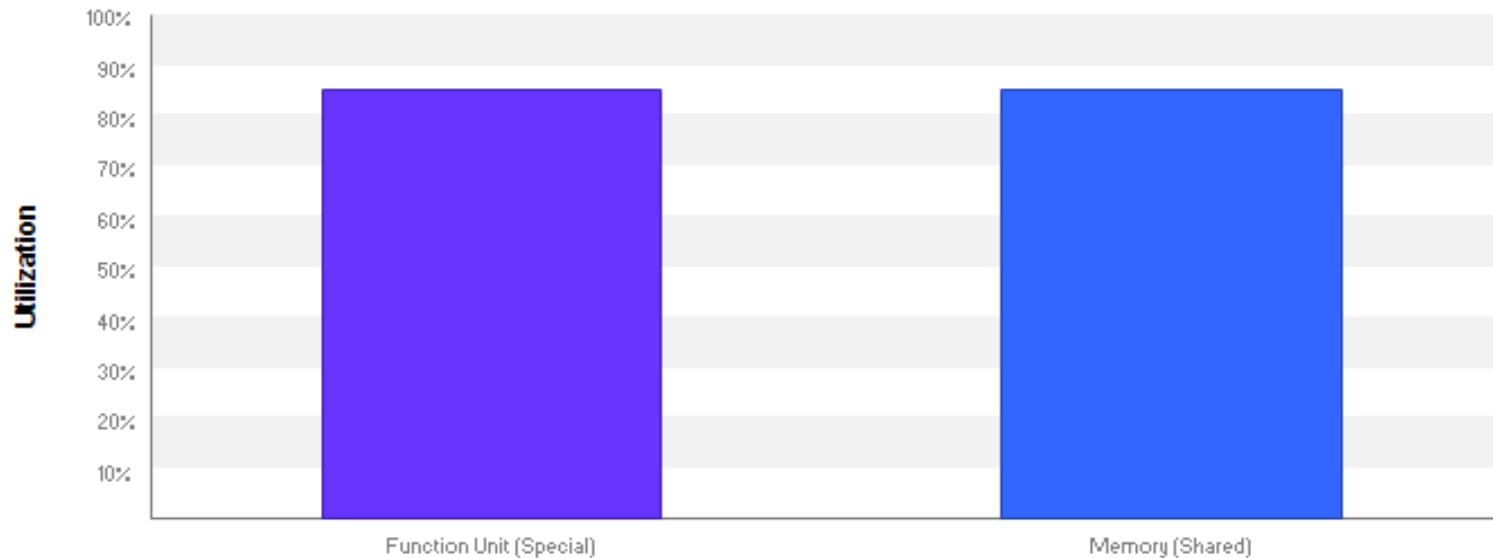
Rank	Description
100	[1 kernel instances] sobel_filter_3x3_v0(int, int, unsigned char const *, unsigned char*)
93	[1 kernel instances] gaussian_filter_7x7_v3_bis(int, int, unsigned char const *, unsigned char*)
49	[1 kernel instances] rgba_to_grayscale_kernel_v0(int, int, uchar4 const *, unsigned char*)

- ❑ Lets look at our Gaussian kernel anyway...

Performance Limiter

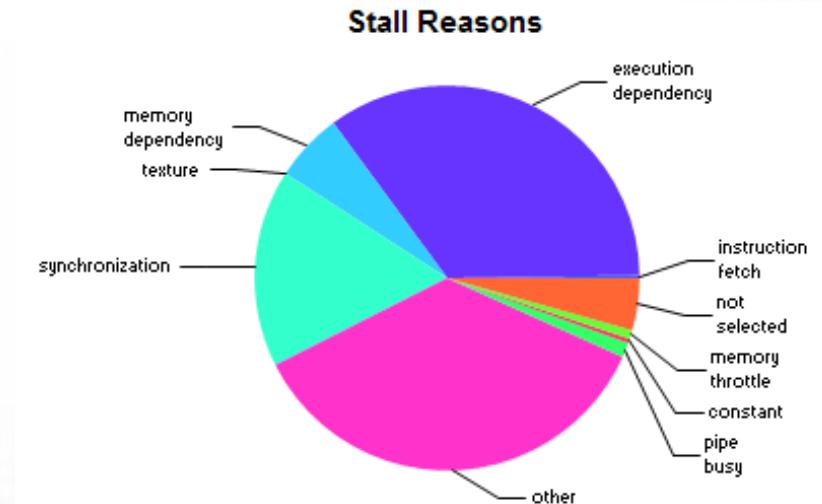
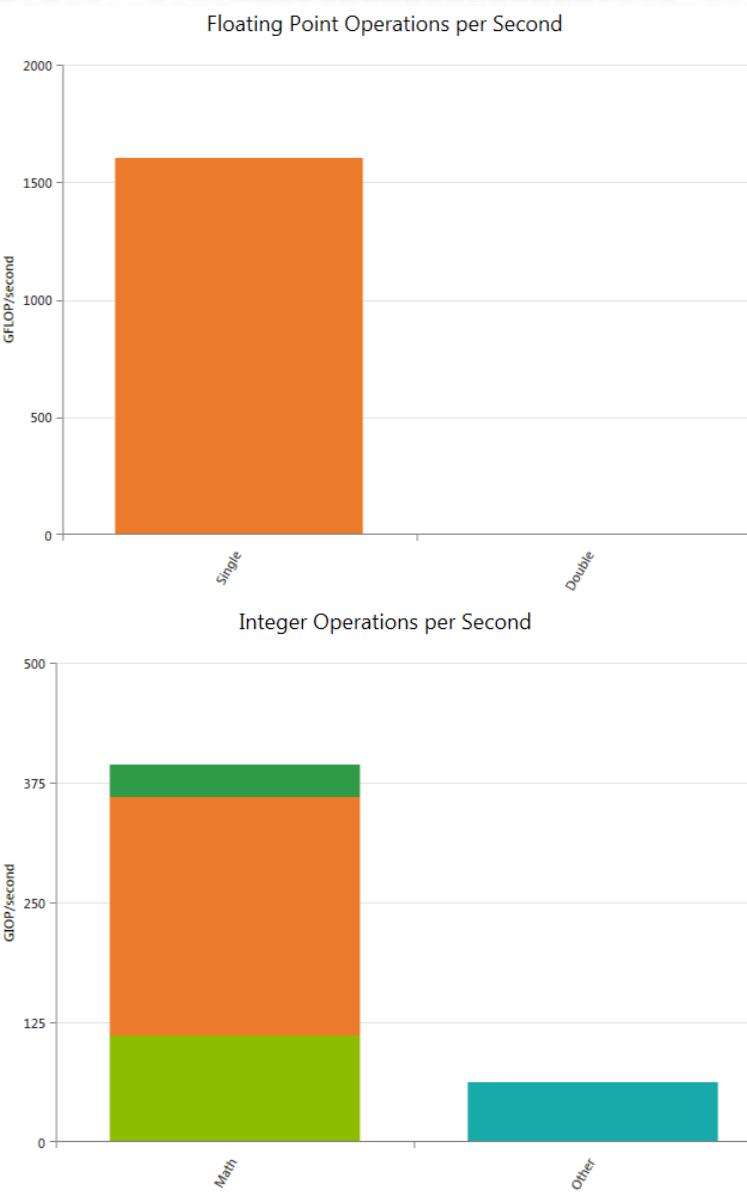
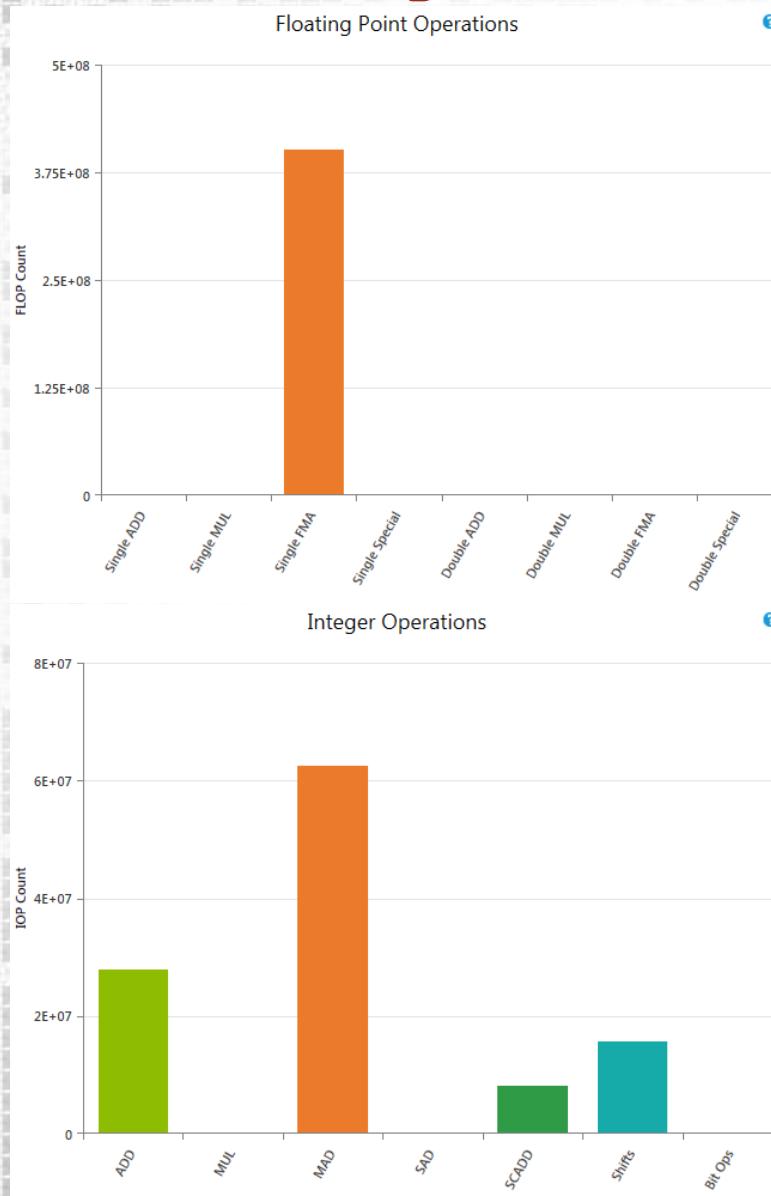
i High Compute And Memory Utilization

The kernel is utilizing greater than 80% of the available compute and memory performance of device "GeForce GTX 980". These utilization levels indicate that additional performance improvement may be difficult to achieve for the kernel.



Looking good

VS NSight IOPS/ FLOPS Metrics



Analysis

- ❑ Our algorithm is making good use of compute and memory
 - ❑ Further improvement will be difficult (but not impossible)
-
- ❑ **Solution:** Optimise a different kernel
 - ❑ sobel_filter_kernel to get the same treatment
 - ❑ **Solution:** Improve Gaussian kernel by changing the technique (parallelise differently)
 - ❑ Separable Filter: Compute horizontal and vertical convolution separately then approximate by binomial coefficients
 - ❑ Ensure we apply the same optimisations to separable filter version

Improvement

Kernel	Time (ms)	Speedup	Rel. Speedup
Gaussian_filter (Step 0)	5.49	1.00x	-
Gaussian_filter (Step 1a)	1.00	5.49x	5.49x
Gaussian_filter (Step 40)	0.49	11.20x	2.04x
Gaussian_filter (Step 5a)	0.28	19.60x	1.75x
Gaussian_filter (Step 9)	0.22	24.95x	1.27x

- ❑ 25x speedup on existing GPU code is pretty good
- ❑ Companion Code: <https://github.com/chmaruni/nsight-gtc>

Summary

- ❑ Profiling with the Visual Profiler will give you guided analysis of how to improve your performance
 - ❑ Show you how to spot key metrics
- ❑ We are trying to achieve good overall utilisation of the hardware (compute and memory engines)
 - ❑ Through an appreciation of memory and compute bounds
- ❑ Follow the APOD cycle
 - ❑ Assess: What is the limiting factor, analyse and profile
 - ❑ Parallelise and improve (apply the knowledge you have learnt over the course)
 - ❑ Optimise
 - ❑ Deploy and Test
- ❑ If in doubt use the lab classes to seek guidance!

Extra Drop-in Sessions/Assignment Help (Labs)

Drop-in session:

- Wed 27th March: 10:00-12:00**
- Thurs 4th April: 12:00-14:00**
- Wed 1st May: 13:00-15:00**
- Wed 8th May: 13:00-15:00**

Assignment help (during lab):

- Tues 7th May: 13:00-15:00 , 16:00-18:00**
- Mon 13th May: 13:00-15:00**

Location: Diamond high spec computer PC room 4

**Normal lab hours: Mondays (15:00-17:00) and Tuesdays (13:00-15:00)*