

Parallel Computing with GPUs: Sorting and Libraries

Dr Mozhgan Kabiri Chimeh

<http://mkchimeh.staff.shef.ac.uk/teaching/COM4521>



The
University
Of
Sheffield.



GPU
RESEARCH
CENTER

Last Week

- ❑ We learnt about Performance optimisation
- ❑ APOD cycle
- ❑ Use of guided analysis to find important kernels
- ❑ Use of guided analysis to find optimisation routes for code

❑ Sorting Networks

❑ Merge and Bitonic sort

❑ Thrust Parallel Primitives Library

❑ Applications of sorting (binning)

Serial Sorting Examples

❑ Insertion Sort

❑ Insert a new element into a sorted list.

❑ E.g. [1 6 3 4 2 5]

❑ [1] -> [1 6] -> [1 3 6] -> [1 3 4 6] -> [1 2 3 4 6] -> [1 2 3 4 5 6]

❑ Bubble Sort

❑ Exchange and Sweep to compare each pair of adjacent elements

❑ $O(n^2)$ worst-case and average case, $O(n)$ best case.

❑ E.g. [1 6 3 4 2 5]

❑ [1 6 3 4 2 5] -> [1 3 6 4 2 5] -> [1 3 4 6 2 5] -> [1 3 4 2 6 5] -> [1 3 4 2 5 6]

❑ [1 3 2 4 5 6]

❑ [1 2 3 4 5 6]

Classifying Sort Techniques/Implementations

☐ Data driven


- ☐ Each step of the algorithm depends on the previous step version
- ☐ Highly serial

☐ Data independent

- ☐ The algorithm performs fixed steps and does not change its processing based on data
- ☐ Well suited to parallel implementations
- ☐ Can be expressed as a sorting network...

Sorting Networks

- ❑ A sorting network is a comparator network that sorts all input sequences
- ❑ Following the same execution of stages
- ❑ Consider the previous Bubble Sort [1 6 3 4 2 5]



[1 6 3 4 2 5] -> [1 3 6 4 2 5] -> [1 3 4 6 2 5] -> [1 3 4 2 6 5] -> [1 3 4 2 5 6]
[1 3 4 2 5 6] -> [1 3 2 4 5 6] -> [1 3 2 4 5 6]
[1 3 2 4 5 6] -> [1 2 3 4 5 6]
[1 2 3 4 5 6]
[1 2 3 4 5 6]

Sweeps

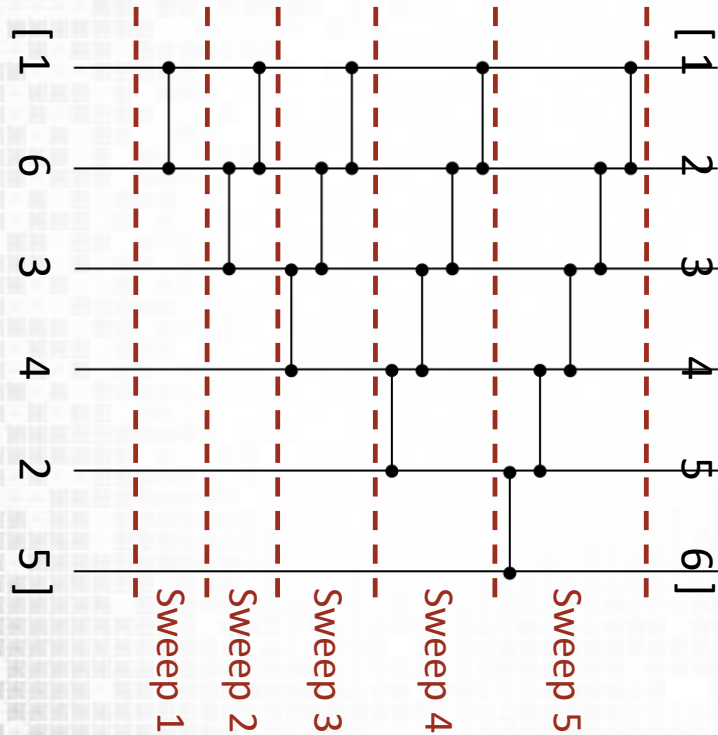
Not considered

Compared not swapped

Compared and swapped

Sorting Networks

□ And Insertion Sort...



$[1 \text{ 6 } 3 \text{ 4 } 2 \text{ 5}]$
 $[1 \text{ 3 } 6 \text{ 4 } 2 \text{ 5}] \rightarrow [1 \text{ 3 } 6 \text{ 4 } 2 \text{ 5}]$
 $[1 \text{ 3 } 4 \text{ 6 } 2 \text{ 5}] \rightarrow [1 \text{ 3 } 4 \text{ 6 } 2 \text{ 5}] \rightarrow [1 \text{ 3 } 4 \text{ 6 } 2 \text{ 5}]$
 $[1 \text{ 3 } 4 \text{ 2 } 6 \text{ 5}] \rightarrow [1 \text{ 3 } 2 \text{ 4 } 6 \text{ 5}] \rightarrow [1 \text{ 2 } 3 \text{ 4 } 6 \text{ 5}] \rightarrow [1 \text{ 2 } 3 \text{ 4 } 6 \text{ 5}]$
 $[1 \text{ 2 } 3 \text{ 4 } 5 \text{ 6}] \rightarrow [1 \text{ 2 } 3 \text{ 4 } 5 \text{ 6}] \rightarrow [1 \text{ 2 } 3 \text{ 4 } 5 \text{ 6}] \rightarrow [1 \text{ 2 } 3 \text{ 4 } 5 \text{ 6}] \rightarrow [1 \text{ 2 } 3 \text{ 4 } 5 \text{ 6}]$


Sweeps ↓

Not considered


Compared not swapped

Compared and swapped


Parallel Sorting Networks



Bubble



Insertion



[1 6 3 4 2 5]
[1 3 6 4 2 5]
[1 3 4 6 2 5]
[1 3 4 2 6 5]
[1 3 2 4 5 6]
[1 2 3 4 5 6]
[1 2 3 4 5 6]
[1 2 3 4 5 6]

Sweeps = 9

❑ Parallel Bubble and Insertion sorting network is still not very efficient



❑ $2n - 3$ sweeps

❑ $n(n - 1)/2$ comparisons - $O(n^2)$ complexity


- ❑ Sorting Networks
- ❑ Merge and Bitonic sort
- ❑ Thrust Parallel Primitives Library
- ❑ Applications of sorting (binning)

Merge Sort

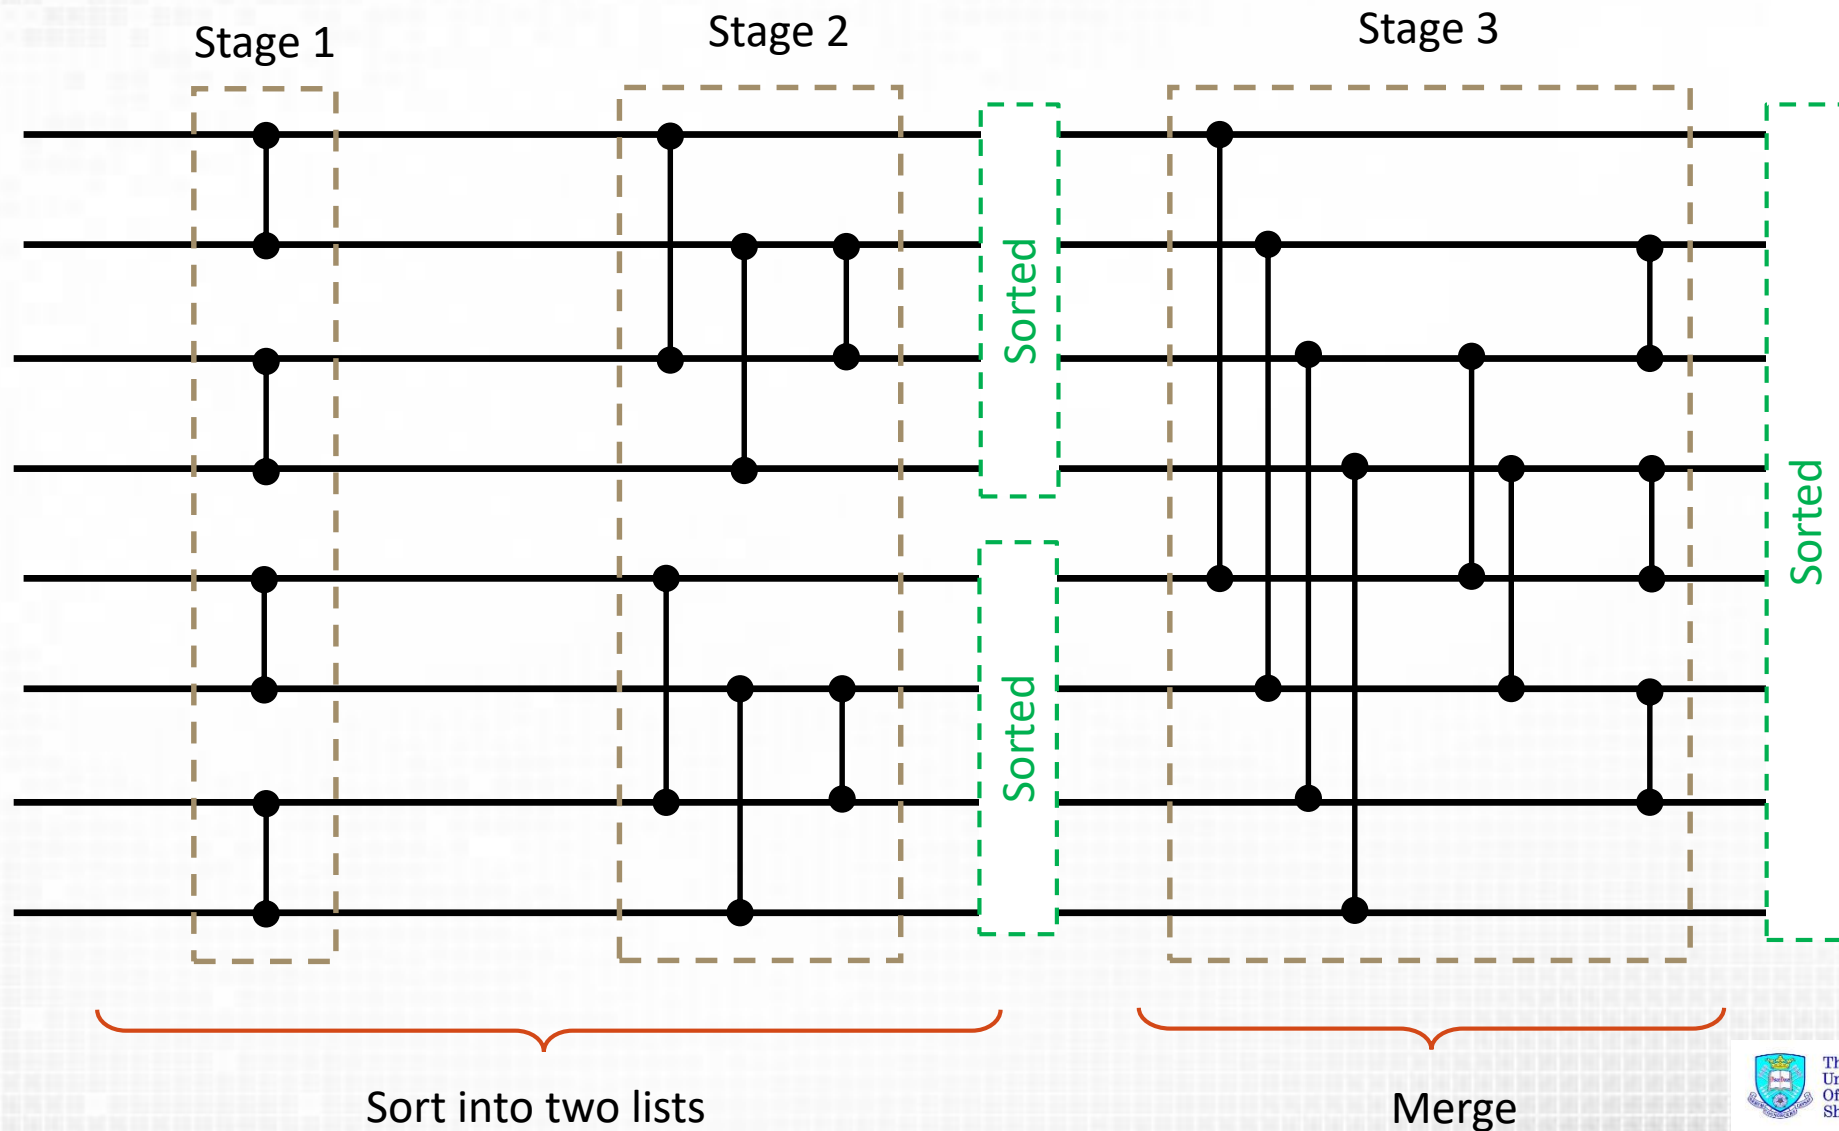
- ❑ To reduce the $O(n^2)$ overhead we need a better sorting network
- ❑ The odd-even merge sort network (for power of 2 n)
 - ❑ Sort all odd and even keys separately and then merge m values of a stage
 - ❑ Merge a sorted sequence of elements on lines $\langle a_1, \dots, a_n \rangle$ with those on lines $\langle a_{n+1}, \dots, a_{2n} \rangle$
 - ❑ Each merge requires $\log(n)$ passes
 - ❑ Total complexity of $O(n \log(n^2) + \log(n))$




Merge Sorting (n=8)



Merge Sorting (n=8)



Merge Sorting (n=8) example




Input	Stage 1		Stage 2		Stage 3						Output
8	1		1		1						1
1	8			5	3					2	2
5	3		3		5				2	3	3
3	5			8				8		4	4
6	2		2			2			5	5	5
2	6			6	4			4		8	6
4	4		4		6					8	8
9	9			9				9			9




Limitations of Merge Sort?

- ❑ What is potentially wrong with a merge sort GPU implementation?
 - ❑ Hint: Think about workload per thread



Limitations of Merge Sort

- ❑ What is potentially wrong with a merge sort GPU implementation?
 - ❑ Irregular memory accesses
 - ❑ Not all values are compared in each pass (uneven workload per thread)




Solution: Bitonic Sort

□ Bitonic sorting network


- Iterative splitting and merging of inputs into increasing large bitonic sequences

- A sequence is bitonic if

- There is an i , such that $a_0 \dots, a_i$ is monotonically increasing and $a_i \dots, a_n$ is monotonically decreasing



Bitonic Sorting Network



- ❑ Sorting and Merging increasing large bionic sequences

- ❑ When $n = 2^k$ there are k levels with $\frac{n}{2}$ comparisons each

- ❑ GPU Implementation

- ❑ Regular access strides :-)

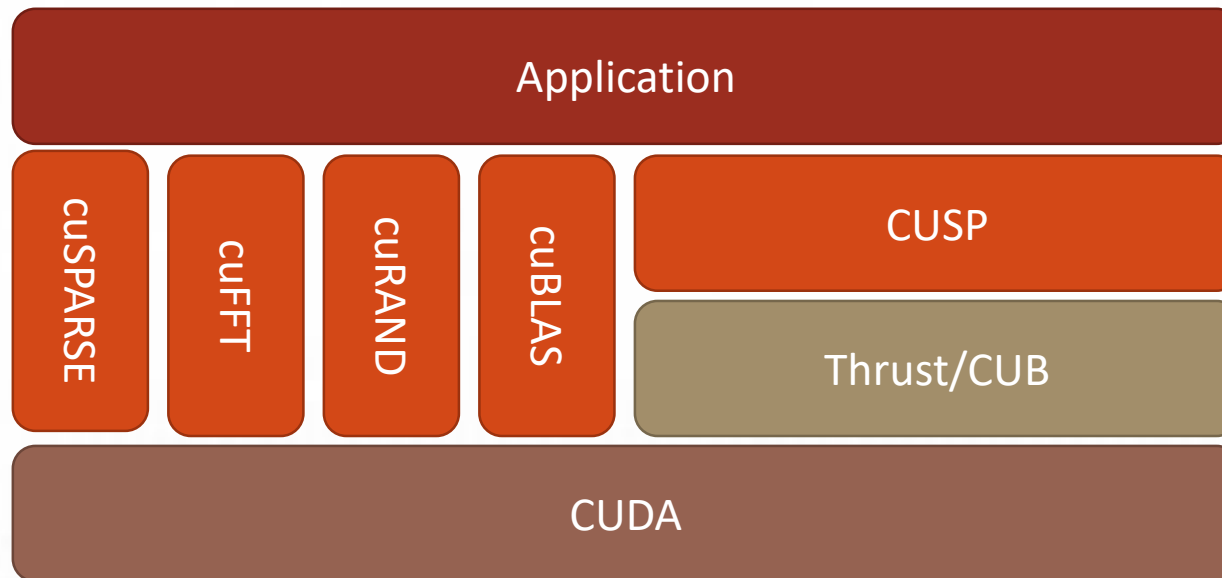
- ❑ Efficiently balanced workload :-)

- ❑ Requires multiple kernel launches to merge over $n > \text{block size}$

- ❑ Sorting Networks
- ❑ Merge and Bitonic sort
- ❑ Thrust Parallel Primitives Library
- ❑ Applications of sorting (binning)

CUDA libraries

- ❑ Abstract CUDA model away from programmer
- ❑ Highly optimised implementations of common tools
 - ❑ Mainly focused on linear algebra



Thrust

- ❑ Template Library for CUDA
 - ❑ Implements many parallel primitives (scan, sort, reduction etc.)
 - ❑ Part of standard CUDA release
 - ❑ Level of Abstraction which hides kernels, mallocs and memcpy's
- ❑ Designed for C++ programmers
 - ❑ Similar in design and operation as the C++ Standard Template Library (STL)
 - ❑ Only a small amount of C++ required..



Thrust containers

❑ Thrust uses only high level *vector* containers

❑ `host_vector`: on host

❑ `device_vector`: on GPU

❑ Other STL containers include

❑ `queue`

❑ `list`

❑ `tack`

❑ `queue`

❑ `priority_queue`

❑ `set`

❑ `multiset`

❑ `map`

❑ `multimap`

❑ `bitset`

❑ STL containers can be used to initialise a Thrust vector

```
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>

int main()
{
    //create a vector on the host
    thrust::host_vector<int> h_vec(10);

    //create a vector on the device
    thrust::device_vector<int> d_vec = h_vec;

    //device data manipulated directly from host
    for (int i = 0; i < 10; i++)
        d_vec[i] = i;

    //vector memory automatically released
    return 0;
}
```

Thrust Iterators

- ❑ They point to regions of a vector
- ❑ Can be used like pointers
 - ❑ Explicit cast when dereferencing very important

```
thrust::device_vector<int>::iterator begin = d_vec.begin();  
thrust::device_vector<int>::iterator end = d_vec.end();  
printf("d_vec at begin=%d", (int)*begin);  
begin++; //move on a single position  
printf("d_vec at begin++=%d", (int)*begin);  
*end = 88;  
printf("d_vec at end=%d", (int)*end);
```

```
d_vec at begin=0  
d_vec at begin++=1  
d_vec at end=88
```

Thrust Iterators

❑ Can be converted to a raw pointer

```
int * d_ptr = thrust::raw_pointer_cast(begin);  
int * d_ptr = thrust::raw_pointer_cast(begin[0]);  
  
kernel<BLOCKS, TPB>(d_ptr);
```

❑ Raw pointers can be used in Thrust

❑ BUT not exactly the same as a vector

```
int* d_ptr;  
cudaMalloc((void**)&d_ptr, N);  
thrust::device_ptr<int> d_vec = thrust::device_pointer_cast(d_ptr);  
//or  
thrust::device_ptr<int> d_vec = thrust::device_ptr<int>(d_ptr)  
cudaFree(d_ptr);
```

Thrust Algorithms

☐ Transformations

- ☐ Application of a function to each element within the range of a vector

☐ Reduction

- ☐ Reduction of a set of values to a single value using binary associative operator
- ☐ Can also be used to count occurrences of a value

☐ Prefix Sum

- ☐ Both inclusive and exclusive scans

☐ Sort

- ☐ Can sort keys or key value pairs

☐ Binary Search

- ☐ Position of a target value

Thrust Transformations

□ Some examples of the many transformations

```
thrust::copy(d_vec.begin(), d_vec.begin() + 10, d_vec_cpy.begin());
```

```
thrust::fill(d_vec.begin(), d_vec.begin() + 10, 0);
```

```
//rand is a predefined Thrust generator
```

```
thrust::generate(d_vec.begin(), d_vec.begin() + 10, rand);
```

```
// fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
```

```
thrust::sequence(d_vec.begin(), d_vec.begin() + 10);
```

```
//all occurrences of the value 1 are replaced with the value 10
```

```
thrust::replace(d_vec.begin(), d_vec.end(), 1, 10);
```

Thrust Algorithms

□ Either in-place or to output vector

```
thrust::device_vector<int> d_vec(10);
thrust::device_vector<int> d_vec_out(10);

//fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
thrust::sequence(d_vec.begin(), d_vec.begin() + 10);

//inclusive scan to output vector
thrust::inclusive_scan(d_vec.begin(), d_vec.end(),
d_vec_out.begin());

//inclusive scan in place
thrust::inclusive_scan(d_vec.begin(), d_vec.end(),
d_vec.begin());

//generate random data (actually a transformation)
thrust::generate(d_vec.begin(), d_vec.end(), rand);

//sort in place
thrust::sort(d_vec.begin(), d_vec.end());
```

Custom Transformations

```
thrust::device_vector<int> d_vec(10);
thrust::device_vector<int> d_vec_out(10);

//fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
d_vec = thrust::sequence(d_vec.begin(), d_vec.begin() + 10);

//declare a custom operator
struct add_5{
    __host__ __device__ int operator()(int a){
        return a + 5;
    }
};

add_5 func;

//apply custom transformation
thrust::transform(d_vec.begin(), d_vec.end(), d_vec_out.begin(), func);

//d_vec is now {5, 6, 7, 8, 9, 10, 11, 12, 13, 14}
```

Thrust Fusion

❑ For best performance it is necessary to fuse operations

```
struct absolute{
    __host__ __device__ int operator()(int a){
        return a < 0 ? -a : a ;
    }
};
absolute func;

//custom transformation to calculate absolute value
thrust::transform(d_vec.begin(), d_vec.end(), d_vec.begin(), func);
//apply reduction, maximum binary associate operator
int result = thrust::reduce(d_vec.begin(), d_vec.end(), 0, thrust::maximum<int>());
```


```
struct absolute{
    __host__ __device__ int operator()(int a){
        return a < 0 ? -a : a ;
    }
};
absolute func;

//apply transform reduction maximum binary associate operator
int result = thrust::transform_reduce(d_vec.begin(), d_vec.end(), func, 0, thrust::maximum<int>());
```

- ❑ Sorting Networks
- ❑ Merge and Bitonic sort
- ❑ Thrust Parallel Primitives Library
- ❑ Applications of sorting (binning)

Sorting and parallel primitives

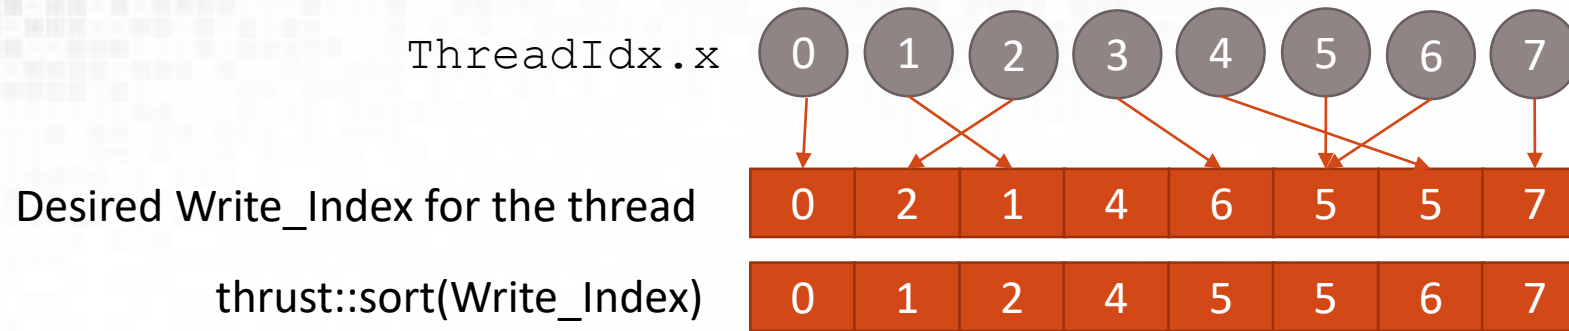
- ❑ Can be very useful for building data structures
 - ❑ We can use prefix sum for writing multiple values per element
- ❑ Remember Gather vs Scatter
 - ❑ What if our outputs are scattered to output
 - ❑ Very common in particle simulations etc.
 - ❑ Outputs might represent spatial bins



Scatter operation

- ❑ Write to a number of locations
 - ❑ Random access write?
- ❑ How to read multiple values afterwards?

Binning and Sorting



Build a data structure

Unique write indices	0	1	2	3	4	5	6	7
Count(Write_Index)	1	1	1	0	1	2	1	1
thrust::inclusive_scan(count)	0	1	2	3	3	4	6	7

← i.e. how many threads want to write to this index

- ❑ We can now read varying values from each bin

- ❑ E.g. for location 5

- ❑ inclusive_scan gives starting index of 4

- ❑ Iterate from index 4 for a count of 2 to find all values of write_index 5

Particle interaction example

- ❑ As with previous slide use sorting
 - ❑ Divide the environment according to some interaction radius
 - ❑ Output particle key value pairs (keys are location determined through some hash function)
 - ❑ Sort Keys
 - ❑ Reorder particles based on key pairs
 - ❑ Generate a partition boundary table
 - ❑ Histogram count and prefix sum
 - ❑ Each particle needs to read all particles in its own location and any neighbouring location
 - ❑ Guarantees particle interactions within the interaction radius

0	1	2 ¹ 2	3
3 ³ 4	5 ⁶ 5	6	7
8	9	10 ⁷	11
12	13 ⁸	14	15

Partition	First agent	Last agent
0		
1		
2	1	2
3		
4	3	4
5	5	6
6		
7		
8		
9		
10	7	7
11		
12		
13	8	8
14		
15		

Summary

- ❑ Sorting networks allow data independent sort algorithms to map easily parallel architectures
- ❑ Choice of a sorting network will dictate the memory access pattern and hence the performance on a GPU
- ❑ Merge sort and Bitonic sort are popular choices for GPUs
- ❑ Thrust implements many parallel primitives
- ❑ Thrust is based on the idea of containers, iterators, transformations and algorithms
- ❑ Sorting can be used to improve complex problems such as particle systems over a fixed range

Acknowledgements and Further Reading

- ❑ Comparison on sorting approaches on GPU

 - ❑ <http://arxiv.org/ftp/arxiv/papers/1511/1511.03404.pdf>

- ❑ <https://devblogs.nvidia.com/parallelforall/expressive-algorithmic-programming-thrust/>